

Second Edition

Coding and Information Theory

RICHARD W. HAMMING

Naval Postgraduate School

PRENTICE-HALL / Englewood Cliffs NJ 07632

Library of Congress Cataloging-in-Publication Data

Hamming, R. W. (Richard Wesley) 1915
Coding and information theory.

Bibliography: p.

Includes index.

1. Coding theory. 2. Information theory.

I. Title.

QA268.H35 1986 519.4 85-17034

ISBN 0-13-139072-4

Editorial/production supervision and
interior design: Nicholas C. Romanelli
Cover design: Frederick Charles Ltd.
Manufacturing buyer: Rhett Conklin

©1986, 1980 by Prentice-Hall
A Division of Simon & Schuster, Inc.
Englewood Cliffs, New Jersey 07632

All rights reserved. No part of this book may be
reproduced, in any form or by any means,
without permission in writing from the publisher.

Printed in the United States of America

10 9 8 7 6 5 4

ISBN 0-13-139072-4 01

Prentice-Hall International (UK) Limited, *London*
Prentice-Hall of Australia Pty. Limited, *Sydney*
Prentice-Hall Canada Inc., *Toronto*
Prentice-Hall Hispanoamericana, S.A., *Mexico*
Prentice-Hall of India Private Limited, *New Delhi*
Prentice-Hall of Japan, Inc., *Tokyo*
Prentice-Hall of Southeast Asia Pte. Ltd., *Singapore*
Editora Prentice-Hall do Brasil, Ltda., *Rio de Janeiro*
Whitehall Books Limited, *Wellington, New Zealand*

Contents

Preface

1 Introduction 1

- 1.1 A Very Abstract Summary 1
- 1.2 History 2
- 1.3 Model of the Signaling System 4
- 1.4 Information Source 5
- 1.5 Encoding a Source Alphabet 7
- 1.6 Some Particular Codes 9
- 1.7 The ASCII Code 10
- 1.8 Some Other Codes 12
- 1.9 Radix r Codes 15
- 1.10 Escape Characters 15
- 1.11 Outline of the Course 17

2 Error-Detecting Codes 20

- 2.1 Why Error-Detecting Codes? 20
- 2.2 Simple Parity Checks 21
- 2.3 Error-Detecting Codes 22
- 2.4 Independent Errors: White Noise 23
- 2.5 Retransmission of Message 25

2.6	Simple Burst Error-Detecting Codes	26
2.7	Alphabet Plus Number Codes: Weighted Codes	27
2.8	Review of Modular Arithmetic	30
2.9	ISBN Book Numbers	32

3 *Error-Correcting Codes* 34

3.1	Need for Error Correction	34
3.2	Rectangular Codes	35
3.3	Triangular, Cubic, and n -Dimensional Codes	37
3.4	Hamming Error-Correcting Codes	39
3.5	Equivalent Codes	42
3.6	Geometric Approach	44
3.7	Single-Error-Correction Plus Double-Error-Detection Codes	47
3.8	Hamming Codes Using Words	49
3.9	Applications of the Ideas	49
3.10	Summary	50

4 *Variable-Length Codes: Huffman Codes* 51

4.1	Introduction	51
4.2	Unique Decoding	52
4.3	Instantaneous Codes	53
4.4	Construction of Instantaneous Codes	55
4.5	The Kraft Inequality	57
4.6	Shortened Block Codes	60
4.7	The McMillan Inequality	62
4.8	Huffman Codes	63
4.9	Special Cases of Huffman Coding	68
4.10	Extensions of a Code	72
4.11	Huffman Codes Radix r	73
4.12	Noise in Huffman Coding Probabilities	74
4.13	Use of Huffman Codes	77
4.14	Hamming–Huffman Coding	78

5 ¹ *Miscellaneous Codes* 79

5.1	Introduction	79
5.2	What Is a Markov Process?	80

5.3	Ergodic Markov Processes	84
5.4	Efficient Coding of an Ergodic Markov Process	86
5.5	Extensions of a Markov Process	87
5.6	Predictive Run Encoding	88
5.7	The Predictive Encoder	89
5.8	The Decoder	90
5.9	Run Lengths	91
5.10	Summary of Predictive Encoding	94
5.11	What Is Hashing?	94
5.12	Handling Collisions	95
5.13	Removals from the Table	96
5.14	Summary of Hashing	96
5.15	Purpose of the Gray Code	97
5.16	Details of a Gray Code	98
5.17	Decoding the Gray Code	99
5.18	Anti-Gray Code	99
5.19	Delta Modulation	100
5.20	Other Codes	101

6 | Entropy and Shannon's First Theorem 103

6.1	Introduction	103
6.2	Information	104
6.3	Entropy	107
6.4	Mathematical Properties of the Entropy Function	114
6.5	Entropy and Coding	119
6.6	Shannon–Fano Coding	121
6.7	How Bad Is Shannon–Fano Coding?	122
6.8	Extensions of a Code	124
6.9	Examples of Extensions	126
6.10	Entropy of a Markov Process	129
6.11	Example of a Markov Process	131
6.12	The Adjoint System	133
6.13	The Robustness of Entropy	136
6.14	Summary	137

7 | The Channel and Mutual Information 138

7.1	Introduction	138
7.2	The Information Channel	139

- 7.3 Channel Relationships 140
- 7.4 Example of the Binary Symmetric Channel 142
- 7.5 System Entropies 145
- 7.6 Mutual Information 148

8 *Channel Capacity* 153

- 8.1 Definition of Channel Capacity 153
- 8.2 The Uniform Channel 154
- 8.3 Uniform Input 156
- 8.4 Error-Correcting Codes 158
- 8.5 Capacity of a Binary Symmetric Channel 159
- 8.6 Conditional Mutual Information 161

9 *Some Mathematical Preliminaries* 164

- 9.1 Introduction 164
- 9.2 The Stirling Approximation to $n!$ 165
- 9.3 A Binomial Bound 170
- 9.4 The Gamma Function $\Gamma(n)$ 173
- 9.5 n -Dimensional Euclidean Space 176
- 9.6 A Paradox 179
- 9.7 Chebyshev's Inequality and the Variance 181
- 9.8 The Law of Large Numbers 183
- 9.9 Other Metrics 189

10 *Shannon's Main Theorem* 191

- 10.1 Introduction 191
- 10.2 Decision Rules 192
- 10.3 The Binary Symmetric Channel 195
- 10.4 Random Encoding 196
- 10.5 Average Random Code 201
- 10.6 The General Case 204
- 10.7 The Fano Bound 205
- 10.8 The Converse of Shannon's Theorem 207

11 *Algebraic Coding Theory* 209

- 11.1 Introduction 209
- 11.2 Error-Detecting Parity Codes Revisited 210

11.3	Hamming Codes Revisited	211
11.4	Double-Error-Detecting Codes Revisited	213
11.5	Polynomials versus Vectors	214
11.6	Prime Polynomials	215
11.7	Primitive Roots	219
11.8	A Special Case	220
11.9	Shift Registers for Encoding	224
11.10	Decoding Single-Error-Correcting Codes	227
11.11	A Double-Error-Correcting Code	228
11.12	Decoding Multiple-Error Codes	234
11.13	Summary	234

Appendix A: Bandwidth and the Sample Theorem 237

A.1	Introduction	237
A.2	The Fourier Integral	238
A.3	The Sampling Theorem	240
A.4	Bandwidth versus Rapid Change	242
A.5	AM Signaling	242
A.6	FM Signaling	244
A.7	Pulse Signaling	245
A.8	Bandwidth Generally	246
A.9	Continuous Signals	247

Appendix B: Some Tables for Entropy Calculations 250

References 253

Index 255

Preface to the Second Edition

The second edition is mainly an expansion of the material in the first edition. The changes were made to (1) increase the clarity of the presentation, (2) extend some of the basic ideas, and (3) indicate further practical details. I have resisted the urgings (both from others and myself) to include more difficult codes, such as Reed–Muller, Viterbi, Fire, and cyclic. They seem to be more suitable for a second course, and in any case require a firmer mathematical background than I have assumed, if they are to be “understood.” I still believe that it is better, in this subject, to master the fundamentals than it is to rush through, with little real understanding, a large body of material. There are now many excellent texts devoted to algebraic coding theory.

Bostrom Management Corporation is thanked again for excellent cooperation.

R. W. HAMMING

Preface to the First Edition

This book combines the fields of coding and information theory in a natural way. They are both theories about the representation of abstract symbols. The two fields are now each so vast that only the elements can be presented in a short book.

Information theory is usually thought of as “sending information from here to there” (transmission of information), but this is exactly the same as “sending information from now to then” (storage of information). Both situations occur constantly when handling information. Clearly, the encoding of information for efficient storage as well as reliable recovery in the presence of “noise” is essential in computer science.

Since the representation, transmission, and transformation of information are fundamental to many other fields as well as to computer science, it is time to make the theories easily available. Whenever and wherever problems of generation, storage, or processing of information arise, there is a need to know both how to compress the textual material as well as how to protect it against possible mutilation. Of the many known encoding methods, we can indicate only a few of the more important ones, but hopefully the many examples in the text will alert the student to other possibilities.

The text covers the fundamentals of the two fields and gives examples of the use of the ideas in practice. The amount of background mathematics and electrical engineering is kept to a minimum. The book

uses, at most, simple calculus plus a little probability theory, and anything beyond that is developed as needed. Techniques that have recently arisen in computer science are used to simplify the presentation and the proofs of many results. These techniques are explained where they are used, so no special knowledge of computer science is required. Many other proofs have been greatly simplified, and when necessary new material has been developed to meet current technological needs. An effort has been made to arrange the material, especially the proof of Shannon's main result, so that it is evident *why* the theorems are true, not just that they have been proved mathematically.

Chapter 11, on algebraic codes, develops the needed mathematics of finite fields. Because of its mathematical difficulty, this chapter is placed last and out of logical order. It can follow Chapter 3, if desired. There is a deliberate repetition in the text; important ideas are usually presented at least twice to ensure that the reader understands them.

The text leaves out large areas of knowledge in the belief that it is better to master a little than to half know a lot. Thus more material may easily be added (at the discretion of the teacher) when it seems appropriate for the class.

I have followed custom in referring to *Hamming codes* and *Hamming distance*; to do otherwise would mislead the student and be false modesty.

Acknowledgments

It is difficult for the author to recall all his indebtedness to others, since he has known about the material from many years of working at the Bell Laboratories. Teaching a course at the Naval Postgraduate School, based on N. Abramson's elegant, small book, *Information Theory and Coding* (Ref. [A]), rearoused this author's interests in the two fields. Short courses on the topic at other places further developed many of the simplifications, elaborations, and examples, and the help of many students is gratefully acknowledged. The help of Albert Wong is especially appreciated.

In the actual production of the book, thanks are due to Bostrom Management Corporation, especially to Doug Thompson. However, as always, all faults are to be assigned to the author.

R. W. HAMMING

Introduction

1.1 A Very Abstract Summary

Although the text uses the colorful words “information,” “transmission,” and “coding,” a close examination will reveal that all that is actually assumed is an information source of symbols $s_1, s_2, \dots, s_q$. At first nothing is said about the symbols themselves, nor of their possible meanings. All that is assumed is that they can be uniquely recognized.

We cannot define what we mean by a symbol. We communicate with each other in symbols, written or spoken (or even gestures), and we may or may not think in terms of symbols, but since we must use symbols to define what we mean by a symbol, we see the circularity of the process. Thus the meaning of a symbol must remain at the intuitive level.

Next we introduce (Sections 4.8 and 6.2) the probabilities $p_1, p_2, \dots, p_q$ of these symbols occurring. How these p_i are determined is not part of the abstract theory. One way to get estimates of them is to examine past usage of the symbol system and hope that the future is not significantly different from the past. For any discrete probability distribution there is the value of the *entropy function*:

$$H = \sum_{i=1}^q p_i \log \frac{1}{p_i}$$

The function H of the probability distribution p_i measures the amount

of uncertainty, surprise, or information that the distribution contains. This function plays a leading role in the theory and provides a lower bound on the average code length. Later we examine more complex probability structures involving the symbols s_i .

The problem of representing the *source alphabet* symbols s_i in terms of another system of symbols (usually the binary system consisting of the two symbols 0 and 1) is the main topic of the book. The two main problems of representation are the following.

1. *Channel encoding*: How to represent the source symbols so that their representations are far apart in some suitable sense. As a result, in spite of small changes (noise) in their representations, the altered symbols can, at the receiving end, be discovered to be wrong and even possibly corrected. This is sometimes called “feed forward” error control.
2. *Source encoding*: How to represent the source symbols in a minimal form for purposes of efficiency. The average code length

$$L = \sum_{i=1}^q p_i l_i$$

is minimized where l_i is the length of the representation of the i th symbol s_i . The entropy function provides a lower bound on L .

Thus, in principle, the theory is simply an abstract mathematical theory of the representation of some undefined source symbols in terms of a fixed alphabet (usually the binary system) with the representation having various properties. In this abstract theory there is no transmission through the channel, no storage of information, and no “noise is added to the signal.” These are merely colorful words used to motivate the theory. We shall continue to use them, but the reader should not be deceived; ultimately this is merely a theory of the representation of symbols. Clearly, this is a fundamental theory for much of human symbol manipulation.

1.2 History

The beginnings of both coding and information theory go far back in time. Many of the fundamental ideas were understood long before 1948, when information theory was first established on a firm basis. In 1948 Claude E. Shannon published two papers on “A Mathematical Theory of Communication” in the *Bell System Technical Journal* (re-

printed in Ref. [S]). The papers almost immediately popularized the field of information theory, and soon additional papers on information theory appeared in the journals, and courses were taught on the subject in electrical engineering and other departments of many universities.

As in most fields that suddenly open up, many of the early applications were ill advised; but how else are the limitations of a new field to be discovered? As a result of the overexpectations of what information theory could do, disenchantment gradually set in, and a consequent decrease in the number of courses taught. Now, perhaps, a more just evaluation can be made, somewhere between the wild enthusiasm of the first days and the sad disappointment that slowly followed.

Information theory sets bounds on what can be done, but does little to aid in the design of a particular system. The idea that it is therefore useless to know information theory is false, as the following analogy shows. Consider the theory of evolution as taught in biology. Few students will ever apply it directly in their lives, yet it is a valuable constellation of ideas. In spite of this lack of direct application, when the ideas behind

1. Small changes in the species (variations)
2. Survival of the fittest (selection)

are mastered, the ideas can profitably be used in many other situations, often far removed from biology. For example, when looking at a social institution, such as a computer science department, a university, a military organization, a banking system, a government, or even a family relationship, one asks: “How did the present situation arise?” and “What forces selected this particular realization for survival?”

A little more appreciation of the power of the theory suggests the questions: “Given the present forces on the social institution, what are its possible variations (its ability to respond to the forces)?” and “How will it evolve (what will survive)?” Thus the ideas in the theory of evolution can be used in many situations far removed from biology.

When asking “How will it evolve?” it is often worth the effort to consider three similar questions:

- What can evolve?
- What should evolve?
- What will evolve?

With the answers in mind, you are in a position to consider how to bring the third answer more into line with the second—how to improve things,

rather than merely settling for the inevitable that will occur if you do nothing active to alter the world.

Similarly, information theory has ideas that are widely applicable to situations remote from its original inspiration. The applicability of the ideas is often not exact—they are often merely suggestive—but the ideas are still very useful.

At about the time that information theory was created, and in about the same place, coding theory was also created. The basic paper, however, was delayed by patent requirements until April 1950, when it appeared in the *Bell System Technical Journal* (reprinted in Refs. [B1], [B1a], and [T]). In the case of coding theory, the mathematical background was at first less elaborate than that for information theory, and for a long time it received less attention from the theorists. With the passing of time, however, various mathematical tools, such as group theory, the theory of finite fields (Galois theory), projective geometry, and even linear programming have been applied to coding theory. Thus coding theory has now become an active part of mathematical research (Refs. [B1], [B3], [B4], [C], [Gu], [L1], [L2], [MS], [Mc], [P], [W]).

Most bodies of knowledge give errors a secondary role, and recognize their existence only in the later stages of design. Both coding and information theory, however, give a central role to errors (noise) and are therefore of special interest, since in real life noise is everywhere.

Logically speaking, coding theory leads to information theory, and information theory provides bounds on what can be done by suitable encoding of the information. Thus the two theories are intimately related, although in the past they have been developed to a great extent quite separately. One of the main purposes of this book is to show their mutual relationships. For further details on the history of coding theory, see Refs. [B2] and [T]. By treating coding theory first, the meaning of the main results of information theory, and the proofs, are much clearer to the student.

1.3 Model of the Signaling System

The conventional signaling system is modeled by:

1. An information source
2. An encoding of this source
3. A channel over, or through, which the information is sent
4. A noise (error) source that is added to the signal in the channel
5. A decoding and hopefully a recovery of the original information from the contaminated received signal
6. A sink for the information

This model, shown in Figure 1.3-1, is the one that we will use for our signaling system. It has many features of signaling systems now in use.




Figure 1.3-1 Standard signaling system

The boxes ENCODE and DECODE should both be thought of as being divided into two parts. In the ENCODE box, we first encode the *source*, taking advantage of any structure in the symbol stream. The more structure there is, the more, typically, we can compress the encoding. These encoded symbols are then further encoded to compensate for any known properties of the *channel*. Typically, this second stage of encoding expands the representation of the message. Thus the ENCODE box first does *source encoding*, followed by *channel encoding*. The DECODE box must, of course, reverse these encodings in the proper order to recover the original symbols.

This separation into two stages is very convenient in practice. The various types of sources are separately encoded into the standard interface between the source and the channel, and then are encoded for the particular channel that is to be used. Thus there is great flexibility in the entire system.

1.4 Information Source

We begin, therefore, with the *information source*. The power of both coding and information theory is to a great extent due to the fact that we *do not define* what information is—we *assume* a source of information, a sequence of symbols in a *source alphabet* $s_1, s_2, \dots, s_q$, having q symbols. When we get to Chapter 6 we will find that information theory uses the entropy function H as a *measure of information*, and by implication this defines what is meant by “the amount of information.” But this is an abstract definition that *agrees only partly* with the commonly accepted ideas concerning information. It is probably this ap-

pearance of treating what we mean by “information” that made information theory so popular in the early days; people did not notice the differences and thought that the theory supplied the proper meaning in all cases. The implied definition does give the proper measure to be used for information in many situations (such as the storage and transmission of data), and people hoped that their situation was one of these (they were often wrong!). Information theory *does not* handle the *meaning* of the information; it treats only the *amount* of information, and even the amount depends on how you view the source.

The source of information may be many things; for example, a book, a printed formal notice, and a company financial report are all information sources in the conventional alphabetic form. The dance, music, and other human activities have given rise to various forms (symbols) for representing their information, and therefore can also be information sources. Mathematical equations are still another information source. The various codes to which we turn next are merely particular ways of representing the information symbols of a source.

Information also exists in continuous forms; indeed, nature usually supplies information to us in that form. But modern practice is to sample the continuous signal at equally spaced intervals of time, and then to digitize (quantize; Ref. [J]) the amount observed. The information is then sent as a stream of digits. Much of the reason for this use of digital samples of the analog signal is that they can be stored, manipulated, and transmitted more reliably than can the analog signal. When the inevitable noise of the transmission system begins to degrade the signal, the digital pulses can be sensed (detected), reshaped, and amplified to standard form *before* relaying them down the system to their final destination. At the destination the digital pulses may, if necessary, be converted back to analog form. Analog signals cannot be so reshaped, and hence the farther the signal is sent and the more it is processed, the more degradation it suffers from small errors.

A second reason that modern systems use digital methods is that large-scale integrated circuits are now very cheap and provide a powerful method for flexibly and reliably processing and transforming digital signals.

Although information theory has a part devoted to analog (continuous) signals, we shall concentrate on digital signals both for simplicity of the theory and because, as noted above, analog signals are of decreasing importance in our technical society. Almost all of our large, powerful computers are now digital, having displaced the earlier analog machines almost completely in information-processing situations. There are almost no large, general-purpose hybrid systems left (1985). Most of our information transmission systems, including the common telephone and hi-fi music systems, are also rapidly going to digital form.

1.5 Encoding a Source Alphabet

It is conventional to represent information (digital signals or more simply symbols) as being in one of two possible states: a switch up or down, on or off, a hole punched or not, a relay closed or open, a transistor conducting or not conducting, a magnetic domain magnetized N-S or S-N, and so on. Currently, devices with two states, called *binary devices*, are much more reliable than are multistate devices. As a result, binary systems dominate all others. Even decimal information-processing systems, such as hand calculators, are usually made from binary parts.

It is customary to use the symbols “0” and “1” as the names of the two states, but any two distinct symbols (marks), such as a circle and a cross, will do. It is often useful to think of the 0 and 1 as only a pair of arbitrary symbols, not as numbers.

First we consider merely the problem of representing the various symbols of the source alphabet. Given two binary (two-state) devices (digits), we can represent four distinct states (things, symbols):

00
01
10
11

For three binary digits we get $2^3 = 8$ distinct states:

000	100
001	101
010	110
011	111

For a system having k binary digits—usually abbreviated as *bits*—the total number of distinct states is, by elementary combinatorial theory,

$$2^k$$

In general, if we have k different independent devices, the first having n_1 states, the second n_2 states, . . . , the k th having n_k states, then the total number of states is clearly the product

$$n_1 n_2 \cdots n_k$$

For example, if $n_1 = 4$, $n_2 = 2$, and $n_3 = 5$, we can represent $4 \times 2 \times 5 = 40$ distinct items (states). This provides an upper bound on the number of possible representations.

This is the maximum number of source symbols we can represent *if all we consider* is the number of distinct states. In Chapter 4 we discuss how, using variable-length codes, to take advantage of the *probabilities* of the source symbols occurring. For the present we merely consider the number of distinct states—and this is the same as assuming that all the source symbols are equally likely to occur.

We have emphasized two-state devices with two symbols; other multistate devices exist, such as a “dead-center switch,” with three states. Some signaling systems also use more than two states, but the theory is most easily presented in terms of two states. We shall therefore generally use a binary system in the text and only occasionally mention systems with r states.

Human beings often attribute meanings to the sequences of 0's and 1's, for example in the ASCII code (Table 1.7-1). The computer (or signaling system), however, merely regards them as sequences of 0's and 1's. In particular, a digital computer is a processor of streams of the two symbols—it is the user (or possibly the input or output equipment) that assigns meanings; the computer merely combines 0's and 1's according to how it is built and how it is programmed. *The logic circuits of a computer are indifferent to the meanings we assign to the symbols; so is a signaling system.*

This is the reason that information theory ignores the meaning of the message, and by so doing it enables us to understand what the equipment does to messages. The theory provides an intellectual tool for understanding the *processing* of information without paying attention to meaning.

Typically, the channel encoding of the message increases the redundancy (to be defined accurately later), as shown in Chapters 2 and 3. The source encoding (mentioned at the end of Section 1.3) usually decreases the redundancy, as shown in Chapters 4 and 5.

We need to think of the source as a random, or stochastic, source of information, and ask how we may encode, transmit, and recover the original information. Specific messages are, of course, actually sent, but the designer of the system has no way of knowing which of the *ensemble* of possible messages will be chosen to be sent. The designer must view the particular message to be sent as a random sample from the population of all possible messages, and must design the system to handle any one of the possible messages. *Thus the theory is essentially statistical*, although we will use only elementary statistics in this book.

Exercises

1.5-1 Compute the number of license plates of the form
 number number number letter letter letter Ans. $(26)^3(10)^3$

1.5-2 How many instructions can there be in a computer with eight binary digits for each instruction? Ans. $2^8 = 256$

1.5-3 Discuss the biquinary system, which uses alternate two- and five-state devices.

1.6 Some Particular Codes

The binary code is awkward for human beings to use. Apparently, people prefer to make a single discrimination among many things. Evidence for this is the size of the usual alphabets, running from about 16 to 36 different letters (in both upper- and lowercase) as well as the decimal system, with 10 distinct symbols. Thus for human use it is often convenient to group the binary digits, called *bits*, into groups of three at a time, and call them the *octal code* (base 8). This code is given in Table 1.6-1.

TABLE 1.6-1 *Octal Code*

Binary	Octal
000	0
001	1
010	2
011	3
100	4
101	5
110	6
111	7

In the octal representation numbers are often enclosed in parentheses with a following subscript 8. For example, the decimal number 25 is written in octal as

$$(31)_8$$

Thus in America, Christmas is Halloween:

$$\text{Dec } 25 = \text{Oct } 31$$

$$(25)_{10} = (31)_8$$

As an example, in Table 1.7-1 of the ASCII code in the next section we have written the octal digits in the left-hand column rather than the

binary digits. The translation from octal to binary is so immediate that there is little trouble in going either way.

Occasionally, the binary digits are grouped in fours to make the *hexadecimal code* (Table 1.6-2). Since computers usually work in *bytes*, which are usually 8 bits each (there are now 9-bit bytes in some computers), the hexadecimal code fits into the machine architecture better than does the octal code, but the octal seems to fit better into human psychology. Thus neither code has a clear victory over the other in practice.

TABLE 1.6-2 *Hexadecimal Code*

Binary	Hexadecimal
0000	0
0001	1
0010	2
0011	3
0100	4
0101	5
0110	6
0111	7
1000	8
1001	9
1010	A
1011	B
1100	C
1101	D
1110	E
1111	F

Exercises

1.6-1 Since $2^8 \approx 3^5$, compare base 2 and base 3 computers.

1.6-2 Make a multiplication table for octal numbers.

1.6-3 From Table 1.6-2, what is the binary representation of D6?

Ans. 11010110

1.7 The ASCII Code

Given an information source, we first consider an encoding of it. The standard ASCII code (Table 1.7-1), which represents alphabetic, numeric, and assorted other symbols, is an example of a code. Basically,

TABLE 1.7-1 *Seven-Bit ASCII Code*

Octal Code	Char-acter	Octal Code	Char-acter	Octal Code	Char-acter	Octal Code	Char-acter
000	NUL	040	SP	100	@	140	`
001	SOH	041	!	101	A	141	a
002	STX	042	"	102	B	142	b
003	ETX	043	#	103	C	143	c
004	EOT	044	\$	104	D	144	d
005	ENQ	045	%	105	E	145	e
006	ACK	046	&	106	F	146	f
007	BEL	047	'	107	G	147	g
010	BS	050	(	110	H	150	h
011	HT	051	)	111	I	151	i
012	LF	052	*	112	J	152	j
013	VT	053	+	113	K	153	k
014	FF	054	,	114	L	154	l
015	CR	055	-	115	M	155	m
016	SO	056	.	116	N	156	n
017	SI	057	/	117	O	157	o
020	DLE	060	0	120	P	160	p
021	DC1	061	1	121	Q	161	q
022	DC2	062	2	122	R	162	r
023	DC3	063	3	123	S	163	s
024	DC4	064	4	124	T	164	t
025	NAK	065	5	125	U	165	u
026	SYN	066	6	126	V	166	v
027	ETB	067	7	127	W	167	w
030	CAN	070	8	130	X	170	x
031	EM	071	9	131	Y	171	y
032	SUB	072	:	132	Z	172	z
033	ESC	073	;	133	[	173	{
034	FS	074	<	134	\	174	
035	GS	075	=	135	]	175	}
036	RS	076	>	136	^	176	~
037	US	077	?	137	—	177	DEL

this code uses 7 binary digits (bits). Since (as noted already) computers work in *bytes*, which are usually blocks of 8 bits, a single ASCII symbol often uses 8 bits. The eighth bit can be set in many ways. It is usually set so that the total number of 1's in the eight positions is an even number (or else an odd number—see Chapter 2). Sometimes it is always set as a 1 so that it can be used as a timing source. Finally, it may be left arbitrary and no use made of it. To convert to the modified ASCII code used by the LT33 8-bit teletype code, use 7-bit ASCII code + (200)₈.

The purpose of the even number of 1's in the eight positions is that then any single error, a 0 changed into a 1 or a 1 changed into a 0, will be detected, since after the change there will be, in total, an odd number of 1's in the whole eight positions. Thus we have an error-detecting code that gives some protection against errors. *Perhaps more important*, the code enables the maintenance to be done much more easily and reliably since the presence of errors is determined by the machine itself, and to some extent the source of the errors can actually be located by the code.

We shall frequently use the check of an even (or odd) number of 1's. It is called a *parity check*, since what is checked is only the parity (the evenness or oddness) of the number of 1's in the message. Many computers have a very useful instruction that computes the parity of the contents of the accumulator.

We are now in a better position to understand the ASCII code (Table 1.7-1). The ASCII has a source alphabet of

$$2^7 = 128$$

possible characters (symbols). These characters are represented (encoded) inside a computer in the binary code. An even-parity check can be used to set the eighth position of the 8-bit bytes of the ASCII code. The three printed symbols of Table 1.7-1 are in the octal code. As an example

$$127 = 1 \ 010 \ 111$$

(where we have dropped the first 2 bits of the first octal symbol). For an even-parity check this would be $127 = 11 \ 010 \ 111$.

Exercises

1.7-1 Write the letters P and p in binary.

1.7-2 To what does 01 010 010 correspond?

1.7-3 Using base 4 code, write out the lowercase ASCII alphabet.

1.7-4 Write out the uppercase ASCII alphabet in base 16.

1.8 Some Other Codes

Another familiar code is the *Morse code*, which was once widely used. Part of the code is given in Table 1.8-1. The dash is supposed to be three times the length of the dot. Although the Morse code may appear

TABLE 1.8-1 *Morse Code*

A	.-
B	-... .
C	-.-.
D	-.. .
E	.
F	...-.
G	---.
H	
I	..
J	.---
K	-.-
L	.-.. .
M	---
N	-. .
O	---
P	.-.-.
Q	---.-
R	.-.
S	
T	-
U	...-
V	...--
W	-.--
X	-.-.-
Y	---.-
Z	---..

to be encoded in a binary code, it is, in fact, a ternary (radix 3, $r = 3$) code, having the symbols dot, dash, and space. The space between dots and dashes in a single letter is 1 unit of time, between letters it is 3 time units, and between words it is 6 time units.

We now make a brief digression to introduce a notation. We will continually need the binomial coefficients

$$C(n, k) = \frac{n!}{k! (n - k)!}$$

They count the number of ways a set of k items can be selected from a set of n items. We have adopted the old-fashioned notation $C(n, k)$ because it is easily produced on a typewriter, it is easily set in print, and it can be gracefully handled by most computers. The currently popular notation

$$\binom{n}{k}$$

is difficult for most equipment to handle and is awkward when it appears in running text.

The Morse code is clearly a *variable-length code* that takes advantage of the high frequency of occurrence of some letters, such as “E,” by making them short, and the very infrequent letters, such as “J,” relatively longer. However, the problems that arise when trying to recognize the words of a variable-length code are great enough in this case to cause almost complete replacement of the Morse code by the van Duuren code, which uses three out of seven positions filled with 1’s and the other four with 0’s. There are $C(7, 3) = 35$ possible words in this code, and as with the 8-bit ASCII code, the van Duuren code permits the detection of many types of errors since the receiver knows exactly the number of 1’s that should be in the received message unit of seven time slots.

Another widely used simple code is the *2-out-of-5 code*. As the name implies, two out of five positions are filled with 1’s. In this code there are, very conveniently, exactly $C(5, 2) = 10$ possible symbols. One of the many ways of associating the code symbols with the numerical values of the decimal digits is the *01247 code*, which is a *weighted code* where we attach *weights* 0, 1, 2, 4, and 7 to the successive columns of the code. The corresponding decimal digit is the sum of the weights where the 1’s occur, with the sole exception that the combination 4, 7 is taken as 0. The code is given in Table 1.8-2. Again, any single error in a message can be recognized because it will have an odd number of 1’s in it.

Exercises

- 1.8-1** Table 1.8-2 is one assignment of number values to the 10 possible symbols of the 2-out-of-5 code; how many possible 2-out-of-5 codes can there be? Ans. 10!

TABLE 1.8-2 *The 2-Out-Of-5 Code*

0	1	2	4	7	Decimal	Corresponds to:
1	1	0	0	0	1	0 + 1
1	0	1	0	0	2	0 + 2
0	1	1	0	0	3	1 + 2
1	0	0	1	0	4	0 + 4
0	1	0	1	0	5	1 + 4
0	0	1	1	0	6	2 + 4
1	0	0	0	1	7	0 + 7
0	1	0	0	1	8	1 + 7
0	0	1	0	1	9	2 + 7
0	0	0	1	1	0	4 + 7

1.8-2 Write 125 in the 2-out-of-5 code.

Ans. 11000 10100 01010

1.8-3 Of all possible 7-bit odd-parity symbols, how many are not used in the van Duuren code?

Ans. 29

1.9 Radix r Codes

As noted earlier, most systems of representing information in computers (and other machines) use two states, although the Morse code is an example of a signaling system with three symbols in the underlying alphabet. The reason for the dominance of two-symbol signaling systems is that two-state devices tend to be more reliable than multistate devices. On the other hand, human beings clearly work better with multistate systems—witness the letters of the alphabet, together with the various punctuation symbols and the decimal digits. Thus it is necessary at times to consider codes with r symbols in their alphabets. For the Morse code, $r = 3$.

The English language alphabet uses 26 letters, both upper- and lowercase, plus miscellaneous punctuation marks. Occasionally, you see a system in which there are exactly the 26 letters, the 10 decimal digits, and a space, a total of 37 symbols. We examine this important code in Section 2.7.

In more abstract notation we assume a source alphabet S consisting of q symbols $s_1, s_2, \dots, s_q$. These are, in turn, represented in some other symbols, say the binary code. We may, therefore, either think of the ASCII code as having $r = q = 2^7 = 128$ symbols $s_1, s_2, \dots, s_{128}$, or we may think of each symbol as being a block of 8 binary digits with 1 extra bit appended to the 7 necessary bits.

1.10 Escape Characters

When sending information it is usually necessary to control, from the source end, the remote equipment that is being used. For example, we have to tell the equipment what to do with the information that is being sent. Therefore, it is necessary to have a number of *reserved symbols* to which the remote equipment responds. “End Of Message” is such a symbol; another might be “Carriage Return”; another a “Shift Type Font” from lower- to uppercase or back again; and still another might be “Repeat Last Message.” In the ASCII code the entire first column of the table is devoted to special characters.

How can this control be done if we do not wish to restrict the information that can be sent through the system? If, for example, we

were using binary digits as numbers in our message, how do we avoid those combinations of bits that are control symbols for the remote equipment? If by chance our stream of binary numbers happened to contain such a reserved symbol, then the equipment would respond to the reserved symbol when we did not want it to do so. Evidently, some kind of restriction on what can be sent is necessary, but the question is: How can this be done with the least pain to the user of the system?

One way this could be done is to add an extra binary digit to each block of digits, using, say, a 0 if it is message and a 1 if it is a control instruction. As can be seen, this will waste a good deal of channel capacity *if* the control words are rarely used. Furthermore, it does *not* solve the problem of relaying control words *through* one piece of equipment to another piece of equipment.

In one form of FORTRAN, the apostrophe (') is used to delimit a string of characters. For example, the string of symbols ABC is written

'ABC'

In the FORTRAN processing, the apostrophes are stripped off and the string ABC emerges.

If the string happens to contain an apostrophe, such as the string FINNEGAN'S WAKE, you need to write

'FINNEGAN''S WAKE'

The processing strips off the first of the pair but lets through the one *immediately* following.

To see how to write the underlying program of the FORTRAN compiler, we use a (slightly extended) *state diagram*, which is widely used in finite automata theory. Each state consists of READ NEXT SYMBOL followed by WHAT TO DO. We have three states to consider: (1) the normal state of processing FORTRAN text, (2) the usual inside-the-string state, and (3) the inside-the-string state that looks for an immediately succeeding apostrophe. The state diagram is shown in Figure 1.10-1. If the system has more than one reserved symbol, slight variations on this diagram will handle them.

In the field of logic this problem is known as the *metalanguage problem*: How does one talk about the language itself, especially when one has to use essentially the same language to do the talking? In practice we use voice accents (quote marks) to indicate that we are talking about the language itself, and we leave them out when we are merely talking. There are other formal devices for solving this common problem that arises in computer science and in other fields of information




Figure 1.10-1 Reserved symbol automata

handling and processing, and these may also be adapted to this common, vexatious problem.

In controlling a digital computer the failure to distinguish between the various levels of metalanguages is a source of constant confusion. A FORTRAN program and the instruction to COMPILE it are in different languages. A program requires two end marks, one for the metalanguage level of COMPILE to indicate the end of the source program and to stop the compilation process, and one for the FORTRAN program to be used at RUN TIME to end the actual computation.

Exercise

1.10-1 Describe in detail how you would encode 'Help!' to get through the FORTRAN-type system described in the text.

Ans. "Help!"

1.11 Outline of the Course

We have introduced the main topics as well as a number of codes with varying properties, and now have a basis for discussing the material the course is to cover.

In Chapters 2, 3, and 11 we look at ways of channel-encoding information (messages, sources of symbols) so that any errors made in going through a channel, up to a given level, may be detected and/or corrected at the terminal end without recourse to the source. For a

detected but uncorrectable error, we might call for a repeat of the message, hoping to get it correct the next time. In general, the error-detecting and error-correcting capability will be accomplished by adding some digits to the message, thus making the message slightly longer. The main problem is to achieve the required protection against the inevitable errors (noise) in the channel without paying too high a price in adding extra digits. This is encoding for the channel.

In Chapters 4 and 5 we focus on source encoding. The compression of messages is important for efficiency. In transmission through space the shorter message will use the signaling equipment for a shorter time; for storage problems less storage will be needed for the compressed code. This is source encoding. In Chapters 4 and 5 we look at this side of the problem and examine ways to reduce the amount of information being sent. We do this by examining the structure of the messages being sent. When there is a great deal of structure in the information being sent, a good deal of message compression can be achieved, and hence greater efficiency results. Since there are so many properties that messages can have which allow message compression, we can look at only a few of the most common ones. We restrict our attention to general methods and must neglect the many special, trick methods, which are often fairly easy to invent.

In Chapter 6 we introduce the central concept of the *entropy* of a source of information and show how it is connected with the concept of the maximum amount of information that can be sent through a given channel. Thus we get a bit closer to the concept of exactly what information is. The first of Shannon's encoding theorems is for a noiseless channel, and is fairly easy to prove. The second theorem, presented in Chapter 10, discusses sending information through a noisy channel and is much more difficult to prove. Fortunately, the simple *binary symmetric channel* is the realistic case, and for this case the proof is easy to understand. The proof in the more general case is more difficult, and is only sketched.

Shannon's two theorems set bounds on what encoding can accomplish. Unfortunately, the second of his theorems is somewhat nonconstructive and does not tell us how in practice to achieve the bounds indicated. The result is not useless since it indicates where we can expect to achieve large improvements in a signaling system and where we can expect at best only small improvements.

The definition of a channel in information theory is often not a practical one and we need other ways of assessing this central concept of channel capacity for transmitting information. Therefore, Appendix A looks briefly at how, in practice, we measure the channel capacity of a signaling system, especially for continuous signals.

We remind the reader that while we often use the language of “signaling from here to there,” all of it is applicable to “signaling from now to then” through some storage medium. In this case it is often the first of Shannon’s theorems (the noise-free one) that is important.

Traditionally, the development of both coding and information theory is done in as elegant and general a way as possible—usually highly abstract, with many fancy mathematical symbols, and devoid of practical aspects. This approach is not necessary, and we regularly pause to show how what we are talking about is reflected in common sense and in actual practice, and how it suggests going about the design of future systems. We also try to make the proofs of the results as intuitively obvious as possible rather than merely mathematically elegant. Along the way we introduce many small, practical details that are important when designing a whole system.

But let us be clear about the approach we are adopting. When there are so many different highly evolved and highly involved signaling systems now in use, and so many new ones over the immediate horizon, it is impractical to take them up one at a time. It is necessary to take an overview approach and concentrate on the fundamentals that seem most likely to be relevant for understanding both past and future signaling systems. Indeed, this is the only reasonable approach to any rapidly changing field of knowledge. The approach through special cases simply leaves the student surprised by tomorrow’s system. Thus, of necessity, the treatment is, at times, abstract and avoids many of the messy details of current systems.

Error-Detecting Codes

2.1 Why Error-Detecting Codes?

We first examine channel encoding. The problem is to design codes that will compensate for the inevitable noise that is in all realistic channels. More and more often we require very reliable transmission through the channel, whether it be through space when signaling from here to there (transmission), or through time when signaling from now to then (storage). Experience shows that it is not easy to build equipment that is highly reliable. By “highly reliable” consider how much computing a modern computer does in 1 hour. At one operation per microsecond it does 3.6 billion (10^9) operations in 1 hour (and each operation involves many individual components). By comparison there are less than 3.16 billion seconds in 100 years (more than your probable lifetime). Similarly, reliable transmission systems require very high reliability of their individual components. Reliability in the transmission of words of human languages is one thing; for transmission of computer programs it is something else—hence the importance of detecting errors. Furthermore, as already noted, error detection is a great aid in high-quality maintenance. Without error detection, a large digital system soon becomes unmaintainable.

If repetition is possible, then it is frequently sufficient merely to detect the presence of an error. When an error is detected we simply repeat the message, the operation, or whatever was being done, and with reasonable luck it will be right the second (or even possibly the third) time (see Section 2.5).

It is obviously not possible to detect an error if every possible symbol, or set of symbols, that can be received is a legitimate message. *It is possible to catch errors only if there are some restrictions on what is a proper message.* The problem is to keep these restrictions on the possible messages down to ones that are simple. In practice, “simple” has in the past tended to mean “easily computable,” but it could in the future mean “easily looked up.” In this chapter we investigate the problem of designing codes such that *at the receiving end*, any single, isolated error can be detected. In Chapter 3 we consider correcting at the receiving end the errors that occur in the message.

2.2 Simple Parity Checks

The simplest way of encoding a binary message to make it error detecting is to count the number of 1's in the message, and then append a final binary digit chosen so that the entire message has an even number of 1's in it. The entire message is therefore of even parity. Thus to $(n - 1)$ message positions we append an n th parity-check position. At the receiving end the count of the number of 1's is made, and an odd number of 1's in the entire n positions indicates that at least one error has occurred.

Evidently, in this code a double error cannot be detected. Nor can any even number of errors be detected. But any odd number of errors can be detected. If (1) the probability of an error in any one binary position is assumed to be a definite number p , and (2) errors in different positions are assumed to be independent, then for n much less than $1/p$, the probable number of single errors is approximately np . The probability of a double error is approximately $n(n - 1)p^2/2$, which is approximately one-half the square of single error. From this it follows that the optimal length of message to be checked depends on both the reliability desired (the chance of a double error going undetected) and the probability of a single error in any one position, p . For more details, see Section 2.4.

Counting the number of 1's and selecting an even number is equivalent to working in a *modulo 2 arithmetic*. “Modulo 2” means that every number is divided by 2 (the modulus) and only the remainder is kept. In this arithmetic (which counts 0, 1, 0, 1, . . .) we count modulo 2 the number of 1's in the $(n - 1)$ message positions and then put this sum into the n th position. Thus there are an even number of 1's in the n positions sent. This *even-parity check* is generally used in the theory. In practice, it is occasionally convenient to adopt an *odd-parity check* so that the message of all 0's is not a legitimate message. The changes

necessary in the theory are sufficiently easy to make that we will continue to use even-parity checks throughout the theory.

To actually find the parity of a string of 0's and 1's we can use a two-state finite automaton (Figure 2.2-1). We start in the state 0, and each 1 in the message causes a change in the state. The final state at the end of the message gives the parity count.




Figure 2.2-1 Parity count circuit

Many computers have an instruction that directly counts the number of 1's in the accumulator register. The right-hand digit of this sum is the required count modulo 2. If this or an equivalent instruction is not available, then *logically adding* (exclusive OR, XOR; see Section 2.8) one half of the message to the other half preserves the parity count in the sum. Repeated use of this observation halves at each step the length of the message and will produce the required sum modulo 2 of all the 1's in the accumulator. Thus about “the first integer greater than or equal to $\log_2 n$ ” logical additions will be needed.

2.3 Error-Detecting Codes

It is common practice to break up a long message in the binary alphabet into runs (blocks) of $(n - 1)$ digits each, and to append one binary digit to each run, making the block that is sent n digits long. The final block may need to be padded out with 0's. This produces the *redundancy* of

$$\frac{n}{n - 1} = 1 + \frac{1}{n - 1}$$

where the redundancy is defined as the number of binary digits actually used divided by the minimum necessary. The *excess redundancy* is $1/(n - 1)$.

Clearly, for low redundancy we want to use long messages. But for high reliability short messages are better. Thus the choice of the

length n for the blocks of message to send for a given probability of error p is an engineering compromise between two opposing forces.

The 2-out-of-5 code mentioned in Section 1.8 is an example of an error-detecting code with $n = 5$ that uses an even-parity check. The van Duuren 3-out-of-7 code is another such code, but uses odd parity. Neither uses all the symbols it could. A variant is the word count occasionally used in sending telegrams.

Many other examples of this simple kind of encoding for error detection occur in practice. For example, when making a memory transfer to tape, drum, or disk, the *logical sum* of all the words being transferred can be appended as a check word. See Ref. [W] for hardware realizations.

2.4 Independent Errors: White Noise

We have already mentioned (Section 2.2) the model usually assumed for errors in a message: (1) an equal probability p of an error in each position, and (2) an independence of errors in different positions. This is called “white noise” in (a poor) analogy with white light, which is supposed to contain uniformly all the frequencies that are detected by the human eye. The theory for this case is very easy to understand. But in practice there are often reasons for errors to be more common in some positions in the message than in others, and it is often true that errors tend to occur in bursts and not be independent (a common power supply, for example, tends to produce a correlation among errors; so does a nearby lightning strike; see Section 2.6).

It is only fair to note that due to the rapid pace of technological change, it is not possible at the design stage to have much of an idea about the noise patterns of errors that will actually occur, so that white noise is often a reasonable assumption.

The probability of an error in any one position is p ; hence the probability of no error is $1 - p$. From the assumed independence of the errors, that is, for white noise, the probability of no error in the n positions is

$$(1 - p)^n$$

The probability of a single error in the n positions is

$$np(1 - p)^{n-1}$$

The probability of k errors is given by the k th term in the binomial expansion

$$1 = [(1 - p) + p]^n = (1 - p)^n + np(1 - p)^{n-1} + \frac{n(n-1)}{2} p^2(1 - p)^{n-2} + \dots + p^n$$

For example, the probability of exactly two errors is

$$\frac{n(n-1)}{2} p^2(1 - p)^{n-2}$$

We can get a formula for the probability of an even number of errors (0, 2, 4, . . .) by adding the following two binomial expansions and dividing by 2:

$$1 = [(1 - p) + p]^n = \sum_{k=0}^n C(n, k)p^k(1 - p)^{n-k}$$

$$[(1 - p) - p]^n = \sum_{k=0}^n (-1)^k C(n, k)p^k(1 - p)^{n-k}$$

(The $[]$ in the next equation means “greatest integer in.”)

$$\frac{1 + (1 - 2p)^n}{2} = \sum_{m=0}^{[n/2]} C(n, 2m)p^{2m}(1 - p)^{n-2m} \quad (2.4-1)$$

The probability of an odd number of errors (which will always be detected) is 1 minus this number.

The probability of no errors is the first term ($m = 0$) of the series (2.4-1). Therefore, to get the probability of an undetected error, we drop the first term of (2.4-1), to get

$$\sum_{m=1}^{[n/2]} C(n, 2m)p^{2m}(1 - p)^{n-2m} \quad (2.4-2)$$

Usually, only the first few terms are of any importance in evaluating this formula.

Exercises

2.4-1 If $p = 0.001$ and $n = 100$, what is the probability of no error?

Ans. $\exp(-\frac{1}{10})$

2.4-2 If $p = 0.001$ and $n = 100$, what is the probability of an undetected error?

2.4-3 If $p = 0.01$ and you want the probability of an undetected error to be 0.005, what is the maximum length n that you can use?

Ans. $n = 10$ but almost $n = 11$

2.4-4 For small p and arbitrary n , do Exercise 2.4-1. Ans. $\exp(-np)$

2.5 Retransmission of Message

When an error is detected, it is often possible to ask for a retransmission, a recomputation, or repeat of the process being done. For example, in the Bell Relay Computers the code used for representing the decimal digits was a 2-out-of-5 code, and whenever an error was detected, a second, and even a third trial was made. On the Model 6 Relay Computer, if after the third trial the message was still not acceptable, then the whole problem was dropped and the next one taken up, in the hope that the defective part would not be used in the new problem.

When reading magnetic tapes it is the custom to use at least an error-detecting code, and to call for repeated readings of the tape if the parity checks are violated. How many repetitions to use depends on your model for the error. If you think that the error is due to a slight loss of magnetization, then with luck a subsequent trial will read the tape correctly; but if you think it is due to a more permanent failure, then repetitions will probably occur until another, independent error occurs, so that then the parity check is met. You will therefore get a message with two errors in it, rather than the right message! Thus the strategy of “detect an error, and if one is found, call for one or more repetitions” is sound *only if* the type of error you expect is transient.

Parity checks have long been used in computers, in both hardware and software (Ref. [W]). For example, in the early days of unreliable drum storage, every WRITE on the drum was followed by the logical sum of all the registers that were used for the message being stored on the drum. This check sum was then stored in a final register of the block on the drum. The drum was next read back and checked to see if the message was written properly. Only then was the information being stored released from the machine’s storage registers. When later the drum was read, the parity check (the logical sum of *all* the registers being stored) was again computed to see if the sum was identically zero. If not, retrials were made to read it again. This method depends on the fact that $x + x = 0$ for logical addition.

2.6 Simple Burst Error-Detecting Codes

Noise (errors) often occur in bursts rather than in isolated positions in the received message. Lightning strikes, power-supply fluctuations, and loose flakes on a magnetic surface are all typical causes of a burst of noise.

Suppose that from measurements in the field we agree on the maximum length L of any burst we are to detect. For ease of discussion (since the changes necessary to handle other cases are easy to invent) assume that the burst length L is the accumulator word length of the computer. We have only to select the appropriate error-detecting code (or error-correcting code of Chapter 3) and instead of computing parity checks over the bit positions, we compute parity word checks over the corresponding word positions. In effect, we work in words, not bits, and have L independent (interleaved) codes, one over each bit position in the word (see Section 3.8). If a burst covers the end of one word and the beginning of another, still no two errors will be in the same code since we *assumed* that any burst length k was ($0 \leq k \leq L$).

Thus we can send messages through noise that is “bursty” provided that we recognize the noise pattern and design for it. For reliable hardware, see Ref. [W].

Example. If the message is

Fall 1980

this can be encoded in ASCII in a burst code as follows (no parity check is used here):

$$\begin{array}{rcl}
 \text{F} & = & 106 = 01 \quad 000 \quad 110 \\
 \text{a} & = & 141 = 01 \quad 100 \quad 001 \\
 \text{1} & = & 154 = 01 \quad 101 \quad 100 \\
 \text{1} & = & 154 = 01 \quad 101 \quad 100 \\
 \text{sp} & = & 040 = 00 \quad 100 \quad 000 \\
 \text{1} & = & 061 = 00 \quad 110 \quad 001 \\
 \text{9} & = & 071 = 00 \quad 111 \quad 001 \\
 \text{8} & = & 070 = 00 \quad 111 \quad 000 \\
 \text{0} & = & 060 = 00 \quad 110 \quad 000 \\
 \hline
 \text{Check sum} & = & 00 \quad 000 \quad 111 = \text{BEL}
 \end{array}$$

The encoded message is therefore

Fall 1980BEL

where BEL is the single ASCII symbol 00 000 111.

Exercise

2.6-1 Given the message

101	01010
011	00110
000	11110
110	00110
111	10101

add the appropriate check to form a burst error-detecting code of word length 1 byte. Ans. 111 00001

2.7 Alphabet Plus Number Codes: Weighted Codes

The codes we have discussed so far have generally assumed a simple form of white noise. This is very suitable for many types of machines, although in serial transmission the loss of a symbol (or the insertion of an extra one) is a common error in some systems, and is not caught by such codes, hence causes a loss of synchronization.

When dealing with human beings, another type of noise is more appropriate. People have a tendency to interchange adjacent digits of numbers; for example, 67 becomes 76. A second common error is to double the wrong one of a triple of digits, two adjacent ones of which are the same; for example, 667 becomes 677, merely a change of one digit. These are the two most common human errors in arithmetic. In a combined alphabet/number system, the confusion of “oh” and “zero” is very common.

A rather frequent situation is to have an alphabet, plus space, plus the 10 decimal digits as the complete set of symbols to be used. This amounts to $26 + 1 + 10 = 37$ symbols in the sending alphabet. Fortunately, 37 is a prime number and we can use the following method for error checking. We *weight* the symbols with weights 1, 2, 3, . . . beginning with the check digit of the message. We reduce the sum modulo 37 (divide by 37 and take the remainder) so that a check symbol can be selected that will make the sum 0 modulo 37. Note that a “blank” at the end, as a check symbol, is *not* the same as nothing.

It is easy to compute that an interchange of adjacent digits will be detected, and that doubling the wrong digit (which is the changing of a single symbol) will also be detected. If the interchanged digits are the k th and $(k + 1)$ st, their original sum,

$$ks_k + (k + 1)s_{k+1}$$

becomes

$$(k + 1)s_k + ks_{k+1}$$

and the change

$$ks_k + (k + 1)s_{k+1} - (k + 1)s_k - ks_{k+1} = s_{k+1} - s_k$$

cannot be divided by 37 with a remainder 0 unless $s_k = s_{k+1}$. A similar argument applies to the interchange of any two digits, adjacent or not (assuming a message length less than 37). Many other interchanges can also be caught. The change of “oh” to “zero” is the change of a single symbol. Such a code is very useful when human beings are involved in the process. Notice that the length of the message being encoded is not fixed, and can exceed 37 symbols if necessary (although special precautions are then required). This type of encoding can be used on credit cards, the names of some kinds of items in inventory, and so on. The probability that a random encoding will get through the input check is $\frac{1}{37}$. Thus by using this simple parity check, forgery attempts will be caught about 97% of the time.

When the check symbol is the “space,” this can sometimes be awkward. If it is awkward, then frequently the whole label, say an inventory name, can be dropped and a new label chosen in its place (a loss of about $\frac{1}{37}$ of the possible inventory names results).

To find the weighted check sum easily, note that if you compute the running sum of a set of n numbers, and then sum these again, you will have the first number entered into the final sum n times, the next number $n - 1$ times, the next $n - 2$ times, and so on, down to the last number, which will get into the final total only once. Thus you have the required weighted sum of the numbers that correspond to the symbols of the alphabet being used.

If this “summing the sum” seems mysterious, consider the message $w x y z$:

Message	Sum	Sum of sum
w	w	w
x	$w + x$	$2w + x$
y	$w + x + y$	$3w + 2x + y$
z	$w + x + y + z$	$4w + 3x + 2y + z$

where we see the weighted sum emerging. This is sometimes known as “progressive digitig.”

Example. If $0 = 0, 1 = 1, 2 = 2, \dots, 9 = 9, A = 10, B = 11, \dots, Z = 35, b1 = 36$, then encode

$$A6 \ 7$$

We proceed as follows:

	Sum	Sum of Sum
$A = 10$	10	10
$6 = 6$	16	26
$b1 = 36$	52	78
$7 = 7$	59	137
$x = x$	$59 + x$	$196 + x$

$$\begin{array}{r} 5 \\ 37 \overline{)196 + x} \\ \underline{185} \\ 11 + x \end{array}$$

Since $11 + x$ must be divisible by 37, it follows that

$$x = 37 - 11 = 26 = Q$$

The encoded message is therefore

$$A6 \ 7Q$$

To check at the receiver that this is a legitimate encoded message, we proceed as follows:

$$\begin{array}{rcl} A & 10 \times 5 = & 50 \\ 6 & 6 \times 4 = & 24 \\ b1 & 36 \times 3 = & 108 \\ 7 & 7 \times 2 = & 14 \\ Q & 26 \times 1 = & \underline{26} \\ \text{Sum} & = 222 = 37 \times 6 \equiv & 0 \text{ modulo } 37 \end{array}$$

The method of encoding used in this section is an example of encoding for *human beings* rather than for the conventional *white noise*. We see how dependent the encoding method chosen is on the assumed noise. It is necessary to emphasize this point because the theory so often tacitly assumes the presence of white noise that the reader is apt to forget this basic assumption when involved with an elegant, complex code.

Exercises

2.7-1 If $0 = 0, 1 = 1, \dots, 9 = 9, A = 10, B = 11, \dots, Z = 35,$
 $b1 = 36$, encode

B23F mod 37

Ans. B23F9

2.7-2 In the same code as that of Example 2.7-1, is

K9K9

a correct message?

2.8 Review of Modular Arithmetic

Because we are often going to use parity checks, we need to get a firm grasp on the arithmetic of the corresponding manipulations. We have seen that mod 2 addition (“mod” is an abbreviation for “modulo”) is the arithmetic that the simple binary parity checks use and is the same as logical addition (exclusive OR; XOR). The rules for this form of addition are

$$0 + 0 = 0$$

$$0 + 1 = 1$$

$$1 + 0 = 1$$

$$1 + 1 = 0$$

There are no numbers other than 0 and 1 in the system. If we choose to work in normal arithmetic, then we merely divide the result by 2 and take the remainder. When we later come to the algebra of linear equations and polynomials, where the coefficients are from the mod 2 number system, we will have the same table for addition. For multiplication we have the rules

$$0 \times 0 = 0$$

$$0 \times 1 = 0$$

$$1 \times 0 = 0$$

$$1 \times 1 = 1$$

Thus multiplication is the logical AND of computing.

Occasionally, we will work modulo some other number than 2. For example, in the preceding section we used numbers modulo 37. Generally, the theory in any prime base p (such as 37) is very much like that of base 2, and we need not go into the details here. You have only to read the preceding paragraph and make simple changes to understand the corresponding arithmetic and algebra. For addition and subtraction mod p , we divide every number by p and take the positive remainder.

For multiplication mod m (not a prime) we must be more careful. Suppose that we have the numbers a and b congruent to a' and b' modulo the modulus m . This means that

$$a \equiv a' \pmod{m}$$

$$b \equiv b' \pmod{m}$$

or

$$a = a' + k_1m$$

$$b = b' + k_2m$$

for some integers k_1 and k_2 . For the product ab we have

$$ab = a'b' + a'k_2m + b'k_1m + k_1k_2m^2$$

$$ab \equiv a'b' \pmod{m}$$

Now consider the particular case

$$a = 15, \quad b = 12, \quad m = 10$$

We have

$$a' = 5, \quad b' = 2$$

and

$$ab \equiv a'b' \equiv 0 \pmod{10}$$

But neither a nor b is congruent to zero! *Only for a prime modulus do we have the important property that if a product is zero, then at least one factor is zero.* Hence the importance of a prime modulus. Now we see why 37 was so convenient a number in Section 2.7.

Modular arithmetic should be clearly understood, especially the need for a prime modulus, because in Chapter 11 we will face the problem of constructing a corresponding modular algebra. In the next section, we therefore give yet another example of this type of encoding.

2.9 ISBN Book Numbers

The International Standard Book Number (ISBN) now appears on most textbooks. It is (usually) a 10-digit code that publishers assign to their books. A typical book will have

0-1321-2571-4

although the hyphens may appear in different positions. (The hyphens are of no importance.) The 0 is for United States and some other English-speaking countries. The 13 is for Prentice-Hall, the publisher. The next six digits, 21-2571, are the book number assigned by the publisher, and the final digit is the weighted check sum, as in Section 2.6. Modulo 10 will not work, since 10 is a composite number. Thus they use the check sum modulo 11, and are forced to allow an X if the required check digit is 10. As before, this could be avoided by giving up $\frac{1}{11}$ of the possible book names.

Searching around, I found the ISBN

0-1315-2447-X

To check that this number is a proper ISBN, we proceed as follows:

code	sum	sum of sum
0		
1	1	1
3	4	5
1	5	10
5	10	20
2	12	32
4	16	48
4	20	68
7	27	95
$X = 10 \quad 37 \quad \underline{132} = (11) \times (12) = 0 \pmod{11}$		

It checks!

Here again we see a simple error-detecting code designed for human use rather than for computer use. Evidently, such codes have wide applicability.

Exercises

2.9-1 Check the ISBN 0-13165332-6.

2.9-2 Check the ISBN 0-1391-4101-4.

2.9-3 Consider the ISBN 07-028761-4. Does it make sense?

2.9-4 Check the ISBNs in your current textbooks.

$$\begin{array}{r}
 3 \quad 4 \\
 2 \\
 1 \quad 7 \\
 \quad 2 \\
 5 \\
 7 \quad 2 \\
 1 \quad 2 \quad 2 \\
 4 \quad 2 \quad 6 \quad 110 = 11.10
 \end{array}$$

$$\begin{array}{r}
 4 \quad 4 \quad 4 \\
 7 \quad 1 \quad 15
 \end{array}$$

$$\begin{array}{r}
 6 \\
 2 \quad 20 \quad 65 \\
 1 \quad 21 \quad 86 \\
 8 \\
 7 \quad 151 \\
 0 \quad 1 \quad 1 = 11.17 \quad ('ADAPTK !')
 \end{array}$$

Error-Correcting Codes

3.1 Need for Error Correction

For channel encoding, simple error detection with repetition is often not enough. The immediate question is: If a computer can find out that there is an error, why can it not find out where it is? The answer is that by proper encoding this can, indeed, be done. For example, we can write out, or do, each thing three times, and then take a vote. But as we shall see, much more efficient methods are available. As an example of the use of an error-correcting code, when information is sent from space probes to Earth, the signaling time is so long that by the time an error can be discovered, the source may long since have been erased. The codes currently used are very sophisticated and can correct many simultaneous errors in a block.

Error correction is also very useful in the typical storage (memory) system. Many organizations, for example, store much of their information on magnetic tapes. Over the years the quality of the recording gradually deteriorates, and the time comes when some of it cannot be read reliably. When this happens, the original source of the information is often no longer available and therefore the information stored on the tape is forever lost. If some error-correction ability is incorporated into the encoding of the information, the corresponding types of errors can be corrected.

Error correction has long been used and can be incorporated into the hardware. For example, the old NORC computer had a simple type of error correction for isolated errors that occurred on its cathode-

ray-tube storage device. Often, it is only the storage device that uses the error correction. The STRETCH computer that IBM built “to *stretch* the state of the art” had single-error correction and double-error detection (Section 3.7) throughout much of the computer. On the STRETCH acceptance test it has been claimed that an error in a circuit developed in the first few minutes and that the error-correcting circuits fixed up the errors for the rest of the hour. Currently, most high-density magnetic recording uses some form of error correction, as does hi-fi digital recording of music. As the components get smaller, solid-state memories generally go to error-correcting codes. See Ref. [W] for hardware checking.

The error correction capability can be in the software as well as in hardware. We mentioned earlier that error-detection programs for drum storage were in the software. Error correction could also be used there. In many other places, correcting codes are built into the software. The software approach has the advantage that the error correction can more easily be put on the important parts of the information and omitted from the less important parts. Furthermore, only as experience reveals the weaknesses of the computer system is it known where error protection is most needed as well as how much is needed. Only after simulation of error protection in software has proved that it handles the problem should it be added to the system in the form of hardware. However, error correction in software will clearly use a lot of computer time.

3.2 Rectangular Codes

The first, and simplest, error-correcting codes are the *triplication codes*, where every message is repeated three times and at the receiving end a majority vote is taken. It is obvious that this system will correct isolated single errors. However, it is also obvious that this encoding system is very inefficient. Three computers in parallel with the associated “compare and interrupt if necessary circuits” is expensive. Therefore, we look for better methods of encoding information so that a single error can be corrected.

The next simplest error-correcting codes are the *rectangular codes*, where the information is arranged, in principle, although not necessarily in fact, in the form of an $m - 1$ by $n - 1$ rectangle (Figure 3.2-1). A parity-check bit is then added to each row of $m - 1$ bits to make a total of m bits. Similarly, a check bit is added for each column. Whether or not the final corner bit is used has little importance; for even-parity checking, the parity sum of its row and its column are the same. Thus




Figure 3-2.1 Rectangular codes

the original code of $(m - 1)(n - 1)$ bits of the rectangle becomes an array of mn bits. The redundancy is therefore $1 + 1/(m - 1) + 1/(n - 1) + 1/(m - 1)(n - 1)$. For a given size mn , the redundancy will be smaller the more the rectangle approaches a square. For square codes of side n , we have $(n - 1)^2$ bits of information and $2n - 1$ bits of checking along the sides.

A rectangular code, for example, was the code that the NORC computer used for its cathode-ray storage tubes. Rectangular codes often are used on tapes (Refs. [B2], [W]). The longitudinal parity checks are on the individual lines, and an extra line is put at the end of the block to form the vertical parity check. Unfortunately, the machine designer rarely lets the user get at the parity-check information so that by using suitable software programs the isolated error could be corrected.

Exercises

3.2-1 Discuss various possible rectangular codes for 24 message bits.

3.2-2 Discuss the use of the corner check bit. Prove that for even-parity checks it checks *both* row and column. Discuss the use of odd-parity checks in this case.

3.2-3 Compare an n^2 code with an $(n - 1)(n + 1)$ code.

3.3 Triangular, Cubic, and n -Dimensional Codes

Thinking about rectangular codes soon suggests a *triangular code*, where each element on the diagonal (and there are only n of them if the entire triangle has n bits on a side) is set by a parity check covering *both* its own row and its own column (Figure 3.3-1). Thus for a given size we




Figure 3.3.1 Triangular codes

lower the redundancy we use. The triangular array of bits of information has $n(n - 1)/2$ bits with n bits added for the checks. Thus the redundancy is $1 + 2/(n - 1)$.

A triangular code of side n has

$$\frac{n(n - 1)}{2} \text{ message and } n \text{ check bits}$$

When we compare a triangular code with a square code for a given number of parity check-bits we find that the triangular code has more message bits. We have included the cubic code in Table 3.3-1.

TABLE 3.3-1 Comparison of Three Codes

n	Square		Triangular		Cubic	
	Message $(n - 1)^2$	Check $2n - 1$	Message $n(n - 1)/2$	Check n	Message $n^3 - 3n + 2$	Check $3n - 2$
2	1	3	1	2	4	4
3	4	5	3	3	20	7
4	9	7	6	4	54	10
5	16	9	10	5	112	13
6	25	11	15	6	200	16
7	36	13	21	7	324	19
8	49	15	28	8	490	22
9	64	17	36	9	704	25
10	81	19	45	10	972	28

Once we have found a better code than the simple rectangular ones, we are led to ask the question: What is the best that we can do with respect to keeping down the redundancy? For a cube we can check each plane in three dimensions (Figure 3.3-2) by a single bit on the




Figure 3-3.2 Cubic code, showing only one plane and its check position

corresponding edge. The three intersecting edges of check bits use $3n - 2$ positions out of the n^3 total positions sent. The excess redundancy is roughly $3/n^2$ extra bits.

If the use of three dimensions is better than two dimensions, then would not going to still higher dimensions be better? Of course, we do not mean to arrange the bits in three-dimensional or higher space; we only imagine them so arranged to calculate the proper parity checks. A little thought about checking over the three-dimensional planes in four-dimensional space will lead to (approximately) an excess redundancy of $4/n^3$. We soon come to the conclusion that the highest possible dimensional space might be the best, and we want a total of 2^n bits of which $(n + 1)$ are checks (one check in each direction plus the common corner bit). There would be, therefore, $(n + 1)$ parity checks. If they were suitably arranged in some order, then they would give a number called the *syndrome*, which we would get by writing a 1 for each failure of a parity check, or a 0 when it is satisfied. These $(n + 1)$ bits form an $(n + 1)$ -bit number, and this number (the syndrome) can specify any of 2^{n+1} things—more than the 2^n locations of the corresponding single error *plus* the fact that no error occurred. Notice that the check bits, as in the simple error-detecting codes, are *equally protected* with the message bits. All the bits sent enter equally—none is treated specially.

However, we see that there are almost twice as many states in the syndrome as are needed to indicate when the entire message is correct plus the position of the error if one occurs. Thus there is almost a

factor of 2 in the loss of efficiency. Nevertheless, this suggests a new approach to the problem of designing an error-correcting code, which we take up in the next section.

Exercises

3.3-1 Make a triangular code of 10 message positions.

3.3-2 Discuss in detail the four-dimensional $2 \times 2 \times 2 \times 2$ code.

3.3-3 Compare the three codes of Table 3.3-1 that use approximately 10 check bits.

3.4 Hamming Error-Correcting Codes

In this section we adopt an algebraic approach to the problem that was started in the preceding section—finding the *best* encoding scheme for single-error correction *for white noise*. Suppose that we have m independent parity checks. By “independent” we mean, of course, that no sum of any combination of the checks is any other check—remember that they are to be added in our modulo 2 number system! Thus the three parity checks over positions

Check 1:	1, 2, 5, 7
2:	5, 7, 8, 9
3:	1, 2, 8, 9

are dependent checks since the sum of any two rows is the third. The third parity check provides no new information over that of the first two, and is simply wasted effort.

The *syndrome* which results from writing a 0 for each of the m parity checks that is correct and a 1 for each failure can be viewed as an m -bit number and can represent at most 2^m things. We need to represent the state of all the message positions being correct, plus the location of any single error in the n bits of the message. Thus we must have the inequality

$$2^m \geq n + 1 \quad (3.4-1)$$

Can such a code be found? We will show how to build ones that exactly meet the equality condition. They represent a particular so-

lution to the problem of designing a suitable code and are known as *Hamming codes*.

The simple, underlying idea of the Hamming code is that the syndrome shall give the actual position (location) of the error, with all 0's in the syndrome meaning no error. What positions should the parity checks use? Let us look at the binary representations of the position numbers (Table 3.4-1). Evidently, if the syndrome is to indicate the position of an error when it occurs, every position that has a 1 in the last position of its binary representation must be in the first parity check (see the extreme right-hand column of the table). *Think carefully why this must be true.* Similarly, the second parity check must be tripped by those positions that have a 1 in the second-lowest position in its binary representation; and so on.

TABLE 3.4-1 *Check Positions*

Position	Binary Representation
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111
8	1000
9	1001
10	1010
.	
.	
.	

Thus we see that the first parity check covers positions 1, 3, 5, 7, 9, 11, 13, 15, . . . ; the second, 2, 3, 6, 7, 10, 11, 14, 15, . . . ; the third, 4, 5, 6, 7, 12, 13, 14, 15, . . . ; the next, 8, 9, 10, 11, 12, 13, 14, 15, 24, 25, . . . ; and so on.

To illustrate what we have just said (see Table 3.4-2), let us design a simple error-correcting code for four binary digits. We must have $2^m \geq n + 1$ and we easily see that $2^3 \geq (4 + 3) + 1 = 8$. Therefore, we will need to have $m = 3$ parity checks, and this gives $7 = n$ = total of message plus parity-check positions. The positions to use to set the parity checks will for convenience be picked as 1, 2, and 4. Thus the information message positions are 3, 5, 6, and 7.

TABLE 3.4-2 Encoding a 4-Bit Message and Locating Error

<i>Encode</i>								
1	2	3	4	5	6	7	Position	
–	–	1	–	0	1	1	Message	
0	1	1	0	0	1	1	Encode	
×							Error	
0	1	0	0	0	1	1	Receive	

<i>Locate error</i>							
Check 1		1	3	5	7		
		0	0	0	1		Fails → 1
Check 2		2	3	6	7		
		1	0	1	1		Fails → 1
Check 3		4	5	6	7		
		0	0	1	1		Correct → 0
Syndrome =	0	1	1	=			3 → position of error

<i>Correct (add 1 into the error position)</i>							
		1					Correct error
0	1	1	0	0	1	1	Corrected message

To encode the information message, I write the message in the positions 3, 5, 6, 7 and compute the parity checks in positions, 1, 2, 4. Let this message be, say, __1_011; the spaces are where the parity checks are to go. The first parity check, which goes in position 1, is computed over positions 1, 3, 5, 7, and looking at the message I see that position 1 gets a 0. I now have 0_1_011. The second parity check, over positions 2, 3, 6, 7, sets position 2 as a 1. I then have 011_011 as the partially encoded message. The third parity check is over positions 4, 5, 6, 7, and looking at what I have, I see that position 4 must get a 0. Thus the final encoded message is 0110011.

To see how the code corrects an isolated error, suppose that when the message 0110011 is sent, the channel removes the 1 in the third position from the left. The corrupted message is then 0100011. For the receiving end, you apply the parity checks in order. The first, over positions 1, 3, 5, 7, evidently fails, so the lowest-order digit of the syndrome is a 1. The second parity check, over 2, 3, 6, 7, fails, so the second digit of the syndrome is 1. The third parity check, over 4, 5, 6, 7, succeeds, so the highest-order digit is a 0. Looking at the syndrome as a binary number, you have the decimal number 3. Thus you change (logically add 1) the symbol in position 3 from the received 0 to a 1. The resulting sequence of 0's and 1's is now correct, and when you strip

off the checking bits 1, 2, and 4, you have the *original information message 1011 in positions 3, 5, 6, and 7.*

Note that the check positions are equally corrected with the message positions. The code is *uniform in its protection; once encoded there is no difference between the message and the check digits.* The cute part of the Hamming code is the numerical ease of both the encoding and the code correction based on the syndrome at the received end. Also, note that the syndrome indicates the position of the error *regardless* of the message being sent; logically adding a 1 to the bit in the position given by the syndrome corrects the received message, where, of course, the syndrome “all 0’s” means that the entire message has no errors.

The redundancy in the example of four message bits with three check bits seems high. But if we took 10 check bits, we have, from equation (3.4-1), $2^{10} \leq 10 + \text{message positions} + 1$, or $1024 - 11 = 1013 \geq \text{message positions}$. This shows that the excess redundancy rises like $\log_2$ of the number of message positions.

Exercises

3.4-1 Discuss the Hamming code using four checks.

3.4-2 Discuss the Hamming code for two checks. Ans. Triplicate code

3.4-3 Show that 1111111 is a correct message.

3.4-4 Correct 1001111 and decode.

3.4-5 Locate the error in 011100010111110, where we have used four checks as in Exercise 3.4-1. What is the correct message?

3.4-6 What is the probability that a random sequence of $2^m - 1$ 0’s and 1’s will be a code word? Ans. $1/2^m$

3.5 Equivalent Codes

The example above is one way of encoding a message. There are many other *equivalent codes*. It should be obvious that any interchange of the positions of the code will leave us with a code that is only trivially different. Similarly, if we were to *complement* (change 0’s to 1’s and 1’s to 0’s) all the digits appearing in a certain position, then again we would have a trivially different code.

We can, if we wish, use this observation about the interchange of positions to move all the check bits to the end of the message (see Figure 3.5-1). For the arrangement we have the permutation

1 $\longrightarrow$ 5
 2 $\longrightarrow$ 6
 3 $\longrightarrow$ 1
 4 $\longrightarrow$ 7
 5 $\longrightarrow$ 2
 6 $\longrightarrow$ 3
 7 $\longrightarrow$ 4




Figure 3-5.1 Hamming code positions

Of course, the syndrome that emerges must also be changed by the same table. This permutation makes the “masking” of the message to get the parity checks, and to make the correction, less simple in appearance, but it is exactly the same amount of work by the computer, since for each check bit we still mask the received message to get the check bit by the “parity sum.” Where we place the check bit is also changed. Although the meaning of the syndrome has its position changed, the syndrome is still essentially the same and refers to a unique place. The position to logically add the correcting bit is therefore changed, but that can be found either by a formula or by a table lookup, which enters the table with the syndrome, and the entry in the table is the word with a 1 in the position to be corrected. Of course, for a long code a message might be two or more words long. The table size is proportional to n . The original Hamming code used computing power rather than table lookup, and, in general, for very long codes this is necessary.

One advantage of putting the check bits at the trailing end rather than distributed throughout the message is that sorting on the encoded message will give the same result as on the original messages.

As an alternative to decoding a message, you could, in these days of cheap, large memories, simply use the message as the address and look up the encoded message. Similarly, the encoding process could be a table lookup using addresses as long as the message of the encoded

block. Such methods, being rapid, tend to emphasize short blocks with higher redundancy than the more elaborate process of encoding (as given above), which works on arbitrarily long messages (although, of course, the longer the message, the more parity checks must be computed—in parallel if you have the available equipment).

Exercise

3.5-1 In the $n = 15$ code, discuss in detail the encoding and decoding if the four parity checks are moved to positions 12, 13, 14, and 15.

3.6 Geometric Approach

We have just given an algebraic approach to error-correcting codes. A different, equivalent approach to the topic is from n -dimensional geometry. In this model we consider the string of n 0's and 1's as a point in n -dimensional space. Each digit gives the value of the corresponding coordinate in the n -dimensional space (where we are assuming that the encoded message is exactly n bits long). Thus we have a cube in n -dimensional space; each vertex is a string of n 0's and 1's. The space consists *only* of the 2^n vertices; there is nothing else in the space of all possible messages except the 2^n vertices. This is sometimes called a “vector space.”

Each vertex is a possible received message, but only selected vertices are to be original messages. A *single* error in a message moves the message point along one edge of the imagined cube to an immediately adjacent point. If we require that every possible originating message be at least a *distance* of two sides away from any other message point, then it is clear that any single error will move a message only one side away and leave the received message as an illegitimate message. If the minimum distance between message points is three sides of the cube, then any single error will leave the received message *closer* to the original message than to any other message and thus we can have single error correction.

Effectively, we have introduced a *distance function*, which is the minimum number of sides of the cube we have to traverse to get from one point to another. This is the same as the number of bits in the representations of the two points that differ. Thus the distance can be looked on as being the number of 1's in the *logical* difference, or sum, of the two points. It is a legitimate distance function since it satisfies the following three conditions:

1. The distance from a point to itself is 0.
2. The distance from point x to point y is the same as the distance from y to x and is a positive number.
3. The triangle inequality holds—the sum of two sides (distance from a to c plus the distance from c to b) is at least as great as the length of the third side (the distance from a to b).

This distance function is usually called the *Hamming distance*. It is the distance function for *binary white noise*.

Using this distance function, we can define various things in the space. In particular, the *surface of a sphere* about a point is the set of points a given distance away. The surface of sphere of radius 1 about the point $(0, 0, 0, \dots, 0)$ is the set of all vertices in the space which are one unit away, that is, all vertices that have only one 1 in their coordinate representation (see Figure 3.6-1). There are $C(n, 1)$ such points.




Figure 3-6.1 Three-dimensional spheres about $(0, 0, 0)$ and $(1, 1, 1)$

We can express the minimum distance between vertices of a set of message points in terms of the error correctability possible. The minimum distance must be at least 1 for uniqueness of the code (Table 3.6-1). A minimum distance of 2 gives single-error detectability. A minimum distance of 3 gives single-error correctability; any single error leaves the point closer to where it was than to any other possible message. Of course, this minimum-distance code could be used instead for double-error detection. A minimum distance of 4 will give both single-error correction plus double-error detection. A minimum distance of 5 would allow double-error correction. *Conversely*, if the required de-

TABLE 3.6-1 *Meaning of Minimum Distance*

Minimum Distance	Meaning
1	Uniqueness
2	Single-error detection
3	Single-error correction (<i>or</i> double-error detection)
4	Single-error correction plus double-error detection (<i>or</i> triple-error detection)
5	Double-error correction
etc.	

gree of detection or correction is to be achieved, the corresponding minimum distance between message points must be observed.

In the case of single-error correction with the minimum distance of 3, we can surround each message point by a unit sphere and not have the spheres overlap. The *volume* of a sphere of radius 1 is the center plus the n points with just one coordinate changed, a volume of $1 + n$. The total volume of the n -dimensional space is clearly 2^n , the number of possible points. Since the spheres do not overlap, the maximum number of message positions k must satisfy

$$\frac{\text{Total volume}}{\text{Volume of a sphere}} \geq \text{maximum number of spheres}$$

or

$$\frac{2^n}{n + 1} \geq 2^k \quad (3.6-1)$$

Since

$$\begin{aligned} n &= m + k \\ 2^{m+k} &\geq 2^k(n + 1) \end{aligned}$$

or

$$2^m \geq (n + 1)$$

This is the same inequality as we derived from the algebraic approach, equation (3.4-1).

We can now see the restrictions for higher error correction. Thus for double-error correction we must have a minimum distance of 5, and we can put nonoverlapping spheres of radius 2 about each message position. The volume of a sphere of radius 2 is the center position plus the n positions a distance 1 away, plus those with two coordinates out of the n changed, which is the binomial coefficient $C(n, 2) = n(n - 1)/2$. Dividing the total volume of the space, 2^n , by the volume of these spheres gives an upper bound on the number k of possible code message positions in the space

$$\frac{2^n}{1 + n + n(n - 1)/2} \geq 2^k \quad (3.6-2)$$

It does not mean that this number can be achieved, only that this is an upper bound. Similar inequalities can be written for larger spheres.

When the spheres about the message points completely exhaust the space of 2^n points, leaving no points outside some sphere, the code is called a *perfect code*. Such codes have a high degree of symmetry (in the geometric model each point is equivalent to any other point) and a particularly simple theory. In only a comparatively few cases do perfect codes exist. They require that the inequalities be equalities.

;

Exercises

- 3.6-1** Extend the bounds of (3.6-1) and (3.6-2) to higher error-correcting codes.
- 3.6-2** Make a table for the bound for double-error correction of equation (3.6-2) ($n = 3, 4, 5, \dots, 11$).

3.7 Single-Error-Correction Plus Double-Error-Detection Codes

It is seldom wise to use only single-error correction, because a double error would then fool the system in its attempts to correct, and the system would use the syndrome to correct the wrong place; thus there would be three errors in the decoded message. Instead, a single-error-correction plus a double-detection code makes a reasonably balanced system for many (but not all) situations. The condition for the double detection is that the minimum distance must be increased by 1; it must be 4.

To make a double-error-detecting code from a single-error-correcting code, we add one more parity check (and one more position): this check is over the whole message. Thus any single error will still produce the right syndrome and the added parity check will give a 1. A double error will now cause a nonzero syndrome, but leave this added parity check satisfied. This situation can then be recognized as a double error; namely, some syndrome appears, but the extra parity check is still correct (Table 3.7-1). It is easy to see that two points which were

TABLE 3.7-1 *Double-Error Detection*

Original Syndrome	New Parity Check	Meaning
0	0	Correct
0	1	Error in added position
Something	1	Original meaning
Something	0	Double error

at the minimum distance of 3 from each other in the original code had a different number of 1's in them, modulo 2. Thus their corresponding extra parity checks would be set differently, increasing the distance between them to 4.

The argument we have just given applies to both the algebraic and geometric approaches, and shows how the two tend to complement each other. We have not yet given any constructive method for finding the higher error-correcting codes; we have only given bounds on them. Chapter 11 is devoted to the elements of their construction. The full theory for error-correcting codes has been developed over the years and is very complex, so we will only indicate the general approach.

Notice that the theory *assumes* that the correcting equipment is working properly; it is only the errors in the received (computed) message that are being handled.

Exercise

3.7-1 Show that the argument to get extra error detection can be applied to any odd minimum distance to get the next-higher (even) minimum distance.

3.8 Hamming Codes Using Words

On many computers it is difficult to get at the individual bits of information, but the computer words are readily available. We can therefore think of the parity checks as being whole words; we logically add, not bits, but words. The logical sum over all the selected words plus the check should give a word with all 0's, and any failure marks the failure of the parity check. If we have, for example, a single-error-correcting code, then we get the location of the wrong word and can reconstruct it easily by, if nothing else, looking at the bits of the parity checks. It should be clear that we could simply add the parity-check failure word to the word in the location of the erroneous word to get the correct word. In this fashion we can correct the errors in any one word, whether it is one, two, . . . , or all the bits of the word that are wrong.

A mixed system results when you use the parity check on the ASCII code bits. The ASCII parity check locates the word with the error and the single-error-detecting parity word over all the words gives the *position in the word* that is wrong. Again, you simply add the parity check to the word that is wrong to do the correction. This amounts to a rectangular code. It will correct any odd number of errors in a single word.

3.9 Applications of the Ideas

In Section 1.2 we claimed that not only were the ideas in the text directly useful but they also had wider general application—we used evolution as an example.

The central idea of error detection and correction is that the meaningful messages must be kept far apart (in the space of probable errors) if we are to handle errors successfully. If two of the possible messages are not far enough apart, one message can be carried by an error (or errors) into the other, or carried at least so close that at the receiving end we will make a mistake in identifying the source.

In assigning names to variables and labels in a FORTRAN, COBOL, ALGOL, Pascal, Ada, and other high-level language, the names should be kept far apart; otherwise, an intellectual “slip of the pen” or a typical error in keying in the name can transfer one name to another meaningful name. If the names are made to differ in at least two positions, single typos will be caught by the assembler. Thus the use of short mnemonic names should be tempered by the prudent need to protect oneself against small slips.

If you estimate the time to locate an error in miswriting a name against the time to add one or more characters to all the names while writing the original program, you will see that using longer-than-minimum names is probably a very wise idea—but it is unlikely that this computation will convince many people to do so! Minimal-length names are the source of much needless confusion and waste of time, yours and the machine's.

The distance function in the Hamming codes is based completely on white noise. Sections 2.7 and 2.9 are included to emphasize that the proper distance function to use depends in some cases on the *psychological distance* between the names as well as more uniformly random keystroke errors.

3.10 Summary

We have given the fundamental nature of error detection and error correction for white noise, namely the minimum distance between message points that must be observed. We have given methods for constructing codes for:

Single-error detecting	min. dist. = 2
Single-error correction	min. dist. = 3
Single-error correction + double-error detection	min. dist. = 4

Their design is easy, and they are practical to construct in software or in hardware chips. They can compensate for weak spots in a system, or can be used throughout an entire system, to get reliable performance out of unreliable parts. One need not put up with poor equipment performance, but the price is both in storage (or time of transmission) and equipment (or time) to encode and possibly correct. You do not get something for nothing! The codes also make valuable contributions to maintenance since they pinpoint the error, and the repairmen will not try fixing the wrong things (meaning “fix” what is working right and ignore what is causing the error!). A more widespread use of the idea of distance between messages was sketched in Section 3.9.

The use of such codes, and more highly developed codes, is rapidly spreading as we get faster and smaller integrated-circuit components. Increasingly, in VLSI (very large system integration) chips the code is part of the hardware.

Variable-Length Codes: Huffman Codes

4.1 Introduction

We now turn to the topic of source encoding. The codes we have looked at so far have all used a fixed length, and they are called *block codes* from the fact that the messages are of fixed block lengths in the stream of symbols being sent. The Morse code mentioned in Chapter 1 is an example of a *variable-length* code, and we now examine variable-length codes in more detail. The advantage of a code in which the message symbols are of variable length is that sometimes the code is more *efficient* in the sense that to represent the same information we can use fewer digits on the average. To accomplish this, we need to know something about the statistics of the messages being sent. If every symbol is as likely as every other one, then the block codes are about as efficient as any code can be (see Section 4.6). But if some symbols are more probable than others, then we can take advantage of this to make the most frequent symbols correspond to the shorter encodings, and the less frequent symbols correspond to the longer encodings. This is exactly what the Morse code does. The letter E in the English language occurs most frequently, and corresponds to the encoded symbol “dot.”

However, variable-length codes bring with them a fundamental problem: At the receiving end, how do you recognize each symbol of the code? In a binary system, for example, how do you recognize the end of one code word and the beginning of the next? If the probabilities of the frequencies of occurrence of the individual symbols are sufficiently

different, then variable-length encoding can be significantly more efficient than block encoding.

In this chapter we take advantage only of the frequencies of occurrence of the individual symbols being sent, and neglect any larger structure the messages may have. An example of a slightly larger scale structure is the fact that in English the letter Q is usually followed by the letter U; in Chapter 5 we take advantage of such correlations in the message being sent.

Exercise

4.1-1 List advantages and disadvantages of variable-length codes.

4.2 Unique Decoding

We need to make clear when we are talking about the symbols to be sent and when we are talking about the symbols used by the signaling system. We will refer to the *source symbols* when we refer to the symbols (such as the letters of the English alphabet) that are to be sent, and to the *code's alphabet* when we refer to the symbols used in sending (such as 0 and 1 in the binary system). In general, we will assume that the source alphabet being sent has q symbols, $s_1, s_2, \dots, s_q$, and that the code's alphabet has r symbols (r for the radix of the system).

The first property that we need is *unique decodability*—the received message must have a single, unique possible interpretation. Consider a code in which the source alphabet S has four symbols, and they are to be encoded in binary as follows:

$$s_1 = 0$$

$$s_2 = 01$$

$$s_3 = 11$$

$$s_4 = 00$$

The particular received message 0011 could be one of these two:

$$0011 = \begin{cases} s_4, s_3 \\ s_1, s_1, s_3 \end{cases}$$

Thus the code is not uniquely decodable. Although this property of unique decodability is not always absolutely essential, it is usually highly desirable.

To get our thoughts clear, we make a formal definition:

Definition. The n th extension of a code is simply all possible concatenations of n symbols of the original source code.

This is also called the n th *direct product* of the code. There are q^n symbols in the n th extension. The definition is necessary because the messages I send look to you, the receiver, as concatenations of the encoded symbols of the source alphabet, and you need to decide which sequence of source symbols I sent. For unique decodability, no two encoded concatenations can be the same, even for different extensions. Clearly, only if every distinct sequence of source symbols has a corresponding encoded sequence that is unique can we have a uniquely decodable signaling system. This is a necessary and sufficient condition, but it is hardly usable in this form.

Exercises

4.2-1 Is the code 0, 01, 001, 0010, 0011 uniquely decodable?

4.2-2 Is the code 0, 01, 011, 111 uniquely decodable?

4.3 Instantaneous Codes

Consider the following code:

$$s_1 = 0$$

$$s_2 = 10$$

$$s_3 = 110$$

$$s_4 = 111$$

Now consider how you as the receiver would decode messages sent in this code. You would set up what is equivalent to a *finite automaton*, or if you prefer, a *decision tree*. Starting in the initial state (Figure 4.3-1), the first binary digit received will cause a branch, either to a terminal state s_1 if the digit is 0, or else to a second decision point if it is a 1. For the next binary digit this second branch would go to the terminal state s_2 if a 0 is received, and to a third decision point if it is a 1. The third would go to the terminal state s_3 if the third digit is a 0, and to the terminal symbol s_4 if it is a 1. Each terminal state would,




Figure 4-3.1 Decoding tree

of course, emit its symbol and then return control to the initial state. Note that each bit of the received stream is *examined only once*, and that the terminal states of this tree are the four source symbols s_1 , s_2 , s_3 , and s_4 .

In this example the decoding is *instantaneous* since when a complete symbol is received, you, the receiver, immediately know this, and do not have to look further before deciding what message symbol you received. No encoded symbol of this code is a prefix of any other symbol. This particular type of code is what is called a *comma code* since the binary digit 0 indicates the end of a symbol, *plus* the fact that in this case no symbol is longer than three digits.

From this you see the basic equivalence of the existence of the decoding tree and the instantaneous decodability; each implies the other. Notice again that using the decoding tree means that each received symbol is looked at only once in the decoding process.

The following code is uniquely decodable but is not instantaneous because you do not know when one symbol is over without looking further:

$$s_1 = 0$$

$$s_2 = 01$$

$$s_3 = 011$$

$$s_4 = 111$$

(It is the previous code with the bits reversed.) We can see the trouble—some code words are *prefixes* of other words; that is, they are the same

as the beginning part of some other symbol. Now consider the string

$$0111 \dots \underbrace{1111}_{s_4} \underbrace{1111}_{s_4}$$

It *can only* be decoded by first going to the end and then identifying runs of three 1's as each being an s_4 until the first symbol is reached. Only then can it be identified as s_1 , s_2 , or s_3 . Thus you cannot decide whether you have the word that corresponds to the prefix, or if you must wait until more is sent to complete the word. The simplest way to decode messages in this particular code is always to start at the back end of the received message! This puts a severe burden on the storage and also causes a time delay.

In view of the existence of the decoding tree, *it is clearly both necessary and sufficient that an instantaneous code have no code word s_i which is a prefix of another code word, s_j .* If we had to deal with a noninstantaneous code, then the attempted decoding tree would have a structure which instead of returning to the start would, when it realized (finally) that it had received a complete word some time ago, emit the proper code word and then go to an appropriate place in the tree, not necessarily the start. As in the example above, the tree could require potentially infinite storage.

4.4 Construction of Instantaneous Codes

It is clear that of all uniquely decodable codes, the instantaneous codes are preferable to ones that are not, and since it will turn out (Section 4.7) that they cost us nothing extra, it is worth concentrating on them. Let us, therefore, explore the problem of constructing them.

Given that we are to construct a code with five symbols s_i in the source code S , and that the code alphabet is binary, we can assign

$$s_1 = 0$$

$$s_2 = 10$$

$$s_3 = 110$$

$$s_4 = 1110$$

$$s_5 = 1111$$

to get an instantaneous (comma) code (Figure 4.4-1).




Figure 4-4.1 Decoding tree

In this construction the use of the 0 for the first symbol reduced the number of possibilities available later. Instead of this, let us use two digits for the first two symbols, $s_1 = 00$ and $s_2 = 01$. We can now use $s_3 = 10$. With two symbols yet to go, we cannot use $s_4 = 11$; instead, we must try $s_4 = 110$, which leaves us with $s_5 = 111$. We get the code

$$s_1 = 00$$

$$s_2 = 01$$

$$s_3 = 10$$

$$s_4 = 110$$

$$s_5 = 111$$

This code is clearly instantaneous since no symbol is a prefix of any other symbol, and the decoding tree is easily constructed (Figure 4.4-2).

Which of these two codes is better (more efficient)? This depends on the frequency of occurrence of the symbols s_i . By “better” we mean, of course, that the sent messages *on the average* will be shorter, more efficient. We will investigate this more closely in Section 4.8. Evidently, “efficient” *must depend on the probability of the various symbols being used in the messages being sent.*




Figure 4-4.2 Decoding tree

Exercises

4.4-1 Devise a similar example using six words.

4.4-2 Find probabilities that favor one or the other of the two codes in this section. Ans. $p_3 + p_4 + p_5$ vs. p_1

4.5 The Kraft Inequality

The Kraft inequality, which we now examine, gives conditions on the existence of instantaneous codes; it tells when the *lengths* of the code words permit forming an instantaneous code, but it does not discuss the code itself.

Theorem. A necessary and sufficient condition for the existence of an instantaneous code S of q symbols s_i ($i = 1, \dots, q$) with encoded words of lengths $l_1 \leq l_2 \leq l_3 \leq \dots \leq l_q$ is

$$\sum_{i=1}^q \frac{1}{r^{l_i}} \leq 1 \quad (4.5-1)$$

where r is the radix (number of symbols) of the alphabet of the encoded symbols.

This inequality is asserting that we cannot have too many short encoded symbols. Most of the l_i must be reasonably large.

It is easy to prove the Kraft inequality from the decoding tree, whose existence follows from the instantaneous decodability. We proceed by induction. For simplicity consider first the binary case (see Figure 4.5-1). For a tree whose maximum length is 1, we have either




Figure 4-5.1 Proof of Kraft inequality

one or two branches of length 1. Thus we have either (for one symbol)

$$\frac{1}{2} \leq 1$$

or (for two symbols)

$$\frac{1}{2} + \frac{1}{2} \leq 1$$

We next assume that the Kraft inequality is true for all trees of length less than n . Now given a tree of maximum length n , the first node leads to a pair of subtrees of length at most $n - 1$, for which we have the inequalities $K' \leq 1$ and $K'' \leq 1$, where K' and K'' are the values of their respective sums. Each length l_i in a subtree is increased by 1 when the subtree is joined to the main tree, so an extra factor of $\frac{1}{2}$ appears. We have, therefore,

$$\frac{1}{2}K' + \frac{1}{2}K'' \leq 1$$

For radix r instead of binary, we have at most r branches at each node—at most r subtrees each with an extra factor of $1/r$ when joined to the main tree. Again the theorem is true.

When can an inequality occur? A moment's inspection shows that if every terminal node of the tree is a code word, then $K = 1$. It is only when some terminal nodes are not used that the inequality occurs. But if any terminal node is not used *for a binary code alphabet*, the preceding decision is wasted and that corresponding digit can be removed from every symbol that passes through this node in its decoding. Thus if the inequality holds, the code is inefficient, and how to correct for this is immediately evident *for binary trees*. Thus $K = 1$ for binary trees with all the terminals used. It is only for radix $r > 2$ that it is reasonable to have unused terminals and hence a K less than 1. Since the theorem gives a condition on the lengths only, the main use is in questions of the existence of a code with a given set of lengths.

Note again that the theorem refers to the existence of such a code, and does not refer to a particular code. A particular code may obey the Kraft inequality and still not be instantaneous, but there will exist codes that have the l_i and are instantaneous.

In the next two examples we will assume that only the code word lengths are given since this is what matters in the theorem, not the actual code words. If the lengths when encoded in binary symbols are 1, 3, 3, 3, the Kraft sum will be $\frac{1}{2} + \frac{3}{8} = \frac{7}{8}$, and an instantaneous code with those lengths is possible. One of the words of length 3 could be shortened to 2 bits. But if the lengths were 1, 2, 2, 3, the sum would be $\frac{1}{2} + 2(\frac{1}{4}) + \frac{1}{8} = \frac{9}{8}$, and such an instantaneous code could not exist.

Let us apply the theorem to yet another example. Suppose that we pick the radix $r = 3$ and want words of lengths 1, 2, 2, 2, 2, 2, 3, 3, 3, 3. The Kraft inequality gives on the left side $\frac{1}{3} + 5(\frac{1}{9}) + 4(\frac{1}{27}) = \frac{28}{27}$, so we cannot hope to find an *instantaneous code* with these constraints. If we drop the last code word of length 3, the sum would be exactly 1 and we can find such a code. To find an instantaneous code with these lengths, we proceed systematically: $s_1 = 0, s_2 = 10, s_3 = 11, s_4 = 12, s_5 = 20, s_6 = 21, s_7 = 220, s_8 = 221, s_9 = 222$. In this construction we have followed the decoding tree (Figure 4.5-2), and have systematically increased the ternary (base 3) number equivalents of the codes so that the argument is easily followed. The reader should realize that the three symbols 0, 1, and 2 are arbitrary and are not numbers. Thus any interchanges of the three symbols at any stage would leave the code essentially the same; it merely interchanges the order of the branches at a node. Indeed, one could interchange the order at one node without doing so at another and still have an equivalent code.




Figure 4-5.2 Decoding tree

Exercises

- 4.5-1** Does the infinite-length comma code $l_1 = 1, l_2 = 2, \dots, l_k = k, \dots$ satisfy the Kraft inequality? (Use $r = 2$.)
- 4.5-2** Generalize Exercise 4.5-1 to radix r .
- 4.5-3** Carry out the computation of K for a comma code.

4.6 Shortened Block Codes

Returning for the moment to the earlier block codes, if we had exactly 2^m code words in a binary system (r^m in a radix r system), we could use m digits to represent each symbol. But suppose that we do not have an exact power of the radix but still want the maximum length to be as short as possible. To see what can happen, consider the case of five symbols. Of the eight binary symbols

000
001
010
011
100
101
110
111

we can drop any three. But if we drop 001, 011, and 101, we can shorten three branches of the decoding tree and still have instantaneous decodability. We will have

$$s_1 = 00$$

$$s_2 = 01$$

$$s_3 = 10$$

$$s_4 = 110$$

$$s_5 = 111$$

(See Figure 4.6-1.) Instead of this choice, we can drop 001, 010, and 011 and shorten only one branch of the tree to the code.




Figure 4-6.1 Decoding tree

$$s_1 = 0$$

$$s_2 = 100$$

$$s_3 = 101$$

$$s_4 = 110$$

$$s_5 = 111$$

(See Figure 4.6-2.) In both cases there now are no unused terminals and therefore $K = 1$. We will call these *shortened block codes*; they are essentially block codes with small modifications.




Figure 4.6-2

Exercises

4.6-1 Discuss the case of two symbols in a ternary base.

4.6-2 Discuss the case of five symbols in the ternary base.

4.7 The McMillan Inequality

The Kraft inequality applies to instantaneous codes, which are a special case of *uniquely decodable codes*. McMillan showed that the same inequality applies to uniquely decodable codes. The underlying idea of the proof of the necessity is that very high powers of a number greater than 1 grow rapidly. If we can tightly bound this growth, then we know that the number is not greater than 1. The proof of the sufficiency follows from the fact that we can do it for instantaneous codes which are special cases of uniquely decodable codes.

The necessity part of the proof begins by taking the n th power of the Kraft expression

$$\left[\sum_{i=1}^q \frac{1}{r^{l_i}} \right]^n \equiv K^n$$

When we expand the left-hand term we find that we have a sum of many terms to various powers: the exponents running from n , the lowest possible power, to nl , the highest, where l is the length of the longest symbol. Thus we have the expression

$$K^n = \sum_{k=n}^{nl} \frac{N_k}{r^k}$$

where N_k is the number of code symbols (of radix r) of length k . Since the code is uniquely decodable, N_k cannot be greater than r^k , which is the number of distinct sequences of length k in our code alphabet of radix r . Therefore, we have the bound

$$K^n \leq \sum_{k=n}^{nl} \frac{r^k}{r^k} = nl - n + 1 < nl$$

(The $+1$ comes from the fact that both end terms in the sum are counted.) This is the inequality we need, since for any $x > 1$ a sufficiently large n makes the number $x^n > nl$. But n can be chosen as large as we please, and it follows therefore that the number K (the Kraft sum) must be ≤ 1 .

From this we see that there is very little to gain from avoiding instantaneously decodable codes and using the more general uniquely decodable codes; both have to satisfy the same Kraft inequality on the lengths of the encoded symbols.

4.8 Huffman Codes

Now for the first time we will make use of the *probabilities* of the various symbols being sent. As in the Morse code, we want the most frequent symbols to have the shortest encodings. If the probability of the i th symbol is p_i and its length is l_i , then the average length of the code is

$$L_{av} = \sum_{i=1}^q p_i l_i$$

With no loss in generality the p_i may be taken in decreasing order. If the lengths l_i are not in the opposite order, that is, we do not have both

$$p_1 \geq p_2 \geq p_3 \geq \dots \geq p_q$$

and

$$l_1 \leq l_2 \leq l_3 \leq \dots \leq l_q$$

then the code is not *efficient* in the sense that we could have a shorter average length by reassigning the code representations of the symbols $s_1, s_2, s_3, \dots, s_q$. To prove this assertion, suppose that for $m < n$ we have both conditions (for some m and n)

$$p_m > p_n \quad \text{and} \quad l_m > l_n$$

In computing the average length, we have, among others, the two terms

$$\text{Old: } p_m l_m + p_n l_n$$

By interchanging the encoded symbols for s_m and s_n we get the corresponding terms

$$\text{New: } p_m l_n + p_n l_m$$

Subtracting the *Old* from the *New* we have the change due to this reassignment:

$$\text{New} - \text{Old: } p_m(l_n - l_m) + p_n(l_m - l_n) = (p_m - p_n)(l_n - l_m) < 0$$

From the foregoing assumptions this is a negative number; we will decrease the average code length if we interchange the encoded symbols for s_m and s_n . We therefore assume that the two running inequalities both hold.

We begin our examination of Huffman coding with encoding into the binary alphabet. In Section 4.11 we look at the base r alphabet. We will use “source symbol” for the input s_i and “code alphabet” for the alphabet into which we are encoding.

The first thing to prove is that the two least frequent source symbols of an efficient code have the same encoded lengths. Suppose that the maximum coded symbol has length l . If there is only one of such length, then since the code is instantaneous, any shorter coded symbol of length $l - 1$ (or shorter) is not a prefix of the maximum-length coded symbol. Therefore, the last part of the longest symbol could be dropped with no loss of information in decoding. Thus the two longest symbols must have the same length, and because of the running inequalities, they must be the two least probable.

The proof of the encoding properties, as well as the method of encoding, is a reduction at each stage to a shorter code. We simply

4.8-2 shows the corresponding splitting (expansion) process. The general case should be obvious from this.

How do we know that this process generates an efficient code? Suppose that there were a shorter code with code length L' with

$$L' < L$$

Let us compare the two decoding trees. In an efficient binary code, all the terminals are occupied and there are no “dead branches.” (A dead branch would enable us to shorten the code by deleting the corresponding binary digit in all the terminals that pass through this useless decision point.)

If there are only two symbols of maximum length in a tree, they must have their last decision node in common, and they must be the two least probable symbols. Before we reduce a tree, the two symbols contribute

$$l_q(p_q + p_{q-1})$$

and after the reduction they contribute

$$(l_q - 1)(p_q + p_{q-1})$$

so that the code length is reduced by

$$p_q + p_{q-1}$$

If there are more than two symbols of the maximum length, we can use the following proposition: Symbols having the same length may be interchanged without changing the average code length. Using this, we can bring the two least probable symbols so that they share their final decision node. Thus after the reduction, we have shortened the code by the amount

$$p_q + p_{q-1}$$

Therefore, in either code we shorten the code and decrease the average code length *by the same amount*.

We apply this to both the decoding trees we are comparing. Since both are decreased by the same amount, the amount of inequality between their lengths is preserved. Repeated application of this will reduce both trees to two symbols. In the Huffman code the length is 1; for the other it must be less than 1, which is impossible. Therefore, for binary codes the Huffman code is the shortest possible code.

The encoding process is not unique in several respects. First, the assignment of the 0 or 1 symbols to the two source symbols at each splitting stage is arbitrary, but this produces only trivial differences. Second, when two probabilities are equal, it appears to be a matter of indifference which we put above the other in the table, although the resulting codes can have different lengths of words. However, in both cases the *average* length of the encoding of messages in these codes will be the same.

As an example of two different Huffman encodings of the same source, let

$$p_1 = 0.4$$

$$p_2 = 0.2$$

$$p_3 = 0.2$$

$$p_4 = 0.1$$

$$p_5 = 0.1$$

If at each stage we put the merged states as low as possible, then in Figure 4.8-1 we get lengths (1, 2, 3, 4, 4) and the average length is

$$L = 0.4(1) + 0.2(2) + 0.2(3) + 0.1(4) + 0.1(4) = 2.2$$

On the other hand, if we push the merged states up as high as possible (Figure 4.8-3), we will get lengths (2, 2, 2, 3, 3) and the average length is

$$L = 0.4(2) + 0.2(2) + 0.2(2) + 0.1(3) + 0.1(3) = 2.2$$




Figure 4-8.3 Alternative coding

Both codes have the same efficiency (average length) but not the same set of lengths of the symbols.

Which one of these two codes should we choose? A very reasonable choice is the one whose average length would vary least over the ensemble of messages. We therefore compute the variances in the two cases:

$$\begin{aligned} \text{Var(I)} &= 0.4(1 - 2.2)^2 + 0.2(2 - 2.2)^2 + 0.2(3 - 2.2)^2 \\ &\quad + 0.1(4 - 2.2)^2 + 0.1(4 - 2.2)^2 = 1.36 \end{aligned}$$

$$\begin{aligned} \text{Var(II)} &= 0.4(2 - 2.2)^2 + 0.2(2 - 2.2)^2 + 0.2(2 - 2.2)^2 \\ &\quad + 0.1(3 - 2.2)^2 + 0.1(3 - 2.2)^2 = 0.16 \end{aligned}$$

Thus the second code has significantly less variability in use on finite-length messages, and is therefore probably the preferable code. It is true that always moving a combined state as high as possible will give a minimum variance code, although we do not prove it here. By moving a merged symbol higher than it should be, we can greatly reduce the variance for a small increase in the average length.

Exercises

4.8-1 Give the Huffman code for $p_1 = \frac{1}{2}$, $p_2 = \frac{1}{3}$, and $p_3 = \frac{1}{6}$.

4.8-2 Give the Huffman code for $p_1 = \frac{1}{3}$, $p_2 = \frac{1}{4}$, $p_3 = \frac{1}{5}$, $p_4 = \frac{1}{6}$, and $p_5 = \frac{1}{20}$.

4.8-3 Give the Huffman code for $p_1 = \frac{1}{2}$, $p_2 = \frac{1}{4}$, $p_3 = \frac{1}{8}$, $p_4 = \frac{1}{16}$, and $p_5 = \frac{1}{16}$.

4.8-4 Give the Huffman code for $p_1 = \frac{27}{40}$, $p_2 = \frac{9}{40}$, $p_3 = \frac{3}{40}$, and $p_4 = \frac{1}{40}$.

4.8-5 Give the Huffman code for the probabilities $\frac{1}{3}$, $\frac{1}{4}$, $\frac{1}{6}$, $\frac{1}{8}$, $\frac{1}{12}$, and $\frac{1}{24}$.

4.8-6 Show two distinct decoding trees for the case $\frac{1}{3}$, $\frac{1}{3}$, $\frac{1}{6}$, and $\frac{1}{6}$.

4.9 Special Cases of Huffman Coding

There are a number of interesting cases of Huffman coding to look at. First, if all the symbols are equally likely and if there are exactly $q = 2^m$ source symbols, then the (binary) Huffman code will be a block code




Figure 4-9.1

with all symbols having the same length m (see Figure 4.9-1). If there are not exactly 2^m source symbols we will get a shortened block code (Section 4.6).

The second case is more interesting. Suppose that the two least probable symbols have their probabilities such that their sum is greater than the most probable symbol; that is, $p_{q-1} + p_q > p_1$. In the Huffman process (see Figure 4.9-1) the corresponding symbol word will go to the top. It then follows that the next two least probable symbols will combine and will go to the top, and so on. If the number of original symbols q is an exact power of 2, then the repeated combining process into the final code of two symbols will have each symbol go through *the same number* of splitting steps, and each symbol will acquire the same number of binary digits in the expansion process. Thus again we will have a block code. If q is not an exact power of 2, then, of course, there are small modifications to the foregoing process, and we get a shortened block code (Section 4.6).

It is only when the probabilities of the source symbols of the message are very different that we get a significant economy from the Huffman encoding process. This is what we would expect if we but thought a bit—only when there are large differences in the probabilities of occurrence does the variable length of the code symbols pay off.

For example, if the probabilities p_i have great variability,

$$p_j \geq \frac{2}{3} \sum_{k=j+1}^q p_k \quad (4.9-1)$$

for all j , then a comma code will emerge. In words, equation (4.9-1) says that each probability is at least as great as two-thirds of the sum of all the probabilities that follow.*

*Koh Wee Jin observed that (4.9-1) can be improved to $p_j \geq \sum_{k=j+2}^q p_k$




Figure 4-9.2 The two possible cases

To prove this, we begin at the bottom of Figure 4.9-2. The sum $p_q + p_{q-1}$ can be greater or less than p_{q-2} , but it cannot be greater than p_{q-3} . Thus at the next stage we get the sum

$$p_q + p_{q-1} + p_{q-2}$$

and p_{q-2} has been combined only once. Repeating this at each stage, we see that we can get a comma code. The code

$$p_1 = \frac{1}{2}$$

$$p_k = \frac{1}{2^k} \quad (k < q)$$

$$p_q = \frac{1}{2^{q-1}}$$

is such a code and permits a comma code realization.

If at any stage of reduction one of these conditions applies, then you know that beyond that point the corresponding splitting process will apply. This can save time in some cases.

As a practical observation it is often convenient to add one symbol for your own use, an escape character (end of message) (see Section

1.10). In the Huffman decoding tree, each other terminal returns the control to start of the tree, but this escape state allows you to escape the decoding process and resume other tasks. Obviously, you assign a very low probability to this escape character so that it has little effect on the rest of the code.

The *Zipf distribution*, where the k th item ($k = 1, 2, \dots, q$) has a frequency of $1/k$, is very common in language problems. Of course, the distribution is cut off someplace or else there would be an infinite number of items (and the infinite series $1/k$ does not converge). Let us compare the performance of a standard block code where each symbol has the same length with what Huffman encoding can do. Table 4.9-1 gives the results, where the first column gives the number of items, the second the length of each symbol in bits, and the third the Huffman encoding average length. The last two columns give the gain and the percent gain in using Huffman codes even in this rather unfavorable case where there is not a great deal of variation in the probabilities.

If you examine the Huffman encoding method, it soon becomes clear that you can work with any numbers that are proportional to the probabilities. Thus you may use the *frequencies* rather than the probabilities, or you can multiply through by some suitable constant to convert the probabilities to integers to make the work easier. Your escape character should be given a frequency of 1 to minimize its effect.

TABLE 4.9-1 *Zipf Distribution Encoded in a Huffman Code*

n	Block Code Length	Huffman L	Gain	Percent Gain
2	1	1.0	0	0
4	2	1.8	0.2	10
8	3	2.68	0.32	11
16	4	3.43	0.57	14
32	5	4.17	0.83	17
64	6	4.89	1.11	18.5
128	7	5.60	1.50	20
256	8	6.26	1.74	22
512	9	6.90	2.10	23
1024	10	7.54	2.5	25

Exercises

4.9-1 (Difficult) Discuss what happens in cases that lie between the two main results.

4.9-2 Verify first three lines of Table 4.9-1.

4.9-3 Discuss using frequency = 1 for an escape character.

4.10 Extensions of a Code

In Section 4.2 we defined an extension of a code. We are now in a position to see why they are useful. Suppose for the purposes of a concrete example that we have two symbols s_1 and s_2 with probabilities $p_1 = \frac{2}{3}$ and $p_2 = \frac{1}{3}$.

<i>Encoding</i>		
s_1 ;	$p_1 = \frac{2}{3}$	$s_1 \longrightarrow 0$
s_2 ;	$p_2 = \frac{1}{3}$	$s_2 \longrightarrow 1$

The average code length is 1.

The second extension takes two symbols at a time (the first extension of a code is itself):

<i>Encoding</i>		
s_1s_1 ;	$p_{1,1} = \frac{4}{9}$	$s_1s_1 \longrightarrow 1$
s_1s_2 ;	$p_{1,2} = \frac{2}{9}$	$s_1s_2 \longrightarrow 01$
s_2s_1 ;	$p_{2,1} = \frac{2}{9}$	$s_2s_1 \longrightarrow 000$
s_2s_2 ;	$p_{2,2} = \frac{1}{9}$	$s_2s_2 \longrightarrow 001$

Here $p_{1,1}$ is four times as probable as $p_{2,2}$. The average length for Huffman encoding is $\frac{17}{18} = 0.94444 \dots$ for comparison. The next extension takes three symbols at a time gives $\frac{76}{81} = 0.93827 \dots$. Taking four s_i at a time leads to the average code length 0.93827. . . .

Thus we see that extensions increase the variability and hence allow Huffman coding to achieve some compression. We will show in Section 6.9 that a lower bound on what can be achieved for this example is 0.91830. . . .

Exercise

4.10-1 Compute the average code length of the extensions given in this section.

4.11 Huffman Codes Radix r

When we go to a radix r code the modifications to the Huffman encoding scheme are easy to make. In the binary code we always found at each stage the number of items (two) to combine into the merged state. When we are in radix $r > 2$ this is not necessarily true, and we need to think first.

At each stage when we combine r symbols into a single symbol (with the probability that is the sum of the probabilities of the individual symbols) we replace r symbols with a single symbol. Thus each merging of r states reduces the number of states by $r - 1$. But at the last state we will want to have a full complement of r states to assign the encoding symbols $0, 1, 2, \dots, r - 1$. Thus we want to have

$$k(r - 1) + r$$

states when we begin, where k is some integer. However, this is the same as

$$k(r - 1) + 1$$

(the integer k has been increased by 1 in this new notation). Often we will not have that number, so we will add enough states each with probability $p = 0$ until we do have the needed number.

Figure 4.11-1 gives one concrete example for base 4 (from Huff-

s_i		p_i						
s_1	1	0.22	1	0.22	1	0.40	0	1.00
s_2	2	0.20	2	0.20	2	0.22	1	
s_3	3	0.18	3	0.18	3	0.20	2	
s_4	00	0.15	00	0.15	00	0.18	3	
s_5	01	0.10	01	0.10	01			
s_6	02	0.08	02	0.08	02			
s_7	030	0.05	030	0.07	03			
s_8	031	0.02	031					

Figure 4-11.1 Huffman radix 4

man's original paper), where we have added two symbols s_9 and s_{10} , both with probability zero. After two reduction stages we reach four symbols, which we encode as 0, 1, 2, $r - 1 = 3$. The backward path of splitting goes as for radix 2 with the obvious changes.

For the proof of the efficiency of the encoding we first, if necessary, again add enough symbols of probability zero to the alphabet of symbols s_i so that at each stage r symbols can be combined. The trees will therefore have all their terminals filled, and we will be able to argue, as in the case of the binary trees, that no other encoding can be more efficient.

Exercises

4.11-1 Use the data of Figure 4.11-1 and encode in radix 3.

4.11-2 Given $p_1 = \frac{1}{3}$, $p_2 = \frac{1}{3}$, $p_3 = \frac{1}{6}$, $p_4 = \frac{1}{10}$, $p_5 = \frac{1}{12}$, $p_6 = \frac{1}{15}$, and $p_7 = \frac{1}{20}$, encode Huffman for $r = 2, 3$, and 4. Find average code lengths. Compare for $r = 2$ and 4.

Ans. $r_2 = 31/12 \sim 3$

4.11-3 In Exercise 4.11-2, after multiplying each probability by 240 to get equivalent frequencies and an escape character of frequency 1, find code for $r = 2$ and also draw the decoding tree.

4.11-4 Do Exercise 4.11-3 for radix 4. Compare with your answer in Exercise 4.11-3.

4.12 Noise in Huffman Coding Probabilities

Suppose that the estimates of the probabilities p_i are not accurate. How much does the average code length suffer? In short, will small errors in the estimates greatly degrade (or improve) the system?

To answer this question, let the p_i be the original Huffman code design probabilities and

$$p'_i = p_i + e_i \quad (4.12-1)$$

be the probabilities for the source that is actually used. Clearly,

$$\frac{1}{q} \sum_{i=1}^q e_i = 0 \quad (4.12-2)$$

since both p_i and p'_i are probability distributions that sum to 1.

As one measure of the size of the errors e_i , we use the variance

$$\frac{1}{q} \sum_{i=1}^q e_i^2 = \sigma^2 \quad (4.12-3)$$

Multiply (4.12-1) by l_i and divide to get the new average symbol length:

$$\begin{aligned} L' &= \frac{1}{q} \sum l_i p'_i = \frac{1}{q} \sum l_i p_i + \frac{1}{q} \sum l_i e_i \\ &= L + \frac{1}{q} \sum l_i e_i \end{aligned}$$

The last term is the change in the average code length.

With the two conditions (4.12-2) and (4.12-3) on the e_i , we resort to the method of Lagrange multipliers to find the *extreme* cases. The Lagrange expression is

$$\mathcal{L} = \frac{1}{q} \sum l_i e_i - \lambda \frac{1}{q} \sum e_i - \mu \left(\frac{1}{q} \sum e_i^2 - \sigma^2 \right)$$

with Lagrange multipliers λ and μ . The partial derivatives with respect to e_i are to be equated to zero; that is,

$$\frac{\partial \mathcal{L}}{\partial e_i} = \frac{1}{q} \left(l_i - \lambda - 2\mu e_i \right) = 0 \quad (i = 1, 2, \dots, q) \quad (4.12-4)$$

We now sum these equations to get

$$\frac{1}{q} \sum_i l_i - \lambda = 0 \quad (4.12-5)$$

which determines λ . Next we multiply equations (4.12-4) by e_i and sum to get

$$\frac{1}{q} \sum_i l_i e_i - 2\mu \frac{1}{q} \sum_i e_i^2 = 0 \quad (4.12-6)$$

which determines μ . Finally, we multiply equations (4.12-4) by l_i and sum to get

$$\frac{1}{q} \sum_i l_i^2 - \lambda \frac{1}{q} \sum_i l_i - 2\mu \frac{1}{q} \sum_i e_i l_i = 0$$

Using the λ from (4.12-5) and the μ from (4.12-6), we have

$$\frac{1}{q} \sum_i l_i^2 - \left[\frac{1}{q} \sum_i l_i \right]^2 - \left[\frac{1}{q} \sum_i e_i l_i \right]^2 / \left[\frac{1}{q} \sum_i e_i^2 \right] = 0$$

Using (4.12-3), this takes the form

$$\begin{aligned} \left(\frac{1}{q} \sum_i e_i l_i \right)^2 &= \left[\frac{1}{q} \sum_i l_i^2 - \left(\frac{1}{q} \sum_i l_i \right)^2 \right] \sigma^2 \\ &= [\text{variance of } (l_i)] [\text{variance of } (e_i)] \end{aligned} \quad (4.12-7)$$

These are the extreme (Chebyshev) values possible for a fixed σ^2 . The equation shows that the more variable the l_i , the more harm (or good) the errors in the estimates of the p_i could cause in the average of the symbol length.

That errors e_i can make no change in the average length can be seen from the following example. For $q \geq 3$ let, for some e ,

$$\begin{aligned} e_1 &= (l_2 - l_3)e \\ e_2 &= (l_3 - l_1)e \\ e_3 &= (l_1 - l_2)e \\ e_i &= 0 \quad (i > 3) \end{aligned} \quad (4.12-8)$$

Clearly,

$$\frac{1}{q} \sum e_i = 0$$

To find the corresponding σ^2 we have, from equation (4.12-3),

$$[(l_2 - l_3)^2 + (l_3 - l_1)^2 + (l_1 - l_2)^2]e^2 = \sigma^2$$

Now consider the average code-length change,

$$\begin{aligned} l_1 e_1 + l_2 e_2 + l_3 e_3 + 0 &= (l_1 l_2 - l_1 l_3 + l_2 l_3 - l_2 l_1 + l_3 l_1 - l_3 l_2)e \\ &= 0 \end{aligned}$$

Hence these changes e_i [equations (4.12-8)] do not change the average code length.

From equation (4.12-7) we conclude that Huffman coding is fairly robust; small changes in the probabilities do not seriously degrade the system. But we have *not* answered the question of how much a redesign using the new probabilities could have improved matters.

4.13 Use of Huffman Codes

The design of a Huffman code is highly recursive, and therefore is easy to program on a computer. The finite automata scheme (tree) of decoding mentioned in Section 4.2 provides a method for decoding Huffman codes using a single r -way branch per received code symbol of radix r . The method is simple and rapid and each received symbol is examined *only once*.

The encoding of a Huffman code is for reasonably sized source alphabets, a simple table lookup. If we have a lot of information to put into long-term backup storage in a computer, for example, and if the source code has symbols whose probabilities vary greatly, then it may well be worth encoding the information into an appropriate Huffman code to save long-term storage. As we have seen, it depends on the variability of the probabilities p_i .

For a particular hand-held language translator, a dictionary of about 12,000 bytes (8 bits) was compressed into an 8192-byte ROM storage device using compression methods resembling Huffman codes. The encoding of the operations in the Intel 432 used Huffman-like codes to compress the programs that are run.

Programs have been written to *automatically* scan a text, design a code, compress the text, and store the text along with the decoding algorithm. The program counts the frequencies of the source symbols and adds one special symbol called "end of text" of frequency 1 (see Section 4.9). This special symbol is used to end the decoding. Such programs have been known to reduce the storage of data to less than half the original amount.

We may also ask what happens to the decoding process when an error arises. It is evident that in a comma code the first appearance of a comma will resynchronize the decoding process. But in many codes it is not obvious how long you will go before you are resynchronized. A satisfactory theory is slowly emerging, but for this text about the best that can be done is to use Monte Carlo methods on the given code. Encode a number of typical messages, insert a change, deletion, or addition to the coded stream of digits, and then apply the decoding tree to see how long it is until the text is coming out correct again. In a

block code the insertion or deletion of a single bit will indefinitely alter the decoding process; a block code can never recover from a single such error.

Exercises

- 4.13-1** Discuss writing a computer program to find a Huffman code from the p_i .
- 4.13-2** Discuss the use of a decoding tree for Hamming codes rather than the computing methods of Chapter 3.
- 4.13-3** For the comma code $s_1 = 0$, $s_2 = 10$, $s_3 = 11$, show that an error in one bit can be propagated arbitrarily far. [*Hint*: Consider $s_1(s_3)^n s_2$.]

4.14 Hamming-Huffman Coding

How can we combine the source compression of Huffman encoding with the noise protection of Hamming encoding? As far as is now known, you encode in a Huffman code and take the output stream, break it up into blocks, and encode it into a suitable noise-protecting Hamming code. For decoding you reverse the process: decode Hamming, then Huffman.

A reasonable question is to ask if a single encoding process can do a better job. It is left to the reader to try. We note that in the decoding, tree nodes near a terminal must not be used by other branches of the tree.

Miscellaneous Codes

5.1 Introduction

There are many different codes for the many different situations that arise in practice. We can study only a few of the more important ones and hope that you will be able to create your own when the need arises.

We shall first look at *Markov processes*, which enable us to handle some of the larger probability structures that arise typically in languages, larger than the simple frequency of occurrence that is used by the Huffman code.

We will then examine *predictive coding*, which is very general and takes advantage of *any* structure you know about that can help you predict the next symbol accurately. The better your ability to predict the next symbol, the shorter the final message you have to use. Nothing is said about how you predict, only that you have written a program that incorporates all the structure that you are going to use. This is a very general method of code compression. Thus, as we said we would, we have gone to very general methods and have skipped the many special codes for special cases. However, it is hoped that the reader, when faced with a special case and having seen the codes we have given, will be able to invent an appropriate method of suitably encoding messages for compression.

We then discuss the remarkable idea of *random encoding*, both for its practical value and for its theoretical use in proving the main result of information theory. The particular method we examine is called

“hash coding” and is widely used in computer software, but has applications to other situations.

Finally, we will look at *Gray codes*, used to go reliably from analog-to-digital representations and back. Since this occurs so often in practice (because nature provides us basically with analog signals), it is a part of the education of computer science students and others who have to deal with “real” signals. We also show a simple variant useful in testing memories of computers. For more material, see Ref. [J].

5.2 What Is a Markov Process?

Up to this point we have used only the probabilities of occurrence of the source alphabet $s_1, s_2, \dots, s_q$. Familiarity with natural languages shows that there is often much more structure in a language. Thus, as we noted before, in English the letter Q is almost always followed by the letter U; the probability of seeing the letter U depends very much on the letter that came before it. In probability theory the notation

$$p(U \mid Q)$$

means “the probability of seeing U given that ($\mid$) you have just seen the letter Q.” In general, in the notation of a j th-order Markov process,

$$p(s_i \mid s_{i_1}, s_{i_2}, \dots, s_{i_j})$$

is the *conditional probability* of seeing s_i given that you have just seen all the letters $s_{i_1}, s_{i_2}, \dots, s_{i_j}$, in that order. *Note* that the sequence of symbols is

$$s_{i_1}, s_{i_2}, \dots, s_{i_j}, s_i$$

This is clearly a situation that arises in practice—the previous part of the message can often greatly influence the probabilities for the next symbol that you will see from the source.

In computer language a *zero memory source* uses each source symbol independently of what has preceded it. A *j-memory source* uses the previous j symbols, and this idea corresponds exactly to a j th-order Markov process.

We consider the system as being in some *state*. For a one-memory source there are q states, one for each symbol s_i of the source alphabet. The standard weather prediction that “tomorrow will be like today” is an example of a one-memory predictor. For a two-memory source there

are q^2 states, one state for each pair of source symbols. The linear trend predictor

$$y_{n+1} = 2y_n - y_{n-1}$$

is a two-memory source; the next value y_{n+1} is estimated from the two previous symbols y_n and y_{n-1} . In general, for a j -memory source there are q^j states in the Markov process.

As an example of a simple one-memory source, suppose that we have an alphabet of three symbols, a , b , and c . Let the probability that an a is followed by any of the three letters be each $\frac{1}{3}$. Let the probability that the letter b is followed by another b be $\frac{1}{2}$, and either of the other two letters, a or c , be $\frac{1}{4}$. Finally, let the probability that the letter c is followed by a c be $\frac{1}{2}$, and either of the other two, a , or b , be $\frac{1}{4}$. We have

$$\begin{array}{lll} p(a | a) = \frac{1}{3} & p(b | a) = \frac{1}{3} & p(c | a) = \frac{1}{3} \\ p(a | b) = \frac{1}{4} & p(b | b) = \frac{1}{2} & p(c | b) = \frac{1}{4} \\ p(a | c) = \frac{1}{4} & p(b | c) = \frac{1}{4} & p(c | c) = \frac{1}{2} \end{array}$$

It is conventional to use a *transition graph* to illustrate such a Markov process. There are, of course, three states in this case (Figure 5.2-1), labeled a , b , and c and indicated by the circles. Each directed line is a transition from one state to another state, whose probability is indicated by the number alongside the line. For example, $p(a | b)$ is the directed line from b to a and has the probability of transition of $\frac{1}{4}$. In this example each state has three lines “out” and three lines “in.”




Figure 5-2.1 Transition graph

We may assign a concrete interpretation to this diagram. Let a mean “fair weather,” b means “rain,” and c mean “snow.” Given that today is fair, it is equally likely that tomorrow it will be any of the three

states, fair, rain, or snow. But if today is either rain or snow it is a 50–50 chance that it will be the same tomorrow, and only 1 in 4 that it will be either one of the other two.

We may write the graph in a matrix form as follows, where the label of the current state is given by the letter on the left of the matrix and the next state after the transition is given above the matrix. The matrix entry is the *transition probability*. We have

$$\begin{array}{c} \longrightarrow \begin{array}{c} a \\ b \\ c \end{array} \begin{pmatrix} \frac{1}{3} & \frac{1}{3} & \frac{1}{3} \\ \frac{1}{4} & \frac{1}{2} & \frac{1}{4} \\ \frac{1}{4} & \frac{1}{4} & \frac{1}{2} \end{pmatrix} \end{array} \quad (5.2-1)$$

In a transition matrix the sum of the elements in any row must total exactly 1, since the current state must go somewhere.

Now suppose that instead of having a definite state today, we have a probability distribution for today's state (p_a, p_b, p_c) , where, of course, $p_a + p_b + p_c = 1$. It can be that one $p_i = 1$ and the other two are 0, which means that we have a definite state today. If we multiply this row vector on the right by the transition matrix (5.2-1), we get the probability of tomorrow's weather:

$$\begin{aligned} & (p_a, p_b, p_c) \begin{pmatrix} \frac{1}{3} & \frac{1}{3} & \frac{1}{3} \\ \frac{1}{4} & \frac{1}{2} & \frac{1}{4} \\ \frac{1}{4} & \frac{1}{4} & \frac{1}{2} \end{pmatrix} \\ &= \left[\frac{1}{3}p_a + \frac{1}{4}p_b + \frac{1}{4}p_c, \frac{1}{3}p_a + \frac{1}{2}p_b + \frac{1}{4}p_c, \frac{1}{3}p_a + \frac{1}{4}p_b + \frac{1}{2}p_c \right] \end{aligned}$$

To find the probabilities 2 days ahead, we multiply on the right again by the transition matrix (5.2-1). Because of the associativity of matrix (and vector) multiplication, this is the same as first multiplying the matrix by itself to get the square of the matrix, and then multiplying the original distribution on the right by this. We get, for the square of the transition matrix,

$$\begin{pmatrix} \frac{1}{3} & \frac{1}{3} & \frac{1}{3} \\ \frac{1}{4} & \frac{1}{2} & \frac{1}{4} \\ \frac{1}{4} & \frac{1}{4} & \frac{1}{2} \end{pmatrix} \begin{pmatrix} \frac{1}{3} & \frac{1}{3} & \frac{1}{3} \\ \frac{1}{4} & \frac{1}{2} & \frac{1}{4} \\ \frac{1}{4} & \frac{1}{4} & \frac{1}{2} \end{pmatrix} = \begin{pmatrix} \frac{10}{36} & \frac{13}{36} & \frac{13}{36} \\ \frac{13}{48} & \frac{19}{48} & \frac{16}{48} \\ \frac{13}{48} & \frac{16}{48} & \frac{19}{48} \end{pmatrix}$$

Writing this matrix in the form

$$\frac{1}{144} \begin{pmatrix} 40 & 52 & 52 \\ 39 & 57 & 48 \\ 39 & 48 & 57 \end{pmatrix} \quad (5.2-2)$$

we see immediately that the entries in the square of the transition matrix are less variable than were those of the original, (5.2-1). If we square this matrix (5.2-2), we get the matrix for predicting 4 days ahead. For this particular original transition matrix it is not hard to prove that the successive powers converge to a limiting matrix.

Do the successive transition matrices preserve the property that the sum of any row is exactly 1? To prove that they do, we have from the two general matrices P and P' the entries for the next one, P'' :

$$p''_{i,j} = \sum_k p_{i,k} p'_{k,j}$$

where by hypotheses the row sums for both matrices are

$$\sum_k p_{i,k} = 1 \quad \text{and} \quad \sum_j p'_{k,j} = 1$$

We compute, therefore, the row sum of the product matrix

$$\sum_j p''_{i,j} = \sum_j \left(\sum_k p_{i,k} p'_{k,j} \right) = \sum_k \sum_j p_{i,k} p'_{k,j} = \sum_k p_{i,k} = 1$$

as required.

If we start with a matrix with a given distribution, what is the limiting distribution for powers of this matrix? For the limiting product of the matrix, what is the probability distribution of the states a , b , and c ? Clearly, we must have a distribution that does not change its probabilities when going to the next day; that is, we must have, for our particular matrix, (5.2-1),

$$(p_a \ p_b \ p_c) \begin{pmatrix} \frac{1}{3} & \frac{1}{3} & \frac{1}{3} \\ \frac{1}{4} & \frac{1}{2} & \frac{1}{4} \\ \frac{1}{4} & \frac{1}{4} & \frac{1}{2} \end{pmatrix} = (p_a \ p_b \ p_c)$$

This is equivalent to the three equations

$$\frac{1}{3}p_a + \frac{1}{4}p_b + \frac{1}{4}p_c = p_a$$

$$\frac{1}{3}p_a + \frac{1}{2}p_b + \frac{1}{4}p_c = p_b$$

$$\frac{1}{3}p_a + \frac{1}{4}p_b + \frac{1}{2}p_c = p_c$$

or

$$-\frac{2}{3}p_a + \frac{1}{4}p_b + \frac{1}{4}p_c = 0$$

$$\frac{1}{3}p_a - \frac{1}{2}p_b + \frac{1}{4}p_c = 0$$

$$\frac{1}{3}p_a + \frac{1}{4}p_b - \frac{1}{2}p_c = 0$$

Only two of these three equations are independent, so we solve the top two with the side condition (the one nonhomogeneous equation)

$$p_a + p_b + p_c = 1$$

After some algebra we get for the *equilibrium solution*

$$p_a = \frac{3}{11}, \quad p_b = p_c = \frac{4}{11} \quad (5.2-3)$$

Thus, regardless of what state distribution we start with, for this transition matrix (5.2-1) we get the same equilibrium state distribution $(\frac{3}{11}, \frac{4}{11}, \frac{4}{11})$. Note carefully that we will be in some definite state—these numbers are probabilities only!

Exercises

5.2-1 Compute the fourth and eighth powers of the transition matrix (5.2-1) and compare the variabilities of all four matrices.

5.2-2 Starting with the square of the transition matrix (5.2-2), compute the equilibrium state distribution.

5.3 Ergodic Markov Processes

The Markov process above is called *ergodic*—from any state we can get to any other state, and in time the system settles down to a limiting distribution independent of the initial state. This does not mean that for the matrix above the weather in time becomes a mixture of fair,

rain, and snow, but only that in the long run *our expectation* settles down to a probability mixture of the states. The long-term probabilities are independent of today's weather.

There are many other types of Markov processes. For example, Figure 5.3-1 shows a nonergodic graph. We see that once we leave state *a*, we will never get back there. Depending on which path we first take, we may go to the single ending state *d*, or else end up going around three states *g*, *h*, and *i*. We do not expect that this kind of graph will be a good model of an information source, so we will set this type aside.




Figure 5-3.1 Nonergodic graph

Even if we do end up in a set of states for which we can get from any one to any other, it does not mean that the probabilities settle down to a single number. Consider, for example, a checkerboard with the transitions being up or down, or left or right, but not diagonally. If we start on a red square, we will find that after an even number of transitions we are on a red square, but for an odd number of transitions we are on a black square. Thus there is an alternation of probabilities for each color of the current square. The theory of all possible Markov processes is therefore messy in principle, but the Markov processes we care about for information systems are ergodic and do not have these peculiarities. In the long run the message symbols settle down to fixed probabilities, although as noted above, not to a mixture of actual symbols but to a probability distribution. It is easy, for example, to understand the words “terminal state” and the words “a set of terminal states.” And it is easy to see that if the terminal set is not all the states of the system, the system is not ergodic in the sense that we can get from any one to any other. Furthermore, we must avoid any alternation of states. For any information source that we may have to encode, we expect to have an ergodic source, although it is not completely impossible that other systems could occur. The complications, if they do arise, are not difficult

to work out. We shall assume in the future that all our systems are ergodic.

Exercise

5.3-1 Given the symbols 0 and 1, draw the one-, two-, and three-memory Markov diagrams.

5.4 Efficient Coding of an Ergodic Markov Process

It should be immediately evident that we can use the Markov structure of the source to improve the encoding of the source. For each state in the Markov system we can use the appropriate Huffman code obtained from the corresponding transition probabilities for leaving that state. Depending on how variable the probabilities are for each state, the encoding will gain a lot or not. However, as we increase the order of the Markov process, we tend to gain less and less while the number of states approaches infinity rapidly.

Let us do this Huffman encoding for the weather matrix of Section 5.2. For state a (today) we get (tomorrow) the binary encoding

$$\begin{array}{lll} a = 1 & p_a = \frac{1}{3} & \\ b = 00 & p_b = \frac{1}{3} & L_a = \frac{5}{3} = 1.6667 \dots \\ c = 01 & p_c = \frac{1}{3} & \end{array}$$

For state b ,

$$\begin{array}{lll} a = 10 & p_a = \frac{1}{4} & \\ b = 0 & p_b = \frac{1}{2} & L_b = \frac{3}{2} = 1.5 \\ c = 11 & p_c = \frac{1}{4} & \end{array}$$

For state c ,

$$\begin{array}{lll} a = 10 & p_a = \frac{1}{4} & \\ b = 11 & p_b = \frac{1}{4} & L_c = \frac{3}{2} = 1.5 \\ c = 0 & p_c = \frac{1}{2} & \end{array}$$

At equilibrium we get [using the equilibrium probabilities (5.2-3)] the average code length

$$L_M = \frac{3}{11}L_a + \frac{4}{11}L_b + \frac{4}{11}L_c = \frac{17}{11} = 1.5454 \dots$$

On the other hand, let us code the average frequencies of the letters a , b , and c . The probabilities are clearly the equilibrium probabilities (5.2-3)

$$p_a = \frac{3}{11} \quad p_b = \frac{4}{11} \quad p_c = \frac{4}{11}$$

The code is

$$\begin{aligned} a &= 01 \\ b &= 00 \\ c &= 1 \end{aligned}$$

We get, therefore,

$$L = \frac{3}{11}(2) + \frac{4}{11}(2) + \frac{4}{11}(1) = \frac{18}{11}$$

Thus the average code length L_M for the individual state encoding is less than the encoding L of the average frequencies

$$L_M = \frac{17}{11} < \frac{18}{11} = L$$

As we expect, the Markov-state system shortens the average code length. The gain would be much greater had the probabilities been less uniform.

Exercise

5.4-1 If the fair weather a has probabilities $p(a | a) = \frac{3}{4}$, $p(b | a) = p(c | a) = \frac{1}{8}$, and the others are the same in this section, show that $L_M = \frac{11}{8} < \frac{12}{8} = L$.

5.5 Extensions of a Markov Process

We have already (Section 4.2) defined an extension of a source. For example, the octal code is the third extension of the binary code. In a sense, an extension is simply a concatenation of the source symbols into longer symbols. We need to examine the extension of a Markov process.

Definition. If S is an m th-order Markov information source with alphabet $s_1, s_2, \dots, s_q$, and if its transition probabilities from one of

the q^m states to the next is given by

$$p(s_j \mid s_{j_1}, s_{j_2}, \dots, s_{j_m}) \quad (5.5-1)$$

then the n th extension of this source is labeled $S^n = T$ with q^n symbols $t_1, t_2, \dots, t_{q^n}$. Each t is a sequence of n symbols from the alphabet of S . The system S^n is a k th-order Markov process where k is the first integer greater than or equal to (m/n) . Since each t_i corresponds to some sequence of n of the s_j , the conditional probabilities of T ,

$$p(t_i \mid t_{i_1}, t_{i_2}, \dots, t_{i_k}) \quad (5.5-2)$$

can be obtained from the original transitional probabilities of S by predicting one s_i at a time while looking back at m earlier symbols. Rewriting (5.5-2), we have for $(kn \geq m)$ and repeated applications of (5.5-1),

$$\begin{aligned} & p(s_{i_1}, s_{i_2}, \dots, s_{i_n} \mid s_{j_1}, s_{j_2}, \dots, s_{j_m}) \\ &= p(s_{i_1} \mid s_{j_1}, s_{j_2}, \dots, s_{j_m}) p(s_{i_2} \mid s_{j_2}, s_{j_3}, \dots, s_{j_m}, s_{i_1}) \\ & \quad \dots p(s_{i_n} \mid s_{i_{n-m}}, s_{i_{n-m+1}}, \dots, s_{i_{n-1}}) \end{aligned}$$

This product assumes $n > m$; if not, the last term would be

$$p(s_{i_n} \mid s_{j_n}, s_{j_{n+1}}, \dots, s_{i_{n-1}})$$

The formalism is more complicated than the idea! We will need the results of this section later.

5.6 Predictive Run Encoding

We have just shown how to use the transition probability structure of the messages to reduce the length of the encoded messages; the simple frequencies of the separate input symbols can be used to design a Huffman code, and the Markov process can handle the larger correlations. Besides these, there are many other possible structures of the input messages that could be used to shorten the encoded message, but they are far too numerous to investigate in detail in this text. Instead, we examine the general case; we *assume* that you have either a program,

or some hard-wired equipment, that can recognize the patterns to be used to predict the next symbol. For simplicity we assume a binary source. Using this program, we will then encode the input stream of binary messages into a corresponding binary stream which has relatively long runs of 0's separated by comparatively few 1's. The better your ability to predict the next symbol, the longer will be the runs of 0's. The stream of 0's and 1's is followed by a second stage of encoding, where the code compression occurs. We will send a stream of numbers, each number giving the length of the corresponding run of 0's, and the receiver will be able to reconstruct the original message using equipment almost exactly like the encoder.


5.7 The Predictive Encoder

Whatever the pattern of the input message that you can recognize and use to predict the next symbol, we assume that you have written a program to do this. You are merely required to predict the next symbol in the input stream. Your predictor can be a fixed one, or it can be "adaptive" in the sense that it can gather statistics about how well it does, and from them make modifications to itself. It can be linear or nonlinear—nothing is being said about how you do the predicting.


It is necessary to emphasize the generality of the program you are supposed to supply for your particular application. The method we are going to discuss makes the single assumption that the program predicts the next input symbol correctly most of the time. The better the predictions are, the more the input message is finally compressed.

As shown in Figure 5.7-1, there is a source S of input symbols s_i . In part (a) there is a predicting program in the rectangular box labeled P which predicts from the stream of symbols s_i (for $i < n$) the predicted symbol p_n . Assuming that both input and output are using the binary system of 0's and 1's, we *logically add* the two current symbols, input s_n and predicted p_n , to get the error e_n of the n th prediction. The output symbol e_n is a 0 whenever the prediction p_n is the same as the input symbol s_n , and a 1 when the prediction is wrong. The predictor P also gets the error messages e_n and is therefore in a position to correct itself and to know exactly the past message stream up to the next prediction.


As noted, if the predictor is fairly good, there will be a stream of 0's with an occasional 1 here and there. A random predictor would have half 0's and half 1's. Even a poor predictor would have many more 0's than 1's.



(a)



(b)



(c)

Figure 5-7.1 Predictive encoding

5.8 The Decoder

At the receiving end (Figure 5.7-1b) there is a *decoder*, which is essentially the same as the predictor. Assuming that the box labeled P is identically the same as the box in the encoder, then if the two start out

in the same state, it follows that each will make the same predictions. When a 1 comes to the decoder in Figure 5.7-1c, it logically adds this error to its prediction and thus the correct output emerges. Again, the error message e_n is supplied to the predictor. The slight difference between the form of the information at the two ends requires only slight modification to the equipment; thus the two ends of the signaling system are essentially the same.

We need to model the situation in mathematical symbols so that we can evaluate how good it is. Let p be the probability of the predictor P being right for each symbol. (Note the reversal of the role of p ; it is now the probability of being right.) We assume that the probability is a constant since otherwise there is a known structure that you should have incorporated into the predictor. For the same reason we *assume* that each prediction error is independent of the others. Your formula may, of course, be recursive; it is the errors we are assuming are independent. Let $1 - p = q$ be the probability of guessing wrong. Naturally, we assume that $p \geq \frac{1}{2}$, since otherwise we would simply reverse the symbols p and q by complementing the predicted 0's and 1's and have a suitable p .

These probabilities are not to be confused with the errors in the equipment; they are associated with the predictor P . Because of the foregoing assumptions about the predictor, the output stream of 0's and 1's has a fixed probability of error and adjacent positions are independent. The prediction errors are "white noise" (Section 2.4). A large p means a good encoder; a p near $\frac{1}{2}$ means that you are doing hardly better than guessing.

5.9 Run Lengths

The probability $p(n)$ of a run of exactly n 0's is (since each run of 0's ends with a 1)

$$p(n) = p^n q \quad (n = 0, 1, \dots)$$

Summed over all lengths n , we have

$$\sum_{n=0}^{\infty} p^n q = q \sum_{n=0}^{\infty} p^n = \frac{q}{1-p} = \frac{q}{q} = 1$$

as it should be.

Suppose that for the moment we decide to send a k -digit binary number to represent a run length. For a reason we will soon see we can do this only for run lengths 0 to $2^k - 2$. At the receiving end the

receiver copies down the indicated number of 0's and follows this by a 1. If the sent number is 0, we simply copy down a 1. For the binary number $(2^k - 1)$ the receiver copies down $(2^k - 1)$ 0s and does *not* copy down a 1. If the next number sent is 0, then of course a 1 is copied, and for any other number that number of 0's is copied followed by the 1. Thus we can send any length of a run of 0's we need by sending a sufficient number of k -digit binary numbers, $k > 1$.

We shall not attempt to use a Huffman coding system. In the original case we have a potentially infinite input alphabet, since in principle any run length could occur; after "blocking it" into runs at most k digits, we could, if we wished, use Huffman encoding.

Returning to the proposed method, what is the *expected* number of digits sent? To find out, we need to sum the probabilities of occurrence times the number sent in each case. Thus we construct the table

Length of Run	Number of Digits Sent
$0 \leq n \leq (2^k - 1) - 1$	k
$(2^k - 1) \leq n \leq 2(2^k - 1) - 1$	$2k$
$2(2^k - 1) \leq n \leq 3(2^k - 1) - 1$	$3k$
etc.	

The expected length is [setting $2^k - 1 = m$ and $p(n)$ as the probability of a run of exactly n 0's]

$$\begin{aligned}
 & [p(0) + p(1) + \dots + p(m-1)]k \\
 & + [p(m) + p(m+1) + \dots + p(2m-1)]2k \\
 & + [p(2m) + p(2m+1) + \dots + p(3m-1)]3k \\
 & + \dots
 \end{aligned}$$

Since

$$p(n) = p^n q$$

we get

$$\begin{aligned}
 & q[1 + p + \dots + p^{m-1}]k \\
 & + qp^m[1 + p + \dots + p^{m-1}]2k \\
 & + qp^{2m}[1 + p + \dots + p^{m-1}]3k
 \end{aligned}$$

Summing the terms in brackets, we arrange the expression in the form.

$$q \left[\frac{1 - p^m}{q} \right] [1 + 2p^m + 3p^{2m} \dots]k$$

This sum is, in turn,

$$q \left[\frac{1 - p^m}{q} \right] \left[\frac{1}{(1 - p^m)^2} \right] k = \frac{k}{1 - p^m}$$

We need a short table (Table 5.9-1) of the function ($k > 1$)

$$\frac{k}{1 - p^{(2^k - 1)}}$$

so that for a given $p \geq \frac{1}{2}$ we can pick the best k (Table 5.9-2).

TABLE 5.9-1 *Table of* $\frac{k}{1 - p^{(2^k - 1)}}$

k	p							
	0.5	0.6	0.7	0.8	0.9	0.95	0.98	0.99
2	2.29	2.55	3.04	4.10	7.38	14.0	50.0	67.3
3			3.27	3.80	5.75	9.94	34.0	44.2
4				4.15	5.04	7.45	22.7	28.6
5					5.20	6.28	15.3	18.7
6						6.25	10.7	12.8
7						7.01	8.33	9.71
8							7.58	8.67
9							8.05	9.05
10								

TABLE 5.9-2 *Table of Efficiency of Encoding*

p	Optimum k	Expected Number of Digits	Expected Length of Runs	Compression Ratio
0.7	2	3.04	3.33	0.91
0.8	3	3.80	5.00	0.76
0.9	4	5.04	10.00	0.50
0.95	6	6.25	20.00	0.31
0.98	7	7.58	50.00	0.16
0.99	8	8.67	100.00	0.087

The number of encoded bits in the string of 0's and 1's for a run of length is $n + 1$; hence the expected length of the encoded runs is

$$\sum_{n=0}^{\infty} (n + 1)p^n q = q \sum_{n=0}^{\infty} (n + 1)p^n = q \frac{1}{(1 - p)^2} = \frac{1}{q}$$

Suppose that we compute the variance. We get

$$\begin{aligned} \sum_{n=0}^{\infty} \left(n + 1 - \frac{1}{q} \right)^2 p^n q &= q \sum (n + 1)^2 p^n - \frac{1}{q^2} \\ &= q \left[\frac{1 + p}{q^3} \right] - \frac{1}{q^2} = \frac{p}{q^2} \end{aligned}$$

From these tables we see that the efficiency of the run encoding depends on the choice of k , and for optimal design this in turn depends on p . Evidently, it is easy to select an appropriate k when p is given.

Exercise

5.9-1 Discuss the use of Huffman encoding of the run of zeros.

5.10 Summary of Predictive Encoding

Because of the generality of this scheme of message compression it is necessary to emphasize both that it does not do everything, and that its efficiency depends completely on the predictor that is supplied. The better the predictor, the longer the runs of 0's and hence the compression at the second stage of encoding. The generality of the predictor is unlimited in theory, although in practice it must be simple enough so that it can keep up with the stream of the input alphabet. All it has to do is predict the next symbol in the input stream. How it does its job is not a part of the theory. If the encoding is to run in real time, then there are buffer-size problems, but on a computer problem this is not a vital point.

5.11 What Is Hashing?

Hashing is a method of reducing the redundancy of an ensemble of m messages that have been selected from a much larger ensemble of M possible messages. It is, however, a reduction in a particular sense only.

To make matters specific, consider the ensemble of names of people, the names of variables in a FORTRAN program, or any other rich source of names. In any particular application, payroll, FORTRAN compilation, or other processing of the names, we expect to have only a very small sample from the ensemble of possible names. Suppose, to be definite, that we expect to have 500 names. We will assign 10 bits ($2^{10} = 1024$, which is more than twice 500). For each name we will take the first 10 bits, do a logical addition of the next 10 bits, the next 10, and continue to the end of the name (padding out the last 10 with 0's if necessary). This process generates a unique "hash" sum from a given name, but of course two different names could generate the same hash sum.

Given this original name, we place in a hash table at the computed 10-bit hash number location (perhaps with a constant added) a pointer to the original name. As a result we have a table of pointers back to the original names, but the table allows the direct location of an item by simply going, via the hash table, to the calculated location rather than by a search through the table of names to find the required name.

What is the probability of two different names leading to the same hash sum? If the hashed names were random, and that is what we are trying to do by the logical addition of each 10 bits, there are approximately $\frac{1}{2}$ times 500^2 possible *collisions*. Like the famous birthday problem where 23 different people gives a probability of more than half of at least one duplicate birth date, the probability of at least one collision is very high. Yet the probability of any one name having a collision with any other name is less than $\frac{1}{2}$, since we have more than twice as many name locations in the table as there are names. Thus the typical name when hashed (encoded) into the 10 bits has less than half a chance of having a collision with any other name; still some collision in the full 500 names is highly probable. We must, therefore, allow for them.

5.12 Handling Collisions

How shall we handle collisions? A simple method is to go to the calculated (hashed) place and if this is empty, we make the entry; but if it has an occupant, we look down the table and insert the new entry in the first empty location that follows. This means that when we look for an entry at a later time we have to look not only at the computed place, but if the pointer there does not point back to the correct name, we must continue looking at the consecutive nonempty places and checking if their pointers give the name until we come to either a match or else a blank. In the latter case we know that there is no entry corre-

sponding to the given name. If we want to make a new entry, as we would in a FORTRAN compilation, we enter it in the empty location. In this way we can construct a table of hashed names.

Of course, if we knew all the names in advance, we could find a particular encoding that would avoid collisions. But if we either have to handle a growing table of names, or else do not want to spend the time to find a special encoding, this approach, through what can be called “random encoding,” is useful.

5.13 Removals from the Table

In some problems it may occasionally be necessary to remove a name from the file of names. This would not happen in a FORTRAN compilation, but it does happen in some applications. How can we do this? Evidently, we compute the address of the name’s pointer and if it is there (or in some following location), we remove it. If there is an occupied cell following the removed one, we must look to see if when it was entered it had been moved down, and if so we must move it back up. Indeed, we must examine *all* the successive entries until we come to a blank, to see how each should be placed. This is simplest done by constructing a pseudofile of names from the following adjacent pointers, and simply reentering them. In this fashion we can keep a file of computed names up to date.

5.14 Summary of Hashing

We see that given a collection of names that may individually be represented by very long (or possibly only occasionally very long) strings of characters, we have compressed them into a length that is slightly longer than minimum. The use of more than minimal redundancy decreases the frequency of collisions and the attendant waste of computer time, and it is an engineering decision as to how many more than the minimal number of bits to use. In our example we used one extra bit and had half the cells of our table empty. The use of two extra bits would greatly decrease the number of collisions, but would double the space of the table in storage.

We see that the random encoding is based on the number of entries expected but not on their detailed structure, and this allows the use of a greatly compressed table which can be used in a direct lookup fashion rather than a searching fashion. We have substituted computing power (the hashing algorithm) for searching, since we go directly to the com-

puted address, and with at most some local fiddling around we find the entry, if the entry exists.

This random encoding is based on the idea that the hashing algorithm we used gives random encodings. Certainly, two names that differ by a single letter will still be separated. It will only be by chance that two fairly different names will collide. The hashing algorithm we used is a very simple one, and there is an extensive theory for how to design a hashing algorithm for various kinds of data (Ref. [K]). The one we used is clearly fairly effective, but probably could be improved upon. We will not go into this point further. We are exploiting the fact that random encoding produces a degree of uniformity from an unknown but structured source—it tends to spread out bunchings.

The main results we want to emphasize are two. First, this is a useful method of encoding data for some purposes; it reduces the redundancy without actually having the redundant structure of the messages (names). Second, this is an example of random encoding, a concept we will need in Chapter 10.

5.15 Purpose of the Gray Code

Another important code is the Gray code. It is designed to meet a common need which is best described in terms of a particular application, although it has wider uses. Suppose that we have a rotating analog wheel (see Figure 5.15-1) and have a number of brushes rubbing on the




Figure 5-15.1 Eight-sector binary wheel

wheel (or lights shining), each on its own track. In order to record how far the wheel has rotated, the tracks are divided into sectors. For the typical binary system the number of sectors would be a power of 2, thus giving rise to a binary number of n digits. There are, of course, n tracks, one for each digit (the figure has three tracks).

If we use direct binary coding of the sectors, then when the wheel moved from, say the binary number 011 to the next binary number 100, each of the three brushes would find that its value changed. Suppose that we tried to read at the moment of changing sectors. We could expect any number to emerge from this attempt to read.

The Gray code is a method of encoding for which successive numbers change by only one position all around the wheel. Thus the attempt to read when the brushes are between two sectors can give only one of these two numbers, and no large error will occur.

5.16 Details of a Gray Code

We have just effectively defined a Gray code, but the definition does not give a unique code. In practice we use one simple version, usually known as the *Gray code*.

The definition of the Gray code is made inductively. For a one-digit code we have the two states, labeled as 0 and 1 (see Figure 5.16-1). Notice that in going around the circle we also have the transition from the 1 to the 0 that occurs in going from the bottom of the column to the top. Imagine these as a column. Below these two we write the symbols in reverse order, and attach to the first two a leading digit 0 and to the last two a leading digit 1. Examination shows that as we pass from the original set to the new set, we change only the leading digit from a 0 to a 1, the other digits being a copy. All the rest of the new symbols also change by one place only because they are the same




Figure 5-16.1 Gray codes

as the original except in reverse order. The transition from the last to the first is a change in only the first (attached) digit. In the figure we show the process of going from one to two digits, from two to three digits, and part of going from three to four digits.

Exercise

5.16-1 Show that for a four-binary-digit code, there can be different codes with the basic property of changing only one digit in going to the next sector.

5.17 Decoding the Gray Code

The numbers that come from the coding wheel of an analog device are typically in the Gray code, and we usually need them in the binary code if we are to process them in a computer. Reading from left to right we take all the 0's as being correct and also take the first 1. But the next position should be complemented. If this position is a zero, we continue complementing, but if it is a 1, we complement it and stop complementing until we pass another 1. Thus the number of earlier 1's we met controls whether we complement the next digit or not. We complement on an odd number of earlier 1's and do not complement on an even number of 1's. This simple rule, together with the general simplicity, makes the particular code the favorite solution to this problem of having only one digit change between successive numbers. The state diagram is easy to prepare.

Exercises

5.17-1 Make a state diagram for decoding the Gray code.

5.17-2 Describe the rule for going from binary to Gray (for simulation modeling).

5.18 Anti-Gray Code

One day when I was reading a journal, I noticed a paper which observed that when testing a memory device it *might* be a good idea to arrange the memory addresses such that the switching tree has as many changes as possible. Shortly after this point I set the paper down and began to

think about just what the problem was (since I noticed that the paper was rather heavy going).

It is evident that in going from one address to the next, the 1's complement would have the most changes. Next, it is evident that two successive complements would get back to the same address, so the best that could be done is that one stage of complementing be followed by complementing all but one position. This, in turn, means that taking every other address would form a Gray code. On the average of the k possible bits you have $k - \frac{1}{2}$ bits changing in the switching tree. This is a good example of the rich variety of useful codes that are possible, and that once you are reasonably familiar with a wide variety of codes, it is easy to develop others to meet your specific needs.

To fill in the details, if you have a k bits for the address of the memory you are testing, you first construct a Gray code on $k - 1$ bits and append a leading 0. Then you interpolate the complement of each address immediately after the address. No new address can be an old one because all the old ones begin with a 0 and all the complement addresses begin with a 1. No two complement addresses can be the same because then two of the original addresses were the same. Thus you will have 2^k distinct addresses in the final code.

For example, if you want a code for a 4-bit address, you begin with the 3-bit Gray code with leading 0's. You then get

Original Gray		New Code
0000	→	0000
0001	→	1111
0011	→	0001
0010	→	1110
0110	→	0011
0111	→	1100
0101	→	0010
0000	→	1101
		0110
		1001
		0111
		1000
		0101
		1010
		0100
		1011

An alternate idea is simply to reverse the order of the address bits so that adjacent numbers are placed far from each other. This can be useful if you have different storage banks in a very high speed computer.

5.19 Delta Modulation

Another method of encoding is to send the difference between the current signal and the previous signal. For slowly changing signals this

means that typically, you can send fewer digits. When changes occur that are more rapid than the method can keep up with, the transmitted signal falls behind, but with suitable terminal equipment this loss will be made up in time (Figure 5.19-1).




Figure 5-19.1

Thus you send

$$\Delta y_n = y_n - y_{n-1}$$

and use fewer digits in Δy than in y_n . The method is widely used with suitable modification to combat noise.

5.20 Other Codes

A topic of some interest is that of resynchronizing a decoding system if a single symbol is either lost or is added to the stream of digits coming into the receiver. In a Huffman code we could be in some trouble, but usually, the receiver will resynchronize fairly soon. But how long it will take on the average is not a simple question, although it has been investigated in some studies. Evidently, in any particular case a Monte Carlo simulation of a code with errors can be run, and from the simulation suitable conclusions can be made. Within a computer these synchronization errors are not likely to occur—only a changed digit is likely.

One way of avoiding resynchronization problems is to use a parallel channel. Another way, when using serial transmission, is to embed the message in a longer block with synchronization pulses at both ends. This, of course, wastes capacity, but it does catch spurious pulses.

With the current low cost of accurate crystal clocks it is common practice to use clock pulses either from a single clock, or by keeping several clocks in synchronization, so that both the sender and the receiver have the same timing pulses.

The topic is not of enough general interest to do more than point out the problem and the obvious solutions. Many other methods of code compression can be found in Ref. [DG].

Entropy and Shannon's First Theorem

6.1 Introduction

Up to this point we have been concerned with coding theory; now we begin information theory. Coding theory answers the questions of (1) how to design codes for white noise (mainly), and (2) how to compress the message when the probabilities (structure) of the messages are known.

We now need a general method for measuring the structure of the source. To do this we introduce the concept of *entropy*. While entropy was used long ago in many physical situations, and the concept in information theory bears a strong resemblance to the classical definition, we shall study entropy on its own merits in this particular field, and not get involved in the various suggestive analogies. For us, entropy is simply a function of a probability distribution p_i . For some analogies, see Refs. [Ga], [Gu], and [ME].

Information theory combines noise protection and efficient use of the channel into a single theory. However, the simple model of channel noise (white noise) is sometimes unrealistic, and we will occasionally treat more general patterns of errors. This leads to the important concept of *channel capacity*, which is taken up in Chapter 8.

The result we are headed for is Shannon's main theorem in Chapter 10, which gives a relationship between the channel capacity C (to be defined in Chapter 8) and the maximum rate of signaling possible. We will prove the remarkable theorem that we can come arbitrarily close to the maximum rate of signaling and also achieve an arbitrarily low rate of error. In a sense this achieves both *Huffman-type* compression

and *Hamming-type* noise protection in the same code. Unfortunately, the proof is not constructive, so that while information theory sets the bounds on what can be done, it does not tell us how to achieve them. Yet, as we said in Section 1.2, the theory is very useful.

In this chapter we prove a very special case of the main theorem, namely the *noiseless coding theorem*, where we ignore the problem of noise. When this theorem and its proof are understood, the proof of the main theorem will appear a bit clearer—once a lot of preliminary mathematical results are covered in Chapter 9.

6.2 Information

Suppose that we have the source alphabet of q symbols $s_1, s_2, \dots, s_q$, each with its probability $p(s_1) = p_1, p(s_2) = p_2, \dots, p(s_q) = p_q$. When we receive one of these symbols, how much information do we get? For example, if $p_1 = 1$ (and, of course, all the other $p_i = 0$), then there is no “surprise,” no information, since you know what the message must be. On the other hand, if the probabilities are all very different, then when a symbol with a low probability arrives, you feel more surprised, get more information, than when a symbol with a higher probability arrives. Thus information is somewhat inversely related to the probability of occurrence.

We wish to construct this function $I(p)$, which measures the amount of information—surprise, uncertainty—in the occurrence of an event of probability p . We assume three things about $I(p)$.

1. $I(p) \geq 0$ (a real nonnegative measure).
2. $I(p_1 p_2) = I(p_1) + I(p_2)$ for independent events (additive).
3. $I(p)$ is a continuous function of p .

The second of these conditions is known as the *Cauchy functional equation* for the function $I(p)$, meaning that it serves to define $I(p)$. If p_1 and p_2 are both the same number p , not necessarily the same event, then

$$I(p^2) = I(p) + I(p) = 2I(p)$$

Now if $p_1 = p$ and $p_2 = p^2$, then we have

$$I(p^3) = I(p) + I(p^2) = 3I(p)$$

and in general we have

$$I(p^n) = nI(p)$$

We recognize that the standard law of exponents for positive integers applies to the function $I(p)$. Following this clue we adapt for our needs the usual exponent extension to fractional values. We set

$$p^n = y, \quad p = y^{1/n}$$

and hence

$$I(y) = nI(y^{1/n})$$

or, after some further manipulation,

$$I(y^{m/n}) = \frac{m}{n} I(y)$$

Thus for the rational numbers the function of $I(p)$ obeys the same formula as the log function. Assumption 3 of continuity allows us to extend this to all numbers ($0 \leq p \leq 1$), rational or irrational. Thus we have

$$I(p) = k \log p$$

for some constant k and some base of the log system. From the first assumption it is natural to pick the constant k as -1 , and we have, finally,

$$I(p) = -\log p = \log \frac{1}{p}$$

for some base of the log system.

Let us examine and illustrate this *additive* property 2 that we assumed, since it plays so central a role in deriving the measure of surprise. Consider the simultaneous and independent toss of a coin and roll of a die. We *feel* that each outcome (coin and die) has no influence on the outcome of the other. The coin has two equally likely outcomes and the die has 6. The product space of the two independent trials has 12 equally likely outcomes. If c_i is the outcome of the coin and d_j is

the outcome of the die, then the derived formula of $I(p)$ is

$$I(c_i d_j) = I(c_i) + I(d_j)$$

$$\log \frac{1}{12} = \log \frac{1}{2} + \log \frac{1}{6}$$

which is, of course, true.

The words “uncertainty,” “surprise,” and “information” are related. Before the event (experiment, reception of a message symbol, etc.) there is the amount of uncertainty; when the event happens there is the amount of surprise; and after the event there is the gain in the amount of information. All these amounts are the same.

This is an engineering definition based on probabilities and is not a definition based on the meaning of the symbols to the human receiver. The confusion at this point has been very great for outsiders who glance at information theory; they fail to grasp that this is a highly technical definition that captures *only part* of the richness of the usual idea of information.

What base of the log system shall we use? It is simply a matter of convention since any set of logs is proportional to any other set. This follows directly from the fundamental relationship for logs:

$$\log_a x = \frac{\log_b x}{\log_b a} = (\log_a b) \log_b x$$

It is convenient to use the base 2 logs; the resulting unit of information is called a *bit* (binary digit). If we use base e , as we must whenever we get into the calculus, then the unit of information is called a *nat*. Finally, sometimes the base 10 is used and the unit is called a *Hartley*, after R. V. L. Hartley, who first proposed the use of the logarithmic measure of information.

We are using the word “bit” in two different ways, both as a digit in the number base 2 and as a unit of information. They are not the same, and we shall be careful to say “bit of information” when we mean that definition and there could be confusion.

It is easy to convert logs from one base to another using

$$\begin{array}{ll} \log_{10} 2 = 0.30103 \dots & \log_2 10 = 3.32193 \dots \\ \log_{10} e = 0.43429 \dots & \log_e 10 = 2.30259 \dots \\ \log_e 2 = 0.69315 \dots & \log_2 e = 1.44270 \dots \end{array}$$

We have just found one solution to the Cauchy functional equation, namely $I(p) = \log_2 (1/p)$. Is this solution unique, to within the base of the log system? Suppose that there were another solution to the derived functional equation

$$I(p^n) = nI(p)$$

Call it $g(p)$. Hence we have the equation

$$g(p^n) = ng(p)$$

Now take the difference of the two solutions and allow for the constant of proportionality C , which depends on the base

$$g(p^n) - C \log_2 \frac{1}{p^n} = n \left[g(p) - C \log_2 \frac{1}{p} \right]$$

Next, pick C as $(p_0 \neq 0, 1)$

$$C = \frac{g(p_0)}{\log_2(1/p_0)}$$

This choice makes the right-hand side of the equation above equal to zero at p_0 . Now for any number z there is an n such that (assuming p_0 neither 0 nor 1)

$$z = p_0^n$$

Hence the left-hand side of the equation becomes

$$g(z) = C \log_2 \frac{1}{z}$$

and the solution is essentially unique. The continuity assumption covers the two values of p that were omitted, namely 0 and 1.

6.3 Entropy

If we get $I(s_i)$ units of information when we receive the symbol s_i , how much do we get on the average? The answer is that since p_i is the probability of getting the information $I(s_i)$, then on the average we get

for each symbol s_i ,

$$p_i I(s_i) = p_i \log_2 \frac{1}{p_i}$$

From this it follows that *on the average*, over the whole alphabet of symbols s_i , we will get

$$\sum_{i=1}^q p_i \log_2 \frac{1}{p_i}$$

Following custom, we label this important quantity (for radix r)

$$H_r(S) = \sum_{i=1}^q p_i \log_r \left(\frac{1}{p_i} \right) \quad (6.3-1)$$

and call it the *entropy* of the signaling system S having symbols s_i and probabilities p_i . Of course,

$$H_r(S) = H_2(S) \log_r 2$$

This is the entropy function for a distribution *when all that is considered are the probabilities p_i of the symbols s_i* . In Section 6.10 we will consider the entropy of a Markov process.

Corresponding to each distribution $P = (p_1, p_2, \dots, p_q)$ of symbols s_i , there is a single number called the entropy and labeled $H(S)$. This is analogous to the usual idea of an average of a distribution—the average is a single number which, in some sense, summarizes the distribution. The entropy $H(S)$ is the weighted average of the logs of the reciprocals of the probabilities of the distribution. The entropy is a single measure of a distribution; it is the average information of the alphabet S .

As an alternative approach to the entropy function, consider that in a long message of N symbols (from the alphabet S) you expect Np_1 of the first symbol; Np_2 of the second, and so on. The probability P of a message of length N with these numbers of symbols is exactly

$$\begin{aligned} P &= p_1^{(Np_1)} p_2^{(Np_2)} \dots p_q^{(Np_q)} \\ &= [p_1^{(p_1)} p_2^{(p_2)} \dots p_q^{(p_q)}]^N \end{aligned}$$

Hence the information is

$$\log \left(\frac{1}{P} \right) = N \sum_{i=1}^q p_i \log \left(\frac{1}{p_i} \right)$$

and the information *per symbol* of the alphabet S is

$$H(S) = \sum \left(p_i \log \frac{1}{p_i} \right)$$

which is again the entropy (as it should be).

It is important to realize that a remark like “Consider the entropy of the source” can have no meaning unless a model of the source is included. Using random numbers as an example, you have formulas for generating *pseudorandom* numbers. Suppose that we had a table of such numbers. If you do not recognize that they are pseudorandom numbers, then you would probably compute the entropy based on the frequencies of occurrence of the individual numbers. Since pseudorandom-number generators do a good job of simulating random numbers, you would also find that each new number came as a complete surprise. But if you knew the structure of the formula used to generate the table, you would, after a few numbers, be able to predict perfectly the next number—there would be no surprise. Your estimate of the entropy of a source of symbols therefore depends on the model you adopt of the structure of the symbols.

The entropy function (6.3-1) involves only the distribution of the probabilities—it is a function of a probability distribution p_i and does not involve the s_i . For example, if the probabilities are 0.4, 0.3, 0.2, and 0.1, then from the table of Appendix B, column 3, we have

p	$p \log \frac{1}{p}$
0.4	0.52877
0.3	0.52109
0.2	0.46439
0.1	0.33219
Sum	1.84644

The entropy of this distribution is therefore 1.84644 (approximately). Although we should write the entropy function $H(S)$ as a function of




Figure 6-3.1

P , we will continue to refer to the alphabet S and write $H(S)$, although occasionally we will write $H(P)$, when there are only two events, with probabilities p and $(1 - p)$.

The function of $p \log_2(1/p)$ is graphed in Figure 6.3-1. From

$$\frac{d}{dp} \left(p \log_2 \frac{1}{p} \right) = \frac{d}{dp} \left(p \log_e \frac{1}{p} \right) \log_2 e = \log_2 \left(\frac{1}{p} \right) - \log_2 e$$

We see both the infinite slope at $p = 0$ and that the maximum occurs at $p = 1/e$, where the derivative is zero.

We also need the property (since it is equated to zero, we can ignore the factor $\log_e 2$)

$$\lim_{x \rightarrow 0} (x \log_e x) = 0$$

To prove this, we write it as

$$\lim_{x \rightarrow 0} \left(\frac{\log_e x}{1/x} \right)$$

and apply l'Hôpital's rule by differentiating both numerator and denominator separately:

$$\lim_{x \rightarrow 0} \left(\frac{1/x}{-1/x^2} \right) = \lim_{x \rightarrow 0} (-x) = 0$$

As another example of the entropy of a distribution, consider the source alphabet $S = \{s_1, s_2, s_3, s_4\}$, where $p_1 = \frac{1}{2}$, $p_2 = \frac{1}{4}$, $p_3 = \frac{1}{8}$, and $p_4 = \frac{1}{8}$. With $r = 2$ we have the entropy

$$\begin{aligned} H_2(S) &= \frac{1}{2} \log_2 2 + \frac{1}{4} \log_2 4 + \frac{1}{8} \log_2 8 + \frac{1}{8} \log_2 8 \\ &= \left(\frac{1}{2}\right)1 + \left(\frac{1}{4}\right)2 + \left(\frac{1}{8}\right)3 + \left(\frac{1}{8}\right)3 \\ &= 1\frac{3}{4} \text{ bits of information} \end{aligned}$$

As one application of the idea of entropy, consider the toss of a coin when both sides are considered to be equally likely.

$$\begin{aligned} I(s_i) &= \log_2 \left(\frac{1}{2}\right)^{-1} = \log_2 2 = 1 \\ H_2(S) &= \left(\frac{1}{2}\right)I(s_1) + \left(\frac{1}{2}\right)I(s_2) = 1 \end{aligned}$$

The distribution consisting of just two events is very common. If p is the probability of the first symbol (event), then the entropy function is


$$H_2(P) = p \log_2 \left(\frac{1}{p}\right) + (1 - p) \log_2 \left(\frac{1}{1 - p}\right)$$

and is tabulated in the last column of the table in Appendix B. The graph of this function is given in Figure 6.3-2. Note that at $p = 0$ and $p = 1$ it has a vertical tangent, since

$$\begin{aligned} &\frac{d}{dp} \left[p \log_2 \left(\frac{1}{p}\right) + (1 - p) \log_2 \left(\frac{1}{1 - p}\right) \right] \\ &= \left[\log_e \left(\frac{1}{p}\right) - 1 - \log_e \left(\frac{1}{1 - p}\right) + 1 \right] \log_2 e \\ &= \log_2 \left(\frac{1}{p}\right) - \log_2 \left(\frac{1}{1 - p}\right) \end{aligned}$$

An application similar to the coin is the roll of a well-balanced die (singular of dice). We have

$$I(s_i) = \log_2 \left(\frac{1}{6}\right)^{-1} = \log_2 6$$

Figure 6-3.2 Entropy function for two probabilities ▷ 104

and the entropy is

$$\begin{aligned} H(S) &= 6\left[\frac{1}{6}I(s_i)\right] = \log_2 6 \\ &= 2.5849 \dots \text{bits of information} \end{aligned}$$

As we see from this example, whenever all the probabilities are equal, the average over the alphabet is the same as the information for any one event.

A final example of the computation of the entropy of a distribution is the entropy of Zipf's law (Section 4.9), where the frequency of the k th item is proportional to $1/k$. Thus for N items, we have the probabilities

$$p_k = \frac{1/k}{S_N} \quad (k = 1, 2, \dots, N)$$

where

$$S_N = \sum_{k=1}^N \frac{1}{k}$$

The entropy for the N items is

$$\begin{aligned}
 H_2(N) &= - \sum_{k=1}^N p_k \log_2 p_k \\
 &= - \sum_{k=1}^N \frac{1/k}{S_N} \left[\log_2 \left(\frac{1}{k} \right) - \log_2 S_N \right] \\
 &= - \frac{1}{S_N} \left[\sum_{k=1}^N \frac{1}{k} \left(\log_2 \frac{1}{k} - \log_2 S_N \right) \right] \\
 &= \frac{1}{S_N} \left(\sum_{k=1}^N \frac{\log_2 k}{k} + \log_2 S_N \sum_{k=1}^N \frac{1}{k} \right) \\
 &= \frac{T_N}{S_N} + \log_2 S_N
 \end{aligned}$$

where

$$T_N = \sum_{k=1}^N \frac{\log_2 k}{k}$$

Using the recursive definitions

$$S_1 = 1$$

$$T_1 = 0$$

$$S_N = S_{N-1} + \frac{1}{N}$$

$$T_N = T_{N-1} + \frac{\log_2 N}{N}$$

then

$$H_2(N) = \frac{T_N}{S_N} + \log_2 S_N$$

is easily computed.

We have Table 6.3-1 (corresponding to Table 4.9-1).

The entropy function of a distribution summarizes one aspect of a distribution much as the *average* in statistics summarizes a distribution. The entropy has properties of both the arithmetic mean (the average)

TABLE 6.3-1 Entropy of Zipf's Law

N	Entropy	Huffman $L_H(N)$
2	0.918	1.00
4	1.792	1.80
8	2.618	2.68
16	3.403	3.43
32	4.149	4.17
64	4.864	4.89
128	5.553	5.60
256	6.222	6.26
512	6.873	6.90
1024	7.511	7.54

and the geometric mean. We have

$$\begin{aligned}
 H(S) &= \sum p_i \log_2 \left(\frac{1}{p_i} \right) \\
 &= \sum \log_2 \left(\frac{1}{p_i} \right)^{p_i} \\
 &= \log_2 \prod_{i=1}^q \left(\frac{1}{p_i} \right)^{p_i}
 \end{aligned}$$

which is a weighted geometric mean.

Exercises

6.3-1 How many bits of information do we get from one draw of a card from a deck of 52 cards?

6.3-2 For $p_1 = \frac{1}{3}$, $p_2 = \frac{1}{4}$, $p_3 = \frac{1}{6}$, $p_4 = \frac{1}{9}$, $p_5 = \frac{1}{12}$, and $p_6 = \frac{1}{18}$, compute the entropy. Ans. $\frac{8}{9} + \frac{11}{12} \log 3 \approx 2.34177$

6.4 Mathematical Properties of the Entropy Function

The entropy function measures the average amount of uncertainty, surprise, or information that we get from the outcome of some situation, say the reception of a message or the outcome of some experiment. It is therefore an important function of the probabilities of the individual events that can occur. Thus in designing an experiment we usually wish




Figure 6-4.1 Bound on $\log_e x$ function

to maximize the amount of information we expect to get; we wish to maximize the entropy function. To do this we need to at least partially control the probabilities of the individual outcomes of the experiment; we need to “design the experiment” suitably. The maximum entropy is increasingly used as a criterion in many situations. Therefore, the entropy function is worth some study for its own sake, independent of the particular applications we will make of it.

The entropy function has a number of mathematical properties that are very useful, and we examine them before getting deeper into the theory. A first property of the $\log_e x$ function can be seen from Figure 6-4-1.

Fitting the tangent line at the point (1, 0), we find that the slope

is

$$\left. \frac{d(\log_e x)}{dx} \right|_{x=1} = 1$$

so that the tangent line is

$$y - 0 = 1(x - 1)$$

or

$$y = x - 1$$

Thus we have for all x greater than zero the useful inequality

$$\log_e x \leq x - 1 \quad (6.4-1)$$

or

$$\log_2 x \leq (x - 1) \log_2 e$$

The equality holds *only* at the point $x = 1$.

The second result we need is the fundamental relationship between two probability distributions. Let the first probability distribution be x_i with of course $\sum x_i = 1$, and the second probability distribution be y_i with correspondingly $\sum y_i = 1$. Consider, now, the expression involving both distributions

$$\sum_{i=1}^q x_i \log_2 \left(\frac{y_i}{x_i} \right) = \frac{1}{\log_e 2} \sum_{i=1}^q x_i \log_e \left(\frac{y_i}{x_i} \right)$$

Using the previous relation (6.4-1), we get

$$\begin{aligned} \frac{1}{\log_e 2} \sum_i x_i \log_e \frac{y_i}{x_i} &\leq \frac{1}{\log_e 2} \sum_i x_i \left(\frac{y_i}{x_i} - 1 \right) \\ &\leq \frac{1}{\log_e 2} \sum_i (y_i - x_i) \\ &\leq \frac{1}{\log_e 2} \left(\sum_i y_i - \sum_i x_i \right) = 0 \end{aligned}$$

Converting back to logs base 2, we have the *fundamental Gibbs inequality*

$$\sum_{i=1}^q x_i \log_2 \left(\frac{y_i}{x_i} \right) \leq 0 \quad (6.4-2)$$

Note that the equality holds *only* when *all* the $x_i = y_i$.

It is natural to ask for the conditions on a probability distribution that lead to the maximum entropy (the minimum clearly occurs when one $p_i = 1$ and all the others = 0). We have

$$H_2(S) = \sum_i p_i \log_2 \frac{1}{p_i} \quad \text{with} \quad \sum_i p_i = 1$$

We begin by considering the quantity

$$\begin{aligned} H_2(S) - \log_2 q &= \sum_{i=1}^q p_i \log_2 \left(\frac{1}{p_i} \right) - \log_2 q \sum_{i=1}^q p_i \\ &= \sum_{i=1}^q p_i \log_2 \left(\frac{1}{qp_i} \right) \end{aligned}$$

Using the Gibbs inequality (6.4-2) with $y_i = 1/q$, we have

$$H_2(S) - \log_2 q \leq 0$$

Therefore,

$$H_2(S) \leq \log_2 q \quad (6.4-3)$$

For equality in (6.4-3) we must have all the $p_i = 1/q$.

For an alternative derivation of this important result, we use the Lagrange multipliers, and consider the function

$$f(p_1, p_2, \dots, p_q) = \frac{1}{\log_e 2} \sum_i p_i \log_e \left(\frac{1}{p_i} \right) + \lambda \left(\sum_i p_i - 1 \right)$$

$$\frac{\partial f}{\partial p_j} = \frac{1}{\log_e 2} \left[\log_e \left(\frac{1}{p_j} \right) - 1 \right] + \lambda = 0 \quad (j = 1, 2, \dots, q)$$

Since everything in this equation is a constant, it follows that each p_j is the same constant. If all the p_j are equal, each has the value $1/q$, where, as usual, the distribution has q members. Thus the maximum entropy is

$$H_2(S) = \log_2 q$$

For any other distribution than the equally likely distribution, the entropy is less than $\log_2 q$.

Based on this result, we now ask: What book has the maximum information? Of course, the answer depends on the size of the book, the size of the type font, the variety of the type font, and so on. Supposing these to be fixed; then the answer is obviously, the book with uniformly random letters! Each letter will come as a complete surprise!

This shows how far the formal engineering definition of information is from the usual human meaning of information. Yes, the definition of information that Shannon made is proper for computing machines, for the telephone system, and for communication systems, but flies in the face of the normal human meaning of information. It was the failure to notice that the definition is appropriate for machines and for machine-like situations, but is not appropriate for what we normally mean by "information," that led to so many misapplications of information theory in the early days. Now that we understand the limitations of the definition we are better able to apply information theory usefully. The very limitation of the definition to machines is what makes the definition so useful in such situations—the definition eliminates much that is vague, uncertain, and confusing. By limiting the theory to a sharp definition, we can get sharp, useful results in many situations. Perhaps it would have been better had Shannon called it *communication theory*, but the title *information theory* is often appropriate.

A simple, familiar application of this principle—that the uniform distribution contains the most information—is given by the standard grading system of A, B, C, D, and F. Setting aside F as having special properties, namely that the student must take the course again, *if* we want to communicate the maximum amount of information with the grading system, then we should use all the other grades equally often. Of course, we may want especially to distinguish the very best, and give comparatively few A's. This is, however, giving a different meaning to information than we have defined. The common habit in graduate schools of giving only A's and B's is a plain waste of signaling capacity. The extreme of one probability being 1 and all the others therefore being 0 is the case of a constant signal. There is no information transmitted. "The constant complainer is soon ignored."

In dealing with a new situation it is often useful to look at the extremes that can arise. For entropy of a distribution $H_2(S)$ we get the least value, for one extreme, $p_1 = 1$, and all other $p_i = 0$ —no surprise:

$$H_2(1, 0, 0, \dots, 0) = 0$$

As we proved, the maximum occurs when all the probabilities are equal, namely

$$H_2\left(\frac{1}{q}, \frac{1}{q}, \dots, \frac{1}{q}\right) = \ln_2 q$$

6.5 Entropy and Coding

We now prove a fundamental relationship between the average code length L and the entropy $H(S)$. Given *any instantaneous code* it has some definite code word lengths l_i represented in some radix r . From the Kraft inequality (4.5-1), we have

$$K = \sum_{i=1}^q \left(\frac{1}{r^{l_i}}\right) \leq 1 \quad (6.5-1)$$

We now define the numbers Q_i (pseudo probabilities):

$$Q_i = \frac{r^{-l_i}}{K} \quad (6.5-2)$$

where, of course,

$$\sum_{i=1}^q Q_i = 1$$

The Q_i may be regarded as a probability distribution. Therefore, we can use the fundamental Gibbs inequality (6.4-2),

$$\sum_{i=1}^q p_i \log_2 \left(\frac{Q_i}{p_i}\right) \leq 0$$

Upon expanding the log term into a sum of logs, we notice that one

term leads to the entropy function,

$$H_2(S) = \sum_{i=1}^q p_i \log_2 \left(\frac{1}{p_i} \right) \leq \sum_{i=1}^q p_i \log_2' \left(\frac{1}{Q_i} \right)$$

Using (6.5-2) in the right-hand side, we obtain

$$\begin{aligned} &\leq \sum_{i=1}^q p_i (\log_2 K - \log_2 r^{-l_i}) \\ &\leq \log_2 K + \sum_{i=1}^q p_i l_i \log_2 r \end{aligned}$$

By the Kraft inequality $K \leq 1$, so that $\log_2 K \leq 0$. Dropping this term can only strengthen the inequality. We have, therefore,

$$H_2(S) \leq \sum_{i=1}^q (p_i l_i) \log_2 r = L \log_2 r$$

or

$$H_r(S) \leq L \quad (6.5-3)$$

where L is the average code word length,

$$L = \sum_{i=1}^q p_i l_i \quad (6.5-4)$$

This is the fundamental result that we need; the entropy supplies a lower bound on the average code length L for any instantaneous decodable system. By the McMillan inequality of Section 4.7, it also applies to any uniquely decodable system.

For efficient binary codes $K = 1$ and we have $\log_2 K = 0$. Therefore, the inequality occurs (in the binary case) only when

$$p_i \neq Q_i = 2^{-l_i}$$

Table 6.3-1 shows how Huffman coding approaches the entropy for the Zipf distribution.

6.6 Shannon–Fano Coding

In Section 6.5 we assumed that the code word lengths l_i were given. Suppose, as is more likely, that the probabilities p_i are given. Huffman coding gives the lengths, but that method makes each length depend on the *whole* set of probabilities.

Shannon–Fano coding is less efficient than is Huffman coding, but has the advantage that you can go directly from the probability p_i to the code word length l_i . Given the source symbols $s_1, s_2, \dots, s_q$ and their corresponding probabilities $p_1, p_2, \dots, p_q$, then for each p_i there is an integer l_i such that

$$\log_r \left(\frac{1}{p_i} \right) \leq l_i < \log_r \left(\frac{1}{p_i} \right) + 1 \quad (6.6-1)$$

since the two extreme values just span a unit length. Removing the logs, we get

$$\left(\frac{1}{p_i} \right) \leq r^{l_i} < \left(\frac{r}{p_i} \right)$$

Take the reciprocal of each term. We obtain

$$p_i \geq \left(\frac{1}{r^{l_i}} \right) > \frac{p_i}{r}$$

†

Since $\sum p_i = 1$, when we sum this inequality, we get

$$1 \geq \sum_{i=1}^q \left(\frac{1}{r^{l_i}} \right) > \frac{1}{r} \quad (6.6-2)$$

which gives the Kraft inequality. Therefore, there is an instantaneous decodable code having these Shannon–Fano lengths.

To get the entropy of the distribution of p_i , we multiply the equation (6.6-1) by p_i and sum:

$$H_r(S) = \sum_{i=1}^q p_i \log_r \frac{1}{p_i} \leq \sum_{i=1}^q p_i l_i < H_r(S) + 1$$

In terms of the average length L of the code (6.5-4), we have

$$H_r(S) \leq L < H_r(S) + 1 \quad (6.6-3)$$

Thus for Shannon–Fano coding we again have the entropy as a lower bound on the average length of the code. It is also part of the upper bound. It is not as easy to directly find an upper bound for Huffman coding (Chapter 4), but since Huffman coding is optimal, it is at least as good as Shannon–Fano.

How do we find the actual code symbols? We simply assign them in order. For example, from the probabilities

$$p_1 = p_2 = \frac{1}{4}, \quad p_3 = p_4 = p_5 = p_6 = \frac{1}{8}$$

we get the Shannon–Fano lengths,

$$l_1 = l_2 = 2, \quad l_3 = l_4 = l_5 = l_6 = 3$$

We then assign

$$\begin{array}{ll} s_1 = 00 & s_3 = 100 \\ s_2 = 01 & s_4 = 101 \\ & s_5 = 110 \\ & s_6 = 111 \end{array}$$

We are assured by the Kraft inequality, which we showed that Shannon–Fano coding obeys, that there are always enough symbols to assign for an instantaneous code. Our orderly assignment leads readily to the decoding tree and the prefix condition is met.

This shows that we can do fairly well even with Shannon–Fano coding, but it is interesting to see how the left-hand inequality arises. First, if each probability p_i were exactly equal to the reciprocal of a power of the radix r , then we would have equality for this assignment of code word lengths—the average word length would be exactly the entropy. To get the code words we simply assign r -ary numbers in increasing sequence and of the required lengths.

6.7 How Bad Is Shannon–Fano Coding?

Since Huffman coding is optimal, and we have temporarily descended to the less-than-optimal (at times) Shannon–Fano coding, it is reasonable to ask “How bad is Shannon–Fano coding?” Consider the following examples.

For a source alphabet s_1, s_2 with probabilities

$$p_1 = 1 - \frac{1}{2^k}, \quad p_2 = \frac{1}{2^k} \quad (k \geq 2)$$

we get

$$\log_2 \left(\frac{1}{p_1} \right) \leq \log_2 2 = 1$$

Thus $l_1 = 1$. But for l_2 , we have

$$\log_2 \left(\frac{1}{p_2} \right) = \log_2 2^k = k = l_2$$

Hence where the Huffman encoding gives both code words one binary digit, Shannon–Fano has s_1 , a 1-bit word, and s_2 , a k -bit word.

Before getting too worried about this inefficiency, let us compute the average word length. Clearly, Huffman has

$$L_H = 1$$

For Shannon–Fano, we have

$$L_{\text{SF}} = 1 \left(1 - \frac{1}{2^k} \right) + k \left(\frac{1}{2^k} \right) = 1 + \left(\frac{k-1}{2^k} \right)$$

We have the table

k	$1 + \frac{k-1}{2^k}$
2	$1 + \frac{1}{4} = 1.25$
3	$1 + \frac{1}{4} = 1.25$
4	$1 + \frac{3}{16} = 1.1875$
5	$1 + \frac{1}{8} = 1.125$
6	$1 + \frac{5}{64} = 1.078125$
	etc.

As another example of the possible loss of efficiency, suppose that all the code symbols have the same probability, $1/q$. Shannon–Fano would have

$$\log_2 q \leq l < \log_2 q + 1$$

If q is not a power of 2, then Huffman would shorten some of the symbols, but Shannon–Fano would not.

The loss in the average code length for using Shannon–Fano is not great, and we shall later see that it does not matter in the theory; only in practice does Huffman dominate Shannon–Fano coding.

6.8 Extensions of a Code

We introduced the idea of entropy in a simple situation, namely *independent* symbols s_i with fixed probabilities p_i of occurring. We need to develop the idea of entropy for both the extensions (Section 4.2) and Markov processes (Sections 5.2 to 5.5) we studied earlier. This we now do in Sections 6.8 and 6.10.

The loss in Shannon–Fano coding compared to Huffman, when measured by the entropy function, occurs when the reciprocals of the probabilities are not exact powers of the radix r . If we encoded not one symbol at a time, but made up a code for blocks of n symbols of the source code at a time, then we could hope to get a better approximation to the lower bound $H_r(S)$.

But more important, as we noted earlier (Section 4.10), the probabilities of the extension are more variable than the original probabilities; so we expect that the larger the extension, the more we will find that both Huffman and Shannon–Fano encoding are efficient. We first review the concept of the extension of a code.

Definition. The n th extension of a source alphabet $s_1, s_2, \dots, s_q$ symbols of the form $s_{i_1}s_{i_2} \dots s_{i_n}$ having the probabilities $Q_i = p_{i_1}p_{i_2} \dots p_{i_n}$. Each block of the n original symbols now becomes a single symbol t_i with probability Q_i . We label this alphabet $S^n = T$.

We saw this definition in Sections 4.2 and 5.5, so it is not a new idea. The entropy of this new system is easily calculated as follows:

$$\begin{aligned} H_r(S^n) &= H_r(T) = \sum_{i=1}^{q^n} Q_i \log_r \left(\frac{1}{Q_i} \right) \\ &= \sum_{i=1}^{q^n} Q_i \log_r \left(\frac{1}{p_{i_1}p_{i_2} \dots p_{i_n}} \right) \end{aligned}$$

Expand the log term as a sum of logs. We get

$$\begin{aligned}
 &= \sum_{i=1}^{q^n} Q_i \sum_{k=1}^n \log_r \left(\frac{1}{p_{i_k}} \right) \\
 &= \sum_{k=1}^n \sum_{i=1}^{q^n} p_{i_1} p_{i_2} \cdots p_{i_n} \log_r \left(\frac{1}{p_{i_k}} \right)
 \end{aligned}$$

For a typical single term that occurs in the summation over k , we have

$$\begin{aligned}
 \sum_{i=1}^{q^n} p_{i_1} p_{i_2} \cdots p_{i_n} \log_r \left(\frac{1}{p_{i_k}} \right) \\
 = \sum_{i_1=1}^q \sum_{i_2=1}^q \cdots \sum_{i_n=1}^q p_{i_1} p_{i_2} \cdots p_{i_n} \log_r \left(\frac{1}{p_{i_k}} \right)
 \end{aligned}$$

For each summation *not* involving i_k , the sum is 1 (since the sum of the probabilities must be 1). That leaves only the summation over i_k . But

$$\sum_{i_k=1}^q p_{i_k} \log_r \left(\frac{1}{p_{i_k}} \right) = H_r(S)$$

Hence the sum of all the terms, since they all have the same form, is

$$H_r(T) = H_r(S^n) = nH_r(S) \quad (6.8-1)$$

The entropy of the n th extension of an alphabet is n times the entropy of the original alphabet. But note also that the extension has q^n symbols.

We can now apply the result (6.6-3) to the extension T . We have (where L_n is the average code length of the n th extension)

$$H_r(S^n) \leq L_n < H_r(S^n) + 1$$

Apply (6.8-1) and divide by n ,

$$H_r(S) \leq \frac{L_n}{n} < H_r(S) + \frac{1}{n} \quad (6.8-2)$$

Since the extension has n symbols s_i in each t_i , it follows that a better measure is $L = L_n/n$. Thus for a sufficiently large n th extension of the code, we can bring the average code word length L as close as we please to the entropy $H_r(S)$. This is Shannon's *noiseless coding theorem*: *The n th extension of a code satisfies (6.8-2).*

6.9 Examples of Extensions

Suppose that the source alphabet is s_1, s_2 with $p_1 = \frac{2}{3}, p_2 = \frac{1}{3}$. The entropy is

$$\begin{aligned} H_2(S) &= \frac{2}{3} \log_2 \frac{3}{2} + \frac{1}{3} \log_2 3 \\ &= \log_2 3 - \frac{2}{3} = 1.5849 \dots - 0.6666 \dots \\ &= 0.9182958 \dots \end{aligned}$$

Huffman coding gives

$$s_1 = 0, \quad s_2 = 1$$

so that the average length of the code is 1. Shannon–Fano would give

$$l_1 = 1, \quad l_2 = 2$$

and the average length is

$$\left(\frac{2}{3}\right)1 + \left(\frac{1}{3}\right)2 = \frac{4}{3} = 1 + \frac{1}{3}$$

In Section 4.10 we gave the average Huffman code lengths for several extensions. For Shannon–Fano we get, for the second extension,

$$\log_2 \frac{9}{4} \leq l_1, \quad l_1 = 2$$

$$\log_2 \frac{9}{2} \leq l_2, \quad l_2 = 3$$

$$\log_2 \frac{9}{2} \leq l_3, \quad l_3 = 3$$

$$\log_2 \frac{9}{1} \leq l_4, \quad l_4 = 4$$

The average code length is

$$\begin{aligned} L_{\text{SF}} &= \left(\frac{4}{9}\right)2 + \left(\frac{2}{9}\right)(3)(2) + \left(\frac{1}{9}\right)4 \\ &= \frac{8 + 12 + 4}{9} = \frac{24}{9} = 2.666 \dots \end{aligned}$$

It happens that for this case we can compute the extensions for Shannon–Fano coding. In the n th extension there will be

Number of Terms	Probability
$C(n, 0)$	$\frac{2^n}{3^n}$
$C(n, 1)$	$\frac{2^{n-1}}{3^n}$
.	.
.	.
.	.
$C(n, i)$	$\frac{2^{n-i}}{3^n}$
.	.
.	.
.	.
$C(n, n)$	$\frac{1}{3^n}$

To find the corresponding l_i we have

$$\begin{aligned}
 \log_2 \frac{1}{p_i} &= \log_2 \left(\frac{3^n}{2^{n-i}} \right) = n \log_2 3 - (n - i) \\
 &= n(\log_2 3 - 1) + i \leq l_i \quad (i = 1, 2, \dots, n)
 \end{aligned} \tag{6.9-1}$$

Let

$$A_n = \text{first integer} \geq n(\log_2 3 - 1) \tag{6.9-2}$$

Then the average code word length is

$$\begin{aligned}
 L_{\text{SF}}^{(n)} &= \frac{1}{3^n} \sum_{i=0}^n C(n, i) 2^{n-i} (A_n + i) \\
 &= \frac{1}{3^n} A_n \sum_{i=0}^n C(n, i) 2^{n-i} + \frac{1}{3^n} \sum_{i=0}^n C(n, i) i 2^{n-i}
 \end{aligned}$$

For the first sum, we have

$$(2 + 1)^n = \sum_{i=0}^n C(n, i) 2^{n-i} = 3^n$$

For the second sum, we have

$$(2 + x)^n = \sum C(n, i) x^i 2^{n-i}$$

Differentiate with respect to x :

$$n(2 + x)^{n-1} = \sum C(n, i) i x^{i-1} 2^{n-i}$$

Set $x = 1$:

$$n3^{n-1} = \sum C(n, i) i 2^{n-i}$$

Assembling the parts, we have

$$\begin{aligned} L_{\text{SF}}(n) &= \frac{1}{3^n} A_n(3^n) + \frac{1}{3^n} n3^{n-1} \\ &= A_n + \frac{n}{3} \end{aligned}$$

where from (6.9-2),

$$A_n \geq n(\log_2 3 - 1) = n(0.5849625 \dots)$$

We compute a short table of A_n and include in the last column of the table the results for Huffman coding.

n	$n(\log_2 3 - 1)$	A_n	$L_{\text{SF}}(n)/n$	$L_{\text{H}}(n)/n$
1	0.5849 . . .	1	1.33333 . . .	1.00000
2	1.1699 . . .	2	1.33333 . . .	0.94444 . . .
3	1.7549 . . .	2	1.00000 . . .	0.93827 . .
4	2.3398 . . .	3	1.08333 . . .	0.93827 . . .
5	2.9248 . . .	3	0.93333 . . .	0.92263 . . .
6	3.5098 . . .	4	1.00000 . . .	
7	4.0947 . . .	5	1.04762 . . .	
8	4.6797 . . .	5	0.95833 . . .	
9	5.2696 . . .	6	1.00000 . . .	
10	5.8496 . . .	6	0.93333 . . .	

From (6.5-3) the lower bound is $H(S) = 0.91829 \dots$. Shannon–Fano coding approaches the entropy irregularly, but according to the earlier result (Section 6.8), it does approach it finally.

Summarizing our discussion up to this point:

1. From equation (6.5-3), for any instantaneous code the Kraft inequality shows that

$$H_2(S) \leq L \log_2 r$$

2. Equation (6.6-3) shows that

$$H_2(S) \leq L_{\text{SF}} \log_2 r < H_2(S)^* + 1$$

3. In Section 4.10, since Huffman coding is optimal,

$$L_{\text{H}} \log_2 r \leq L_{\text{SF}} \log_2 r$$

4. In equation (6.8-2) the n th extension satisfies

$$H_2(S) \leq \frac{L \log_2 r}{n} \leq H_2(S) + \frac{\log_2 r}{n}$$

Assembling all this we get for the n th extension,

$$H_2(S) \leq \frac{L_{\text{H}}(n) \log_2 r}{n} \leq \frac{L_{\text{SF}}(n) \log_2 r}{n} < H_2(S) + \frac{\log_2 r}{n}$$

where $L_{\text{H}}(n)$ is the average code length for the n th extension for Huffman and $L_{\text{SF}}(n)$ is for Shannon–Fano. The table bears this out experimentally. This running inequality shows clearly why in Section 6.7 we are justified in saying that the theoretical loss in going to Shannon–Fano coding instead of the optimal Huffman coding was not serious *for large n* . For small n the practical loss may be significant.

Exercises

6.9-1 Do example for $p = \frac{1}{2}$.

6.9-2 Do for general value of p .

6.10 Entropy of a Markov Process

In Section 5.2 we introduced the idea of a Markov process to enable us to represent certain types of structure in the input stream of data. In particular, it handles correlations between successive symbols. We now examine the entropy of a Markov process.

Given that we have already seen the previous m symbols,

$$s_{i_1}, s_{i_2}, \dots, s_{i_m}$$

what is the probability that the next symbol will be s_i ? In mathematical symbols this conditional probability is written as

$$p(s_i \mid s_{i_1}, s_{i_2}, \dots, s_{i_m})$$

Note again that the sequence of the symbols is

$$s_{i_1}, s_{i_2}, \dots, s_{i_m}, s_i$$

For the zeroth-order Markov process the probabilities depend only on the symbol s_i , and is written in the earlier notation as p_i . For the first-order Markov source we have the digraph frequencies of pairs of symbols of the alphabet.

Using our “surprise” definition for the amount of information we get when we receive a symbol s_i , we define the information as

$$I(s_i \mid s_{i_1}, s_{i_2}, \dots, s_{i_m}) = \log_2 \left(\frac{1}{p(s_i \mid s_{i_1}, s_{i_2}, \dots, s_{i_m})} \right)$$

given that we have already seen (are in the state)

$$s_{i_1}, s_{i_2}, \dots, s_{i_m}$$

The *conditional entropy* over the alphabet S of symbols s_i is naturally given by the expression

$$\begin{aligned} H(S \mid s_{i_1}, s_{i_2}, \dots, s_{i_m}) \\ = \sum_S p(s_i \mid s_{i_1}, s_{i_2}, \dots, s_{i_m}) \log_2 \left(\frac{1}{p(s_i \mid s_{i_1}, s_{i_2}, \dots, s_{i_m})} \right) \end{aligned}$$

This is the conditional entropy of the source alphabet S *given* that we have the sequence of symbols

$$s_{i_1}, s_{i_2}, \dots, s_{i_m}$$

We next consider the larger system and allow for the probability of the states of the Markov process occurring. Let

$$p(s_{i_1}, s_{i_2}, \dots, s_{i_m})$$

be the probability of being in the state

$$s_{i_1}, s_{i_2}, \dots, s_{i_m}$$

Then it is natural to define the entropy of the Markov system as the probabilities of being in a state times the conditional entropies of the states, that is, by the expression

$$H(S) = \sum_{s^m} p(s_{i_1}, s_{i_2}, \dots, s_{i_m}) H(s_i | s_{i_1}, s_{i_2}, \dots, s_{i_m})$$

Using our earlier definition of the conditional entropy, we get

$$\begin{aligned} H(S) = \sum_{s^m} \sum_S p(s_{i_1}, s_{i_2}, \dots, s_{i_m}) p(s_i | s_{i_1}, s_{i_2}, \dots, s_{i_m}) \\ \times \log \left[\frac{1}{p(s_i | s_{i_1}, s_{i_2}, \dots, s_{i_m})} \right] \end{aligned}$$

But since

$$p(s_{i_1}, s_{i_2}, \dots, s_{i_m}) p(s_i | s_{i_1}, s_{i_2}, \dots, s_{i_m}) = p(s_{i_1}, s_{i_2}, \dots, s_{i_m}, s_i)$$

we have the entropy of the Markov process as

$$H(S) = \sum_{s^{m+1}} p(s_{i_1}, s_{i_2}, \dots, s_{i_m}, s_i) \log_2 \left[\frac{1}{p(s_i | s_{i_1}, s_{i_2}, \dots, s_{i_m})} \right]$$

6.11 Example of a Markov Process

As an example consider the Markov process (in a binary alphabet) shown in Figure 6.11–1. In symbols the second-order Markov process is defined by

$$p(0 | 0, 0) = 0.8 = p(1 | 1, 1)$$

$$p(1 | 0, 0) = 0.2 = p(0 | 1, 1)$$

$$p(0 | 0, 1) = 0.5 = p(1 | 1, 0)$$

$$p(1 | 0, 1) = 0.5 = p(0 | 1, 0)$$




Figure 6-11.1 Markov process

What is the probability of being in the particular states 0,0; 0,1; 1,0; and 1,1? Using the obvious symmetry, we see that

$$p(0, 0) = p(1, 1) \quad \text{and} \quad p(0, 1) = p(1, 0)$$

Therefore, we must have the equations

$$\begin{aligned} p(0, 0) &= 0.5p(0, 1) + 0.8p(0, 0) \\ p(0, 1) &= 0.2p(0, 0) + 0.5p(0, 1) \end{aligned}$$

These equations both come to the same form:

$$0.2p(0, 0) - 0.5p(0, 1) = 0 \quad (6.11-1)$$

Since we must be some place, we also have

$$p(0, 0) + p(0, 1) + p(1, 0) + p(1, 1) = 1$$

or

$$2p(0, 0) + 2p(0, 1) = 1 \quad (6.11-2)$$

The solution of (6.11-1) and (6.11-2) is

$$p(0, 0) = \frac{5}{14} = p(1, 1)$$

$$p(0, 1) = \frac{2}{14} = p(1, 0)$$

These are the *equilibrium probabilities* of the ergodic process.

In the general case, these equilibrium probabilities can be computed in principle, but in practice they are often hard to find. Continuing with the example, we now have the table for the equilibrium state.

s_{i1}	s_{i2}	s_i	$p(s_i s_{i1}, s_{i2})$	$p(s_{i1}, s_{i2})$	$p(s_{i1}, s_{i2}, s_i)$
0	0	0	0.8	$\frac{5}{14}$	$\frac{4}{14}$
0	0	1	0.2	$\frac{5}{14}$	$\frac{1}{14}$
0	1	0	0.5	$\frac{2}{14}$	$\frac{1}{14}$
0	1	1	0.5	$\frac{2}{14}$	$\frac{1}{14}$
1	0	0	0.5	$\frac{2}{14}$	$\frac{1}{14}$
1	0	1	0.5	$\frac{2}{14}$	$\frac{1}{14}$
1	1	0	0.2	$\frac{5}{14}$	$\frac{1}{14}$
1	1	1	0.8	$\frac{5}{14}$	$\frac{4}{14}$

From the table we can now compute

$$\begin{aligned}
 H(S) &= \sum_{2^3} p(s_{i1}, s_{i2}, s_i) \log_2 \left[\frac{1}{p(s_i | s_{i1}, s_{i2})} \right] \\
 &= 2 \left[\frac{4}{14} \log_2 \left(\frac{1}{0.8} \right) \right] + 2 \left[\frac{1}{14} \log_2 \left(\frac{1}{0.2} \right) \right] + 4 \left[\frac{1}{14} \log_2 \left(\frac{1}{0.5} \right) \right] \\
 &= \frac{4}{7} (\log_2 10 - 3) + \frac{1}{7} (\log_2 10 - 1) + \frac{2}{7} (\log_2 2) \\
 &= 0.801377 \dots \text{ bits per binary digit}
 \end{aligned}$$

6.12 The Adjoint System

To find the entropy of the Markov source, we need to find the stationary probabilities of being in each state of the Markov process. But as noted before, they are often very hard to find. Therefore, we need some

method of bounding the entropy. This method, which we now develop, is based on the concept of an *adjoint system*.

How can we find a bound on the entropy of a Markov process? We will use a first-order Markov process for simplicity; nothing except the notation differs for the n th-order system.

Since we have found that the fundamental Gibbs inequality for two distributions is our basic tool, we write it down for the distributions we know, namely $p(s_i)$, $p(s_j)$, and $p(s_j, s_i)$. The symbols $p(s_j, s_i)$ is the probability of the pair of symbols, and is also $p(s_i, s_j)$. We have, already,

$$\sum_j p(s_j, s_i) = p(s_i) = p_i$$

and

$$\sum_i p(s_j, s_i) = p(s_j) = p_j$$

Thus we start with, using (6.4-2),

$$\sum_{s^2} p(s_j, s_i) \log_2 \left[\frac{p(s_i)p(s_j)}{p(s_j, s_i)} \right] \leq 0$$

The equality occurs only when for all i and j we have the equalities $p(s_j, s_i) = p(s_i)p(s_j)$. We now use the conditional probability expression

$$p(s_j, s_i) = p(s_i | s_j)p(s_j)$$

to get

$$\sum_{s^2} p(s_j, s_i) \log_2 \left[\frac{p(s_i)}{p(s_i | s_j)} \right] \leq 0$$

or expanding the log term and transposing, we get

$$\sum_{s^2} p(s_j, s_i) \log_2 \left[\frac{1}{p(s_i | s_j)} \right] \leq \sum_i \sum_j p(s_j, s_i) \log_2 \left[\frac{1}{p(s_i)} \right]$$

Therefore, again using the conditional probability,

$$\begin{aligned} \sum_j p(s_j) \sum_i p(s_i | s_j) \log_2 \left[\frac{1}{p(s_i | s_j)} \right] &\leq \sum_i p(s_i) \log_2 \left[\frac{1}{p(s_i)} \right] \\ &\leq H(\bar{S}) \end{aligned}$$

where $H(\bar{S})$ is the entropy of the original symbols, which we will call the *adjoint system*. We can therefore write

$$\sum_j p(s_j)H(S | s_j) \leq H(\bar{S})$$

and the entropy of the Markov process is bounded by the entropy of the adjoint system, which is the zero memory source with probabilities of the source alphabet $p(s_i) = p_i$. The equality holds only if $p(s_j, s_i) = p_j p_i$. Thus we have proved the result that *constraints can only lower the entropy*.

The entropy of an extension of a Markov process is easily found from our earlier entropy calculations of a Markov source and of an extension. Again we consider, for the sake of simple notation, only a first-order Markov chain. We have

$$\begin{aligned} H(S^n) &= \sum_{S^n} \sum_{s^n} p(t_j, t_i) \log_2 \left[\frac{1}{p(t_i | t_j)} \right] \\ &= \sum_{s^{2n}} p(t_j, t_i) \log_2 \left[\frac{1}{p(t_i | t_j)} \right] \end{aligned}$$

But it is a first-order process; hence

$$p(t_j, t_i) = p(s_{i1}, s_{i2}, \dots, s_{in}, s_{j1})$$

We now apply the original conditional probabilities *repeatedly* to get

$$\begin{aligned} p(t_j | t_i) &= p(s_{i1} | s_j) p(s_{i2} | s_{i1}) \dots p(s_{in} | s_{i_{n-1}}) \\ H(S^n) &= \sum_{S^n} p(t_j, t_i) \log_2 \left[\frac{1}{p(s_i | s_j)} \right] + \dots \\ &\quad + \sum_{s^{2n}} p(t_j, t_i) \log_2 \left[\frac{1}{p(s_{in} | s_{i_{n-1}})} \right] \end{aligned}$$

Each term may be reduced by summing over all the symbols not in the log term:

$$\begin{aligned} \sum_{s^{2n}} p(t_j, t_i) \log_2 \frac{1}{p(s_{i1} | s_{i2})} &= \sum_{s^2} p(s_j, s_i) \log_2 \left[\frac{1}{p(s_i | s_j)} \right] \\ &= H(\bar{S}) \end{aligned}$$

Therefore,

$$H(S^n) = nH(\bar{S})$$

as we expect.

6.13 The Robustness of Entropy

We have just seen the central role entropy plays in coding theory, and it is natural to ask how robust (sensitive) it is with respect to small changes in the distribution $P = (p_1, p_2, \dots, p_q)$. There are two cases of importance; first, the maximum entropy distribution $p_i = 1/q$; and, second, the general distribution. The changes in the distribution can be written in the form

$$p'_i = p_i + \Delta p_i \quad \text{with } \sum \Delta p_i = 0$$

or else in the relative change form,

$$p'_i = p_i(1 + e_i) \quad \text{with } \sum p_i e_i = 0$$

We can use the Taylor expansion about the original distribution:

$$f(x) = f(a) + f'(a)(x - a) + f''(a) \frac{(x - a)^2}{2!} + f'''(a) \frac{(x - a)^3}{3!} + \dots$$

or else the Cauchy form with the mean value θ :

$$f(x) = f(a) + f'(a)(x - a) + f''(\theta) \frac{(x - a)^2}{2} \quad (a < \theta < x)$$

We can also get upper and lower bounds in the various cases, and so on. Thus many different formulas can be derived that measure the robustness of the entropy function. We will derive only one example, that using the additive error.

We begin with the robustness of the maximum entropy distribution, where $p_i = 1/q$ for all i . Using the n variable representation of the Taylor series, and for the moment regarding all the Δp_i as independent variables, we have

$$H(S') = H(S) + \sum \frac{\partial H}{\partial p_i} \Delta p_i + \frac{1}{2} \sum \frac{\partial^2 H}{\partial p_i^2} (\Delta p_i)^2 + \dots$$

We have for the individual terms in the expansion

$$\begin{aligned}
 H_e(S') &= - \sum \left(\frac{1}{q} + \Delta p_i \right) \log_e \left(\frac{1}{q} + \Delta p_i \right) \\
 \frac{H_e(S')}{\partial \Delta p_i} &= - \sum \log_e \left(\frac{1}{q} + \Delta p_i \right) - 1 \\
 \frac{\partial^2 H_e(S')}{\partial (\Delta p_i)^2} &= - \frac{1}{1/q + \Delta p_i}
 \end{aligned}$$

Therefore, evaluating these derivatives at $\Delta p_i = 0$, we have

$$\begin{aligned}
 H_e(S') &= H_e(S) - \left(q \log_e \frac{1}{q} + 1 \right) \sum \Delta p_i - \frac{1}{2} \sum q (\Delta p_i)^2 + \dots \\
 &= H_e(S) - \frac{q}{2} \sum (\Delta p_i)^2 + \dots
 \end{aligned}$$

We recognize the sum as being effectively the variance of the changes from the uniform distribution. As you would expect, at an extreme the value of the entropy changes very little for small changes in the distribution.

6.14 Summary

We have introduced the concept of entropy and have applied it to various sources of information to obtain bounds on what can be done by suitable encoding of the source. Entropy, which measures uncertainty, surprise, and information, is clearly the natural mathematical tool to use in such situations. In practice, entropy has uses far beyond this limited area of application, but we cannot go into them here.

Shannon's *noiseless* coding theorem shows that by using a suitably large extension of the source, the average length of the encoded message can be brought as close to the entropy of the source as desired. We demonstrated this both for a simple source and for the more complicated Markov source. However, we have not done it for all possible types of sources. The noiseless theorem is a forerunner of Shannon's main result that even in the presence of noise (errors), we can find a suitable encoding system. We prove this in Chapter 10.

The Channel and Mutual Information

7.1 Introduction

Let us review where we are. We are concerned with sending information either from here to there (transmission), or from now to then (storage). As shown in Figure 7.1-1, we start with a source of information (symbols) on the left and encode it in some fashion, such as was discussed in Chapters 2 through 5. It is then sent through a channel. In actual practice the channel is noisy, in the sense that the output is not always the same as the input. Next, the message is decoded, and finally the decoded message is sent on to its destination.

It is time to look more closely at the channel. We will use the concept of entropy developed in Chapter 6. How much information actually gets through the channel? How shall we measure the *mutual information* of the two ends of the system? How shall we define the *capacity* of the channel? These are the questions we examine in this and the next chapter.




Figure 7.1-1 Signaling model

Furthermore, we must learn to look at the system from the receiving end, where the decisions are to be made, instead of looking at the front end and following what happens to the signal as it goes through the system.

7.2 The Information Channel

An information channel is a statistical model of the medium through which the signal passes (or in which it is stored). In practice, there are physical limitations on the fidelity with which the transmission can occur. We need to formalize this idea in order to compute how much information goes through a given channel.




Figure 7.2-1 Channel

Figure 7.2-1 illustrates what we mean by a channel. A channel is described by a set of conditional probabilities $P(b_j | a_i)$, which are the probabilities that an input a_i from an alphabet of q letters will appear as some b_j from an alphabet of s letters. We use a capital P for the channel probabilities. The sizes of q and s of the alphabets need not be the same. For example, in an error-correcting code, the alphabet of possible received symbols is much larger than is the alphabet of sent symbols; thus $s \geq q$. On the other hand, a channel may always emit the same symbol when either of two different input symbols are sent, so that $s \leq q$.

In this model the channel is completely described by the matrix of

conditional probabilities

$$P_{i,j} = P(b_j | a_i)$$

Note the customary reversal of the subscripts. A row contains all the probabilities that a particular input symbol a_i becomes the output symbol b_j . This matrix appears in Figure 7.2-1.

This approach evades all questions of signal form, detection equipment, band width, and so on, and concentrates on the information.

The channel transition matrix has the following properties:

1. The i th row corresponds to the i th input symbol a_i .
2. The j th column corresponds to the j th output symbol b_j .
3. The sum of the elements in a row is always 1, that is,

$$\sum_{j=1}^s P_{i,j} = \sum_{j=1}^s P(b_j | a_i) = 1$$

This merely means that for each input a_i we are certain that something will come out, and the $P(b_j | a_i)$ give the distribution of these probabilities.

4. If $p(a_i)$ is the probability of the input symbol a_i of occurring, then

$$\sum_{i=1}^q \sum_{j=1}^s P(b_j | a_i) p(a_i) = 1$$

This means that when something is put into the system, then certainly something comes out.

The probabilities $P_{i,j}$ characterize the channel completely. We are, of course, supposing that the channel is *stationary*, meaning that the probabilities do not change with time, and for the moment we are supposing that the errors that do occur are independent of each other.

7.3 Channel Relationships

In the transmission system the input symbols are selected with probabilities $p(a_i)$, and consequently the received symbols b_j will have some corresponding probabilities of occurring $p(b_j)$. (This notation is clear, although it does not satisfy the strict canons of mathematics.) The two sets of probabilities $p(a_i)$ and $p(b_j)$ are related by the following *channel*

equations:

$$\begin{aligned}
 p(a_1)P_{1,1} + p(a_2)P_{2,1} + \dots + p(a_q)P_{q,1} &= p(b_1) \\
 p(a_1)P_{1,2} + p(a_2)P_{2,2} + \dots + p(a_q)P_{q,2} &= p(b_2) \\
 &\vdots \\
 &\vdots \\
 &\vdots \\
 p(a_1)P_{1,s} + p(a_2)P_{2,s} + \dots + p(a_q)P_{q,s} &= p(b_s)
 \end{aligned} \tag{7.3-1}$$

or in a more compressed notation,

$$\sum_{i=1}^q p(a_i)P(b_j | a_i) = p(b_j) \quad (j = 1, 2, \dots, s)$$

If the channel had no noise, then for each i , one $P_{i,j} = 1$ and all others = 0. Thus one $p(a_i) = p(b_j)$. For all noise each output for a given a_i has the same probability, namely $p(a_i)/s$; hence $p(b_j) = 1/s$.

For a uniform input $p(a_i) = 1/q$; hence

$$p(b_j) = \frac{1}{q} \sum_i P(b_j | a_i)$$

We write the probability of a pair of symbols a_i and b_j occurring as $P(a_i, b_j)$. This has the same meaning as $P(b_j, a_i)$. We can write this probability in two ways: either as

$$P(a_i, b_j) = p(a_i)P(b_j | a_i)$$

or as

$$P(a_i, b_j) = p(b_j)P(a_i | b_j)$$

Equating these two forms of the same thing, we get [$p(b_j) \neq 0$]

$$P(a_i | b_j) = \frac{P(b_j | a_i)p(a_i)}{p(b_j)} \tag{7.3-2}$$

This is known as *Bayes' theorem* on conditional probabilities. It is clearly symmetric in the two sets of symbols. The $P(b_j | a_i)$ are known as the *forward conditional probabilities*, since they start at the front with

the a_i given and express the probabilities of occurrence of the b_j . The $P(a_i | b_j)$ are correspondingly known as the *backward conditional probabilities*; given the output, what symbol probably caused it?

In the Bayes' formula (7.3-2) we may use the channel relationships (7.3-1) for the $p(b_j)$ and put them in the denominator to get the equivalent expression,

$$P(a_i | b_j) = \frac{P(b_j | a_i)p(a_i)}{\sum_{i=1}^q P(b_j | a_i)p(a_i)} \quad (7.3-3)$$

Summing this over all the a_i clearly gives

$$\sum_{i=1}^q P(a_i | b_j) = 1$$

which means that given output b_j , some a_i was certainly put into the channel.

7.4 Example of the Binary Symmetric Channel

A binary channel is probably the single most useful model of a channel. It has two input symbols, 0 and 1, and two output symbols, 0 and 1. Figure 7.4-1 illustrates this channel. The binary channel is said to be *symmetric* if

$$P_{0,0} = P_{1,1}, \quad P_{0,1} = P_{1,0}$$




Figure 7.4-1 Binary symmetric channel

If $P_{1,0} \neq P_{0,1}$, then the engineer would probably adjust things so that they were equal, so that the chance of a 0 becoming a 1 would be equal to the chance of a 1 becoming a 0.

Let the probabilities of the input symbols be (lowercase letters)

$$p(a = 0) = p$$

$$p(a = 1) = 1 - p$$

Next, let the binary symmetric channel probabilities be (upper case letters)

$$P_{0,0} = P_{1,1} = P$$

$$P_{0,1} = P_{1,0} = Q$$

The channel matrix is therefore

$$\begin{pmatrix} P & Q \\ Q & P \end{pmatrix}$$

The channel relationships (7.3-1) become

$$pP + (1 - p)Q = p(b = 0) \tag{7.4-1}$$

$$pQ + (1 - p)P = p(b = 1)$$

Note that these equations can be simply checked by computing their sum:

$$\begin{aligned} p(b = 0) + p(b = 1) &= p(P + Q) + (1 - p)(P + Q) \\ &= (p + 1 - p)(P + Q) = (1)(1) = 1 \end{aligned}$$

Given that we know what symbol we received, what are the probabilities for the various symbols that might have been sent?

We first compute the two denominators for equation (7.3-3):

$$\sum_{i=1}^2 P(b_1 | a_i)p(a_i) = Pp + Q(1 - p)$$

$$\sum_{i=1}^2 P(b_2 | a_i)p(a_i) = Qp + P(1 - p)$$

We then have

	No Noise	All Noise
$P(a = 0 b = 0) = \frac{Pp}{Pp + Q(1 - p)}$	1	p
$P(a = 1 b = 0) = \frac{Q(1 - p)}{Pp + Q(1 - p)}$	0	$1 - p$
$P(a = 0 b = 1) = \frac{Qp}{Qp + P(1 - p)}$	0	p
$P(a = 1 b = 1) = \frac{P(1 - p)}{Qp + P(1 - p)}$	1	$1 - p$

which involve the choice of the probabilities of the source.

For no noise $P = 1$, $Q = 0$, while for all noise $P = \frac{1}{2} = Q$. In the special case of equally likely input symbols ($p = \frac{1}{2}$), we have the very simple equations

$$\begin{aligned} P(a = 0 | b = 0) &= P = P(a = 1 | b = 1) \\ P(a = 1 | b = 0) &= Q = P(a = 0 | b = 1) \end{aligned}$$

As an example of a peculiar channel, suppose that $P = \frac{9}{10}$ and $Q = \frac{1}{10}$ for the binary symmetric channel probabilities. But suppose also that the probability of the input $a = 0$ is sent $\frac{19}{20} = p$ and $a = 1$ is sent $\frac{1}{20} = (1 - p)$. We then have

$$\begin{aligned} P(a = 0 | b = 0) &= \frac{171}{172} \sim 1 \\ P(a = 1 | b = 0) &= \frac{1}{172} \\ P(a = 0 | b = 1) &= \frac{19}{28} \sim 1 \\ P(a = 1 | b = 1) &= \frac{9}{28} \end{aligned}$$

Thus if you receive a $b = 0$, then almost certainly an $a = 0$ was sent. But if a $b = 1$ is received, it is still a $\frac{19}{28} > \frac{2}{3}$ probability that $a = 0$ was sent. Therefore, regardless of what symbol is received, you must assume that an $a = 0$ was sent!

This situation arises whenever both

$$\begin{aligned} P(a = 0 | b = 0) &> P(a = 1 | b = 0) \\ P(a = 0 | b = 1) &> P(a = 1 | b = 1) \end{aligned}$$

From the earlier equations for the binary symmetric channel these conditions are

$$Pp > Q(1 - p)$$

$$Qp > P(1 - p)$$

or both

$$\begin{array}{l} p > Q \\ p > P \end{array}$$

In words, this says that the bias in the choice of input symbols is greater than the bias of the channel. In short, clearly channels can be misused.

7.5 System Entropies

In Chapter 6 we developed the concept of the entropy of an information source (which we now label A). The entropy measures the average amount of information per symbol of the source, or equivalently, the uncertainty associated with the source. $H(A) = 0$ when the source is certain and $H(A)$ is a maximum when all the a_i are equally likely. We have the definition of the entropy as

$$H_r(A) = \sum_{i=1}^q p(a_i) \log_r \left(\frac{1}{p(a_i)} \right) \quad (7.5-1)$$

As noted, the entropy function has the following properties:

1. $H_r(A) \geq 0$.
2. $H_r(A) \leq \log_r q$, where q is the number of input symbols.
3. $H_r(A) = \log_r q$ when all the source symbols are equally likely.

The rest of this chapter is simply the systematic use of the entropy function applied to various distributions of the model of the signaling system.

Clearly, corresponding to (7.5-1), the entropy of the *received* symbols should be defined by

$$H_r(B) = \sum_{j=1}^s p(b_j) \log_r \left[\frac{1}{p(b_j)} \right] \quad (7.5-2)$$

This quantity measures the uncertainty of the output symbols and has the same properties as the entropy of the source, since all the properties were derived from the form of the entropy function.

Similar expressions for the conditional entropy may be derived as follows. We naturally have for the *conditional entropy*, given a particular b_j ,

$$H_r(A | b_j) = \sum_{i=1}^q P(a_i | b_j) \log_r \left[\frac{1}{P(a_i | b_j)} \right] \quad (7.5-3)$$

If we average (7.5-3) over all the b_j , using of course the appropriate probabilities $p(b_j)$ as weights, we get

$$\begin{aligned} H_r(A | B) &= \sum_{j=1}^s p(b_j) H(A | b_j) \\ &= \sum_{j=1}^s \sum_{i=1}^q p(b_j) P(a_i | b_j) \log_r \left[\frac{1}{P(a_i | b_j)} \right] \\ &= \sum_{j=1}^s \sum_{i=1}^q P(a_i, b_j) \log_r \left[\frac{1}{P(a_i | b_j)} \right] \end{aligned} \quad (7.5-4)$$

Doing exactly the same thing *except* that we view it from the sending end, corresponding to (7.5-3), we get the conditional entropy having sent a_i :

$$H_r(B | a_i) = \sum_{j=1}^s P(b_j | a_i) \log_r \left[\frac{1}{P(b_j | a_i)} \right] \quad (7.5-5)$$

Similar to (7.5-4), averaging this time over the input alphabet we get the corresponding conditional entropy,

$$H_r(B | A) = \sum_{i=1}^q \sum_{j=1}^s P(a_i, b_j) \log_r \left[\frac{1}{P(b_j | a_i)} \right] \quad (7.5-6)$$

We now connect two ends of the channel together. In order to measure the uncertainty of the joint event for both source and receiver, we define the *joint entropy* in a similar way:

$$H_r(A, B) = \sum_{i=1}^q \sum_{j=1}^s P(a_i, b_j) \log_r \left[\frac{1}{P(a_i, b_j)} \right] \quad (7.5-7)$$

First we examine the joint entropy for the special case when A and B are *statistically independent* for all i and j (meaning that what comes out does *not* depend on what goes in, all noise). The independence

condition is expressed by the equation

$$P(a_i, b_j) = p(a_i)p(b_j)$$

The joint entropy therefore becomes

$$H_r(A, B) = \sum_{i=1}^q \sum_{j=1}^s p(a_i)p(b_j) \left\{ \log_r \left[\frac{1}{p(a_i)} \right] + \log_r \left[\frac{1}{p(b_j)} \right] \right\}$$

Since

$$\sum_{i=1}^q p(a_i) = 1 \quad \text{and} \quad \sum_{j=1}^s p(b_j) = 1$$

we get (for independent input and output)

$$H_r(A, B) = H_r(A) + H_r(B) \quad (7.5-8)$$

Normally, the output of a channel depends at least partly on the input (otherwise, why are we using it?). In this case can write

$$P(a_i, b_j) = p(a_i)P(b_j | a_i)$$

and the joint entropy is

$$\begin{aligned} H_r(A, B) &= \sum_{i=1}^q \sum_{j=1}^s P(a_i, b_j) \log_r \left[\frac{1}{p(a_i)} \right] \\ &\quad + \sum_{i=1}^q \sum_{j=1}^s P(a_i, b_j) \log_r \left[\frac{1}{P(b_j | a_i)} \right] \end{aligned}$$

But

$$\sum_j P(a_i, b_j) = p(a_i)$$

Therefore,

$$\begin{aligned} H_r(A, B) &= \sum_{i=1}^q p(a_i) \log_r \left[\frac{1}{p(a_i)} \right] \\ &\quad + \sum_{i=1}^q \sum_{j=1}^s P(a_i, b_j) \log_r \left[\frac{1}{P(b_j | a_i)} \right] \quad (7.5-9) \\ &= H_r(A) + H_r(B | A) \end{aligned}$$

The joint entropy is the sum of the source entropy and the conditional entropy given A .

The conditional entropy $H_r(B | A)$ represents the information loss in the channel in going from input to output. It is how much must be added to the source entropy to get the joint entropy. $H_r(B | A)$ is called *the equivocation*. It is also called the *noise entropy* of the channel.

Since A and B play symmetric roles, we clearly have the corresponding expression to (7.5-9),

$$H_r(A, B) = H_r(B) + H_r(A | B) \quad (7.5-10)$$

In summary, in this section we have defined various entropies: $H(B)$ by (7.5-2), $H(A, B)$ by (7.5-7), $H(A | B)$ by (7.5-4), $H(B | A)$ by (7.5-6), $H(A | b_j)$ by (7.5-3), and $H(B | a_i)$ by (7.5-5). Each is derived by using the appropriate probability distribution.

7.6 Mutual Information

Consider again the transmission system of Figure 7.2-1. The input symbols are the a_i and the output are the b_j , and the channel is defined by the conditional probabilities P_{ij} .

Prior to reception the probability of the input symbol a_i was $p(a_i)$. This is the *a priori probability* of a_i . After reception of b_j , the probability that the input symbol was a_i becomes $P(a_i | b_j)$, the conditional probability that a_i was sent given that you received b_j . This is an *a posteriori probability* of a_i . The change in the probability measures how much the receiver learned from the reception of the b_j . In the ideal channel with no noise, the a posteriori probability is 1, since you are certain from the received b_j exactly what was sent. In practical systems there are finite nonzero probabilities that errors will occur and you, the receiver, cannot be absolutely sure what was sent.

The difference between the information uncertainty before (the a priori probabilities) and after reception of a b_j (the a posteriori probabilities) measures the gain in information due to the reception of the b_j . This information is called the *mutual information* and is naturally defined as

$$\begin{aligned} I(a_i; b_j) &= \log_r \left[\frac{1}{p(a_i)} \right] - \log_r \left[\frac{1}{P(a_i | b_j)} \right] \\ &= \log_r \left[\frac{P(a_i | b_j)}{p(a_i)} \right] \end{aligned} \quad (7.6-1)$$

If the two probabilities $p(a_i)$ and $P(a_i | b_j)$ are the same, then we have gained no information and the mutual information is zero. No information has been transmitted. It is only when we have learned something new about the probabilities of the a_i from the received b_j that the mutual information can be positive.

Multiply numerator and denominator of the log term of (7.6-1) by $p(b_j)$. Since

$$P(a_i | b_j)p(b_j) = P(a_i, b_j) = P(b_j | a_i)p(a_i)$$

we have

$$I(a_i; b_j) = \log_r \left[\frac{P(a_i, b_j)}{p(a_i)p(b_j)} \right] = I(b_j; a_i) \quad (7.6-2)$$

The mutual information $I(a_i; b_j)$ has the following properties:

1. From the symmetry of the a_i and b_j , the definition gives

$$I(a_i; b_j) = \log_r \left[\frac{P(a_i | b_j)}{p(a_i)} \right]$$

$$I(b_j; a_i) = \log_r \left[\frac{P(b_j | a_i)}{p(b_j)} \right]$$

and $I(a_i; b_j) = I(b_j; a_i)$.

2. We also have

$$I(a_i; b_j) \leq I(a_i) \quad (7.6-3)$$

This follows from the definition

$$I(a_i) = \log_r \left[\frac{1}{p(a_i)} \right]$$

and

$$I(a_i; b_j) = \log_r P(a_i | b_j) + I(a_i)$$

The maximum for the probability $P(a_i | b_j)$ is 1, and the maximum for its log is therefore 0:

$$I(a_i; b_j) \leq I(a_i)$$

3. If a_i and b_j are independent, that is, if

$$P(a_i | b_j) = p(a_i)$$

or what is the same thing, if

$$P(a_i, b_j) = p(a_i)p(b_j)$$

then

$$I(a_i; b_j) = 0$$

Because of the inevitable noise the behavior of a channel can be understood only on the average. We therefore begin by averaging the mutual information over the alphabets using the appropriate probabilities of the symbols occurring:

$$\begin{aligned} I(A; b_j) &= \sum_i P(a_i | b_j) I(a_i; b_j) \\ &= \sum_i P(a_i | b_j) \log_r \left[\frac{P(a_i | b_j)}{p(a_i)} \right] \end{aligned} \quad (7.6-4)$$

Similarly,

$$I(a_i; B) = \sum_j P(b_j | a_i) \log_r \left[\frac{P(b_j | a_i)}{p(b_j)} \right] \quad (7.6-5)$$

Finally, we get

$$\begin{aligned} I(A; B) &= \sum_{i=1} p(a_i) I(a_i; B) \\ &= \sum_i \sum_j P(a_i, b_j) \log_r \left[\frac{P(a_i, b_j)}{p(a_i)p(b_j)} \right] \\ &= I(B; A) \end{aligned} \quad (7.6-6)$$

by symmetry. The first two, (7.6-4) and (7.6-5), are called the *conditional mutual information*. $I(A; b_j)$ measures the information gain provided by the reception of b_j . $I(a_i; B)$ is the information gain about the alphabet B given that a_i was sent.

The third, $I(A; B)$, which is symmetric in the two alphabets, provides a measure of the information gain of the whole system and does

not depend on the individual input and output symbols but only on their probabilities; it is called the *system mutual information*. The system mutual information has the properties:

1. $I(A; B) \geq 0$ (from fundamental inequality, Section 6.4).
2. $I(A; B) = 0$ if and only if A and B are independent.
3. $I(A; B) = I(B; A)$ (from symmetry).

We can relate these various entropies to each other by the following algebraic manipulations. For example, from the Gibbs inequality (6.4-2), we have

$$\begin{aligned}
 I(A; B) &= \sum_{i=1}^q \sum_{j=1}^s P(a_i, b_j) \log \left[\frac{P(a_i, b_j)}{p(a_i)p(b_j)} \right] \geq 0 \\
 &= \sum_{i=1}^q \sum_{j=1}^s P(a_i, b_j) [\log P(a_i, b_j) \\
 &\quad - \log p(a_i) - \log p(b_j)] \geq 0 \\
 &= - \sum_i \sum_j P(a_i, b_j) \log \left[\frac{1}{P(a_i, b_j)} \right] \\
 &\quad + \sum_i p(a_i) \log \left[\frac{1}{p(a_i)} \right] + \sum_j p(b_j) \log \left[\frac{1}{p(b_j)} \right] \geq 0 \\
 &= H(A) + H(B) - H(A, B) \geq 0
 \end{aligned}$$

Using the results (7.5-9) and 7.5-10), namely

$$\begin{aligned}
 H(A, B) &= H(A) + H(B | A) \\
 &= H(B) + H(A | B)
 \end{aligned}$$

we get the two results

$$\begin{aligned}
 I(A; B) &= H(A) - H(A | B) \geq 0 \\
 &= H(B) - H(B | A) \geq 0
 \end{aligned} \tag{7.6-7}$$

Therefore,

$$\begin{aligned}
 0 &\leq H(A | B) \leq H(A) \\
 0 &\leq H(B | A) \leq H(B)
 \end{aligned} \tag{7.6-8}$$

and

$$H(A, B) \leq H(A) + H(B) \quad (7.6-9)$$

Equation (7.6-9) shows that the joint entropy is maximum when the two alphabets are independent.

As examples of relationships, the average mutual information is

$$I(A; B) = \begin{cases} H(A) + H(B) - H(A, B) \\ H(A) - H(A | B) \\ H(B) - H(B | A) \end{cases}$$

The equivocation is

$$H(A | B) = H(A) - I(A; B)$$

$$H(B | A) = H(B) - I(A; B)$$

The joint entropy is

$$H(A, B) = \begin{cases} H(A) + H(B) - I(A; B) \\ H(A) + H(B | A) \\ H(B) + H(A | B) \end{cases}$$

Channel Capacity

8.1 Definition of Channel Capacity

Given the conditional probabilities $P(b_j | a_i) = P_{i,j}$ which define a channel, what is the maximum amount of information that we can send through the channel? This is the main question addressed in this chapter.

The mutual information connects the two ends of the channel together. It is defined by equation (7.6-7),

$$I(A; B) = H(A) - H(A | B)$$

where the entropy $H(A)$ is the uncertainty of the source *before* the reception of B , and $H(A | B)$ is the average uncertainty *after* the reception of the b_j of B . Thus $I(A; B)$ is the change in the information. An alternative expression for $I(A; B)$ is equation (7.6-6),

$$I(A; B) = \sum_A \sum_B P(a, b) \log_r \left[\frac{P(a, b)}{p(a)p(b)} \right]$$

where we have dropped subscripts and adopted an obvious summation convention.

This formula involves the input symbol frequencies $p(a)$. We saw in the example of a binary symmetric channel (Section 7.4) how a poor match of the $p(a)$ to the channel can ruin a channel. Indeed, we know that if the probability of one symbol is $p(a) = 1$, then all the others

must be zero and the constant signal contains no information. How can we best choose the $p(a)$ to get the most through the channel, and what is that amount?

We postpone the problem of choosing the $p(a)$ by using a standard mathematical trick; we simply define the *channel capacity* as the maximum over all possible assignments of the $p(a)$:

$$C = \max_{p(a)} I(A; B) \quad (8.1-1)$$

This makes the channel capacity depend only on the $P_{i,j}$ and temporarily evades the difficulty of finding the suitable $p(a)$. The probabilities $p(b)$ are determined from both the $p(a)$ and the $P_{i,j}$ via the channel equations (7.3-1). By this definition there is no question of ever having a higher mutual information than the channel capacity C , since by definition C is the maximum.

You want me to measure the capacity in bits per second (or some such units). But remember that storage is the same as transmission. Hence we measure capacity in bits of information per bits sent.

✓

8.2 The Uniform Channel

From equation (7.6-7) we have the mutual information

$$I(A; B) = H_r(B) - H_r(B | A)$$

From equation (7.5-6) this is

$$\begin{aligned} I(A; B) &= H_r(B) - \sum_A \sum_B P(a, b) \log_r \left[\frac{1}{P(b | a)} \right] \\ &= H_r(B) - \sum_A p(a) \sum_B P(b | a) \log_r \left[\frac{1}{P(b | a)} \right] \end{aligned}$$

In a highly symmetric code, such as a parity error-detecting code, a perfect error-correcting code, and many others, each symbol is equivalent, in a sense, to any other symbol in the code—the differences in the probabilities for errors in the various symbols are inessential. For example, we can convert any symbol into the 0, 0, . . . , 0 symbol (in the binary system) by the trivial process of complementing each position that has a 1 in it. If the channel has white noise, and this is what most of the codes were designed to handle, then the entries in any row of

the channel matrix, for the cases above, are merely a permutation (rearrangement) of those in the first row. Each symbol has the same set of transition probabilities as any other symbol.

Under the condition that each row is a permutation of the first row, it follows that the sums

$$\sum_B P(b | a) \log_r \left[\frac{1}{P(b | a)} \right] = W = H(B | a) \quad (8.2-1)$$

are independent of a . No matter which symbol a you pick, this *average will be the same*. We can, therefore, sum over all the symbols of the alphabet B and get for each a

$$I(A; B) = H_r(B) - W \sum_A p(a) = H_r(B) - W \quad (8.2-2)$$

Thus we see the use for the following definition (a special case of common occurrence).

Definition. A *uniform channel* has a matrix

$$\begin{pmatrix} P_{1,1} & P_{1,2} & \dots & P_{1,s} \\ P_{2,1} & P_{2,2} & \dots & P_{2,s} \\ \vdots & \vdots & & \vdots \\ \vdots & \vdots & & \vdots \\ \vdots & \vdots & & \vdots \\ P_{q,1} & P_{q,2} & & P_{q,s} \end{pmatrix}$$

such that each row is a permutation of the first row.

Suppose that we think of our channel as an *extension of the binary symmetric channel*. Each alphabet symbol is a block of n bits and each bit is subject to the same independent probability of success P (and failure Q). This definition makes the error-detecting and perfect error-correcting codes have symmetric channels.

As a trivial example, consider the noiseless channel where each input symbol a goes into exactly one output symbol b ; that is, for each given a , one $P(b | a) = 1$ and all other $P(b | a) = 0$ for that a . Each term, therefore, is

$$P(b | a) \log_r \left[\frac{1}{P(b | a)} \right] = 0$$

Hence $W = 0$ and

$$I(A; B) = H_r(B) = H_r(A)$$

as we expect.

8.3 Uniform Input

Given a uniform channel, clearly the sensible way to use it is to use uniform $p(a)$ for the input alphabet symbols. We have the formula (8.2-2)

$$I(A; B) = H_2(B) - W \quad (8.3-1)$$

Suppose that we apply this formula to an n -bit single-error-detecting binary code. [We know from (6.8-1) what the answer will be, but for exercise we will compute all the details.] The number of input symbols is $q = 2^{n-1}$, since (due to the parity check) only half of the possible sequences of n binary digits are possible inputs. Because of the errors that are possible, any binary sequence could come out. We have the following situation (where P = probability of no error per bit sent):

Number of Cases	Type	Probability
1	No error	P^n
n	One error	$P^{n-1}Q$
$\frac{n(n-1)}{2}$	Two errors	$P^{n-2}Q^2$
.	.	.
.	.	.
.	.	.
$C(n, k)$	k errors	$P^{n-k}Q^k$

But

$$W = \sum_B P(b | a) \log_2 \left[\frac{1}{P(b | a)} \right]$$

Grouping together all the cases having the same probability, we have

$$= \sum_{k=0}^n C(n, k) P^{n-k} Q^k \log_2 \left[\frac{1}{P^{n-k} Q^k} \right] \quad (8.3-2)$$

Expand the product in the log term into two log terms:

$$(n - k) \log_2 \left(\frac{1}{P} \right) + k \log_2 \left(\frac{1}{Q} \right) \quad (8.3-3)$$

From the first term we get for the sum (8.3-2) (to take care of the $n - k$ factor and then the n factor),

$$\begin{aligned} \log_2 \left(\frac{1}{P} \right) \sum_{k=0}^{n-1} \frac{n!}{k! (n-1-k)!} \frac{n-k}{n-k} P^{n-k} Q^k \\ = P \log_2 \left(\frac{1}{P} \right) \sum_{k=0}^{n-1} \frac{n(n-1)!}{k! (n-1-k)!} P^{n-1-k} Q^k \\ = nP \log_2 \left(\frac{1}{P} \right) (P + Q)^{n-1} = nP \log_2 \left(\frac{1}{P} \right) \end{aligned}$$

From the second term of (8.3-3) we get for the sum (8.3-2), correspondingly,

$$nQ \log_2 \left(\frac{1}{Q} \right)$$

Together they give

$$n \left[P \log_2 \left(\frac{1}{P} \right) + Q \log_2 \left(\frac{1}{Q} \right) \right] = nH_2(P)$$

where we have used the natural notation

$$H_2(P) = P \log_2 \left(\frac{1}{P} \right) + Q \log_2 \left(\frac{1}{Q} \right) \quad (8.3-4)$$

The mutual information (8.3-1) now becomes

$$I(A; B) = H_2(B) - nH_2(P) \quad (8.3-5)$$

The entropy of the n th extension is again n times the entropy of one symbol.

8.4 Error-Correcting Codes

In Section 1.3 we observed that sometimes the encoding process was broken up into two stages, one to suitably match the source of information, the second to match the channel. We may, if we wish, combine this second stage of encoding with the definition of the channel and call it part of the channel. Indeed, sometimes it is arranged so that as far as the user is concerned, it appears as if it were the channel. Let us examine what this channel looks like.

For a perfect error-correcting code of n bits and probability P of being correct for each single bit in the white noise channel, then for the channel errors $P_{i,j}$ the user sees the following table. Remember in this channel a single error is correctly corrected, and a double error is wrongly corrected into a triple error.

J	Occurrence of:	Actual Probability	H/A
	No errors	$P^n + nQP^{n-1}$	
	One error	0	
	Two errors	0	
	Three or more errors	$1 - [\text{no error}]$	

Thus the channel errors have a peculiar structure in terms of the original causes of errors. However, we may still apply the methods of the preceding section to this uniform channel.

When the error correction is either in hardware or is *automatically* included in the software, the user sees the above. It also indicates what was earlier observed: In practice, for a variable probability source, we are likely to use a Huffman code to match the source, and then apply some error-correcting code to blocks of this encoding. More and more often computers are being built with error detection and/or correction in the hardware. It is a matter of preference on your part whether you wish to regard this as a single process of encoding, or as two processes, or whether you wish to regard the second stage of encoding as belonging to the channel itself. How to handle the situation should be clear.


Exercises

8.4-1 Apply this section to the 7-bit single-error-correcting code.

8.4-2 Extend this section to minimum distance of four codes.

8.5 Capacity of a Binary Symmetric Channel

If P is the probability of correct transmission, then the channel matrix for a binary symmetric channel (Section 7.4)



is given as

$$\begin{pmatrix} P & Q \\ Q & P \end{pmatrix}$$

Let p = probability of choosing $a = 0$. Then $1 - p$ = probability of choosing $a = 1$ and the mutual information for this uniform channel is equation (8.2-2). Using $\log_2$ and temporarily dropping the subscript 2 on the entropy function H , we have

$$I(A; B) = H(B) - W$$

where, from (8.2-1),

$$\begin{aligned} W &= \sum_B P(b | a) \log \left[\frac{1}{P(b | a)} \right] \\ &= P \log \left(\frac{1}{P} \right) + Q \log \left(\frac{1}{Q} \right) \\ &= H(P) \end{aligned}$$

using (8.3-4). Therefore,

$$I(A; B) = H(B) - H(P) \quad (8.5-1)$$

To find $H(B)$ we note that [equation (7.4-1)]

$$p(b = 0) = pP + (1 - p)Q$$

$$p(b = 1) = pQ + (1 - p)P$$

It is easy to see that

$$p(b = 0) + p(b = 1) = 1$$

For convenience we set

$$p(b = 0) = pP + (1 - p)Q = x \quad (8.5-2)$$

Hence

$$p(b = 1) = 1 - x$$

and using the notation of (8.3-4),

$$\begin{aligned} I(A; B) &= x \log \left(\frac{1}{x} \right) + (1 - x) \log \left(\frac{1}{1 - x} \right) - H(P) \\ &= H(x) - H(P) \end{aligned}$$

To get the channel capacity C , we ask: What choice of the symbol probability p maximizes the mutual information $I(A; B)$? Since we are working with the variable x , this is the same as asking what maximizes

$$H(x) = x \log \left(\frac{1}{x} \right) + (1 - x) \log \left(\frac{1}{1 - x} \right)$$

From section 6.4 the maximum occurs for $x = \frac{1}{2}$, and

$$H\left(\frac{1}{2}\right) = \frac{1}{2} \log 2 + \frac{1}{2} \log 2 = \log 2 = 1$$

The capacity of the channel is therefore

$$C = \max_A I(A; B) = 1 - H(P) \quad (8.5-3)$$

We can find the value of p from (8.5-2) as follows:

$$x = \frac{1}{2} = pP + (1 - p)Q$$

$$\frac{1}{2} = pP + (1 - p)(1 - P)$$

$$= pP + 1 - p - P + pP$$

Transposing and rearranging, we get

$$(1 - 2P)p = \frac{1 - 2P}{2}$$

Since the channel probability $P \neq \frac{1}{2}$ (else no information is sent) it follows, on dividing by $(1 - 2P) \neq 0$, that the symbol probability $p(a_i)$ is

$$p = \frac{1}{2}$$

which is what we would expect from the symmetry.

Thus we have shown that the binary symmetric channel has the capacity

$$\boxed{C = 1 - H_2(P) = 1 - H_2(Q)} \quad (8.5-4)$$

which occurs when the two input symbols are chosen with equal frequency.

$$p_{\text{min}} = \frac{1}{2} \quad !!$$

8.6 Conditional Mutual Information

How can we find the symbol probabilities $p(a)$ to meet the channel capacity in more general cases? This section gives one very useful result that helps answer this question in many situations.

Returning to the general channel, we can write

$$\begin{aligned} I(A; B) &= \sum_A p(a) \sum_B P(b | a) \log \left[\frac{P(b | a)}{p(b)} \right] \\ &= \sum_A p(a) I(a; B) \end{aligned}$$

where we have written the *conditional mutual information* as

$$I(a; B) = \sum_B P(b | a) \log \left[\frac{P(b | a)}{p(b)} \right]$$

In general, $I(a; B)$ depends on a , but we have the theorem: *When the input symbols are picked to meet the channel capacity C , then*

$$I(a; B) = C$$

for every of a [assuming that $p(a) \neq 0$].

To prove this, we assume the contrary, that some symbol, say a_1 , is operating above capacity. Then, since the average is C , there must be some symbol, say a_2 , that is operating below C . Thus we assume

that

$$I(a_1; B) > C$$

$$I(a_2; B) < C$$

For an infinitesimal change that increases $p(a_1)$ and decreases $p(a_2)$ by the same amount (to keep the total probability equal to 1), the direct rate of change in $I(A; B)$ is the difference

$$\frac{\partial I(A; B)}{\partial p(a_1)} - \frac{\partial I(A; B)}{\partial p(a_2)}$$

But the change in the $p(a)$ induces a change through the channel equations in the $p(b)$. The rate of change in each $p(b)$ is given by

$$P(b | a_1) - P(b | a_2)$$

From the expression for the mutual information,

$$I(A; B) = \sum_A p(a) \sum_B P(b | a) \log \left[\frac{P(b | a)}{p(b)} \right]$$

we have, therefore, the instantaneous rate of change in $I(A; B)$ as

$$\begin{aligned} & \sum_B P(b | a_1) \log \left[\frac{P(b | a_1)}{p(b)} \right] - \sum_B P(b | a_2) \log \left[\frac{P(b | a_2)}{p(b)} \right] \\ & + \sum_A p(a) \left\{ - \sum_B P(b | a) \left[\frac{1}{p(b)} \right] [P(b | a_1) - P(b | a_2)] \right\} \\ & = I(a_1; B) - I(a_2; B) \\ & - \sum_B \sum_A p(a) \left[\frac{P(b | a)}{p(b)} \right] [P(b | a_1) - P(b | a_2)] \\ & = I(a_1; B) - I(a_2; B) \\ & - \sum_B [P(b | a_1) - P(b | a_2)] \\ & = I(a_1; B) - I(a_2; B) - (1 - 1) \\ & = I(a_1; B) - I(a_2; B) > 0 \end{aligned}$$

Since we were operating at channel capacity and it is the maximum that can occur, and since the change has an instantaneous rate that is positive, we see that if our hypothesis were true, we could get above the maximum by a suitably small change in $p(a_1)$ and $p(a_2)$. We conclude, therefore, that the hypothesis that any letter operates above the channel capacity must be wrong; *every symbol in the source alphabet operates at capacity when the whole alphabet operates at channel capacity.*

We can often use this result in particular situations to help us find the individual $p(a)$'s that meet channel capacity—we can find them one at a time.

Some Mathematical Preliminaries

9.1 Introduction

Some people believe that a theorem is proved when a logically correct proof is given; but some people believe that a theorem is proved *only* when the student sees why it is inevitably true. The author tends to belong to this second school of thought. Therefore, we make a digression at this point to prove, rather than assume, a number of results that are needed to understand the proof of Shannon's main result. The subjects are presented in sufficient detail so that a general appreciation of them is obtained, not just a mere acquaintance with the result. Thus more than the minimum material is occasionally presented.

The two main ideas are n -dimensional space, and the law of large numbers. We have already seen how n -dimensional space is useful for representing code symbols. Since we are going to use large extensions of codes to get near the entropy, and hence the channel capacity, the dimension of the space of the proof will be very high indeed.

The law of large numbers is needed to handle the fact that when we are near ideal signaling rates, the number of errors in a code word is very high. The manner in which the probable number of errors is distributed is central to the proof of Shannon's main result.

We will also need to bound the tail of a binomial distribution to show that it does not contribute very much. For this we need the famous Stirling approximation to $n!$ This approximation is often used but is rarely presented in a simple manner.

9.2 The Stirling Approximation to $n!$

The Stirling approximation to $n!$ is

$$n! \approx n^n e^{-n} \sqrt{2\pi n} \quad (9.2-1)$$

It is widely used but is seldom completely presented. We need it, so we shall derive it.

The quantity $n!$ is a product, and products are hard to approximate. This suggests approximating the log of $n!$:

$$\log n! = \sum_{k=1}^n \log k$$

(all logs in this section are $\log_e$). The summation in turn suggests the integral (which we do by integration by parts)

$$\begin{aligned} \int_1^n \log x \, dx &= (x \log x - x) \Big|_1^n \\ &= n \log n - n + 1 \end{aligned} \quad (9.2-2)$$




Figure 9.2-1 Trapezoid rule

If we use the trapezoid rule to approximate the integral (see Figure 9.2-1), we will under-approximate the area. Therefore ($\log 1 = 0$),

$$\begin{aligned} \int_1^n \log x \, dx &\geq \frac{1}{2} \log 1 + \log 2 + \log 3 + \dots \\ &\quad + \log(n-1) + \frac{1}{2} \log n \\ &\geq \log n! - \frac{1}{2} \log n \end{aligned}$$

We have, using (9.2-2), and transposing one term,

$$n \log n - n + \frac{1}{2} \log n + 1 \geq \log n!$$

Taking antilogs (which due to the positive growth and monotonicity of the log function preserves the inequality direction) we get, term by term, the *upper bound*,

$$n^n e^{-n} \sqrt{n} e \geq n! \quad (9.2-3)$$

To overapproximate the integral we use the midpoint integration formula and note (Figure 9.2-2) that if the midpoint line segment is




Figure 9.2-2 Midpoint formula

tilted until it is tangent to the curve, then the area under the line segment is not changed. For the first half-interval, $1 \leq x \leq \frac{3}{2}$, we use the tangent line at $x = 1$ to get a bounding triangle of area $\frac{1}{8}$. For the last half-interval we use the maximum, $\log n$, and get a bounding area $\frac{1}{2} \log n$. We have, therefore,

$$\begin{aligned} \int_0^n \log x \, dx &\leq \frac{1}{8} + \log 2 + \log 3 + \dots + \log(n-1) + \frac{1}{2} \log n \\ &\leq \log n! + \frac{1}{8} - \frac{1}{2} \log n \end{aligned}$$

Again use (9.2-2) and rearrange, to get

$$n \log n - n + 1 - \frac{1}{8} + \frac{1}{2} \log n \leq \log n!$$

Take antilogs, we get the *lower bound*,

$$n^n e^{-n} \sqrt{n} e^{7/8} \leq n! \quad (9.2-4)$$

The coefficient C of the approximation

$$n! = n^n e^{-n} \sqrt{n} C \quad (9.2-5)$$

lies, from (9.2-3) and (9.2-4), between

$$\boxed{e^{7/8} \leq C \leq e}$$

or, in numbers,

$$2.39887 \dots \leq C \leq 2.71828 \dots \quad (9.2-6)$$

Since the midpoint formula has approximately one-half the error of the trapezoid rule, and the errors of the two formulas have opposite signs, a better guess at C is the weighted average

$$\frac{2e^{7/8} + e}{3} = 2.50534 \dots \simeq C \quad (9.2-7)$$

It will turn out that the value for C (for infinite n) is

$$\sqrt{2\pi} = 2.50663 \dots$$

If we want to get this coefficient C exactly (for infinite n), we use a trick and consider the integral

$$I_k = \int_0^{\pi/2} \cos^k x \, dx \quad (k \geq 0)$$

A trigonometric substitution gives ($k \geq 2$)

$$\begin{aligned} I_k &= \int_0^{\pi/2} \cos^{k-2} x (1 - \sin^2 x) \, dx \\ &= I_{k-2} + \int_0^{\pi/2} \cos^{k-2} x (-\sin x) \sin x \, dx \end{aligned}$$

Integration by parts gives

$$\begin{aligned} I_k &= I_{k-2} + \frac{\cos^{k-1} x}{k-1} \sin x \Big|_0^{\pi/2} - \frac{1}{k-1} \int_0^{\pi/2} \cos^{k-1} x \cos x \, dx \\ &= I_{k-2} - \left(\frac{1}{k-1} \right) I_k \end{aligned}$$

Solving for I_k , we have the recurrence relation for the integral,

$$I_k = \left(\frac{k-1}{k} \right) I_{k-2} \quad (k \geq 2) \quad (9.2-8)$$

It is easy to show that

$$I_1 = 1$$

$$I_0 = \frac{\pi}{2}$$

We next note that since $\cos x \leq 1$ for all x in the interval,

$$I_{2k-1} > I_{2k} > I_{2k+1}$$

Divide through by I_{2k-1} :

$$1 > \frac{I_{2k}}{I_{2k-1}} > \frac{I_{2k+1}}{I_{2k-1}}$$

Substitute from (9.2-8) for the integrals to get

$$1 > 2k \left[\frac{2k-1}{2k} \frac{2k-3}{2k-2} \cdots \frac{3}{2} \right]^2 \frac{\pi}{2} > \frac{2k}{2k+1}$$

We can make the numerator in the square bracket equal to $(2k)!$ by multiplying both numerator and denominator by the denominator. We also use the fact that

$$2k(2k-2)(2k-4) \cdots 2 = 2^k k!$$

to get

$$1 > \pi k \left[\frac{(2k)!}{2^k (k!) 2^k (k!)} \right]^2 > \frac{2k}{2k+1}$$

We now take the positive square root of each term:

$$1 > \sqrt{\pi k} \left[\frac{(2k)!}{(k!)^2 2^{2k}} \right] > \sqrt{1 - \frac{1}{2k+1}}$$

Finally, we substitute the approximation (9.2-5)

$$n! = n^n e^{-n} \sqrt{n} C$$

into this for each factorial to get

$$1 > \sqrt{\pi k} \left[\frac{(2k)^{2k} e^{-2k} \sqrt{2k} C}{2^{2k} k^k e^{-k} \sqrt{k} C k^k e^{-k} \sqrt{k} C} \right] > \sqrt{1 - \frac{1}{2k+1}}$$

Practically everything cancels and we have

$$1 > \frac{\sqrt{2\pi}}{C} > \sqrt{1 - \frac{1}{2k+1}}$$

Therefore, for large k ,

$$C \longrightarrow \sqrt{2\pi}$$

Thus we have the Stirling approximation. As Table 9.2-1 shows, the ratio of Stirling's approximation to the true value is remarkably good even for small values. The table shows that the approximation was designed to be *relatively* exact for $n = \infty$, but the difference does *not* approach zero.

TABLE 9.2-1 *Stirling Approximation*

n	Stirling	True	Stirling/True
1	0.92214	1	0.92214
2	1.91900	2	0.95950
3	5.83621	6	0.97270
4	23.50618	24	0.97942
5	118.01916	120	0.98349
6	710.07818	720	0.98622
7	4,980.3958	5,040	0.98817
8	39,902.395	40,320	0.98964
9	359,536.87	362,880	0.99079
10	3,598,695.6	3,628,800	0.99170

If this seems strange, consider the two functions

$$f(n) = n + \sqrt{n}$$

$$g(n) = n$$

The limit of the ratio

$$\frac{f(n)}{g(n)} = \frac{n + \sqrt{n}}{n} = 1 + \frac{1}{\sqrt{n}} \longrightarrow 1$$

The difference

$$f(n) - g(n) = \sqrt{n} \longrightarrow \infty$$

9.3 A Binomial Bound

Another mathematical result we need is the number of vertices in the unit cube of Section 3.6 that are in a sphere of radius n . This comes from the Hamming bounds like (3.6-2). Thus we need a bound on a sum of binomial coefficients

$$1 + C(n, 1) + C(n, 2) + \dots + C(n, \lambda n) \quad (9.3-1)$$

where $0 < \lambda < \frac{1}{2}$. (We suppose that λn is an integer.)

Since $\lambda < \frac{1}{2}$, the terms are all monotonically growing in size, and the largest is therefore the last one,

$$C(n, \lambda n) = \frac{n!}{(\lambda n)! (n - \lambda n)!} \quad (9.3-2)$$

We apply the upper and lower Stirling's approximations (9.2-3) and (9.2-4) to each factorial as appropriate to get

$$C(n, \lambda n) \leq \frac{n^n e^{-n} \sqrt{n} e}{[(\lambda n)^{\lambda n} e^{-\lambda n} \sqrt{\lambda n} e^{7/8}] [(n - \lambda n)^{(n - \lambda n)} e^{-(n - \lambda n)} \sqrt{n - \lambda n} e^{7/8}]}$$

Regrouping terms, we get

$$\begin{aligned} C(n, \lambda n) &\leq \left[\frac{n^n}{n^{\lambda n} n^{(n - \lambda n)}} \right] \left[\frac{1}{\lambda^{\lambda n} (1 - \lambda)^{(1 - \lambda)n}} \right] \left[\frac{e^{-n}}{e^{-\lambda n} e^{-(n - \lambda n)}} \right] \\ &\quad \times \frac{e \sqrt{n}}{e^{(7/8)} \sqrt{\lambda n} e^{(7/8)} \sqrt{(1 - \lambda)n}} \end{aligned}$$

The first and third brackets are exactly 1 and the other two become

$$C(n, \lambda n) \leq \left[\frac{1}{\lambda^\lambda (1 - \lambda)^{(1 - \lambda)}} \right]^n \left[\frac{1}{e^{3/2} \lambda (1 - \lambda) n} \right]^{1/2} \quad (9.3-3)$$

The log to the base 2 of the denominator of the first bracket is

$$\log[\lambda^\lambda (1 - \lambda)^{1 - \lambda}] = \lambda \log \lambda + (1 - \lambda) \log (1 - \lambda) = -H(\lambda)$$

Hence the first bracket of (9.3-3) is

$$2^{nH(\lambda)} \quad (9.3-4)$$

The second bracket is a constant times $n^{-(1/2)}$.

We now write the original series (9.3-1) in the reverse direction

$$C(n, \lambda n) + C(n, \lambda n - 1) + \dots + C(n, 1) + 1$$

and we bound the sum by a geometric progression. To get the “rate” of the bounding geometric progression, note that in going from term to term in the *usual way*, each binomial coefficient is found from the

preceding one by multiplying by (term by term)

$$\frac{n}{1}, \frac{n-1}{2}, \frac{n-2}{3}, \dots, \frac{n-\lambda n+1}{\lambda n}$$

If we start from our maximum term $C(n, \lambda n)$ and go *down*, we have the multipliers which are the reciprocals of the previous numbers and are in the reverse order:

$$\frac{\lambda n}{n-\lambda n+1}, \dots, \frac{3}{n-2}, \frac{2}{n-1}, \frac{1}{n}$$

We need to bound these. An upper bound of these numbers is

$$\frac{\lambda}{1-\lambda}$$

so an upper bound on the sum of the binomial coefficient terms is the largest term times

$$\sum_{m=0}^{\infty} \left(\frac{\lambda}{1-\lambda} \right)^m = \frac{1}{1 - \frac{\lambda}{1-\lambda}} = \frac{1-\lambda}{1-2\lambda} \quad (9.3-5)$$

Assembling everything of (9.3-1), we have [use (9.3-4) and (9.3-5)]

$$\sum_{k=0}^{\lambda n} C(n, k) \leq 2^{nH(\lambda)} \left[\frac{1}{e^{3/2}\lambda(1-\lambda)} \right]^{1/2} \frac{1-\lambda}{1-2\lambda} \left(\frac{1}{n} \right)^{1/2}$$

For all n sufficiently large the product of the last three terms is less than 1; hence for sufficiently large n ,

$$\sum_{k=0}^{\lambda n} C(n, k) \leq 2^{nH(\lambda)} \quad (9.3-6)$$

where $H(\lambda)$ is the entropy function of λ . This is our needed inequality.

Note that the sum of *all* the binomial coefficients of order n is

$$\sum_{k=0}^n C(n, k) = 2^n$$

as can be seen from the generating function

$$(1 + x)^n = \sum_{k=0}^n C(n, k)x^k$$

when you set $x = 1$.

9.4 The Gamma Function $\Gamma(n)$

To extend the factorial function to nonintegral arguments we begin with the study of a certain definite integral as a function of a parameter n . Consider the integral

$$\Gamma(n) = \int_0^{\infty} e^{-x} x^{n-1} dx \quad (9.4-1)$$

The exponent $n - 1$ is conventional. For $n > 1$ we can integrate by parts to get

$$\Gamma(n) = x^{n-1} \frac{e^{-x}}{-1} \Big|_0^{\infty} + (n - 1) \int_0^{\infty} e^{-x} x^{n-2} dx$$

or

$$\Gamma(n) = (n - 1)\Gamma(n - 1)$$

For $n = 1$ we have

$$\Gamma(1) = \int_0^{\infty} e^{-x} dx = -e^{-x} \Big|_0^{\infty} = 1$$

and it follows that for integer n ,

$$\Gamma(2) = 1$$

$$\Gamma(3) = 2!$$

$$\Gamma(4) = 3!$$

⋮

$$\Gamma(n) = (n - 1)! \quad (n = 1, 2, \dots)$$

The gamma function provides the natural generalization for the factorial function since the integral also exists for noninteger n .

We consider next the gamma function for argument $\frac{1}{2}$. We have

$$\Gamma\left(\frac{1}{2}\right) = \int_0^{\infty} e^{-x} x^{-1/2} dx$$

The substitution $x = t^2$ rationalizes the integrand and gives

$$\Gamma\left(\frac{1}{2}\right) = \int_0^{\infty} e^{-t^2} \frac{2t}{t} dt = 2 \int_0^{\infty} e^{-t^2} dt$$

Since the integrand is an even function, we can write this as

$$\Gamma\left(\frac{1}{2}\right) = \int_{-\infty}^{\infty} e^{-t^2} dt$$

This is often called the *error integral*.

Consider, now, the product

$$\Gamma\left(\frac{1}{2}\right)\Gamma\left(\frac{1}{2}\right) = \int_{-\infty}^{\infty} e^{-x^2} dx \int_{-\infty}^{\infty} e^{-y^2} dy$$

We transform (formally) to polar coordinates

$$\begin{aligned} [\Gamma\left(\frac{1}{2}\right)]^2 &= \int_0^{2\pi} \int_0^{\infty} e^{-r^2} r dr d\theta \\ &= 2\pi \left. \frac{e^{-r^2}}{-2} \right|_0^{\infty} = \pi \end{aligned}$$

Therefore,

$$\Gamma\left(\frac{1}{2}\right) = \sqrt{\pi}$$

To justify the formal change to polar coordinate, suppose that we consider finite limits

$$I(L) = \int_{-L}^L e^{-x^2} dx \int_{-L}^L e^{-y^2} dy$$




Figure 9.4-1 Rectangular-to-polar coordinates

We have from Figure 9.4-1,

$$\int_0^{2\pi} \int_0^L e^{-r^2} r \, dr \, d\theta \leq I(L) \leq \int_0^{2\pi} \int_0^{L\sqrt{2}} e^{-r^2} r \, dr \, d\theta$$

or

$$\pi(1 - e^{-L^2}) \leq I(L) \leq \pi(1 - e^{-2L^2})$$

As $L \longrightarrow \infty$, we get

$$\pi \leq I(L) = \Gamma^2\left(\frac{1}{2}\right) \leq \pi$$

Thus the transformation from rectangular to polar coordinates is justified and we believe that $\Gamma\left(\frac{1}{2}\right) = \sqrt{\pi}$.

9.5 *n*-Dimensional Euclidean Space

n-Dimensional space merely means that we have *n* independent variables $x_1, x_2, \dots, x_n$. Euclidean space means that we use the Pythagorean distance

$$x_1^2 + x_2^2 + \dots + x_n^2 = r^2 \quad (9.5-1)$$

and we can define a *sphere of radius r* by this expression.

A little reflection on classical Euclidean geometry will suffice to see why the volume of an *n*-dimensional sphere depends on the radius *r* as r^n . Therefore, we have the formula for the volume

$$V_n(r) = C_n r^n \quad (9.5-2)$$

where C_n is some constant depending on *n*. For example,

$$C_2 = \pi \quad \text{and} \quad C_3 = \frac{4\pi}{3} \quad (\text{note that } C_1 = 2)$$

To find the values of C_n , we use the same trick as in Section 9.4 of multiplying the gamma integral by itself and then going to polar coordinates (which worked so well). Let us consider the product of *n* of the integrals. We have [using (9.5-1)]

$$\begin{aligned} [\Gamma(\tfrac{1}{2})]^n &= \pi^{n/2} = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} \dots \int_{-\infty}^{\infty} e^{-r^2} dx_1 dx_2 \dots dx_n \\ &= \int_0^{\infty} e^{-r^2} \frac{dV_n(r)}{dr} dr \end{aligned}$$

This last step takes a moment's thinking about the shell of thickness Δr and comparing this with our result in two dimensions,

$$\int_0^{\infty} e^{-r^2} \frac{d(\pi r^2)}{dr} dr = 2\pi \int_0^{\infty} e^{-r^2} r dr$$

as we really did in Section 9.4 (although it was disguised somewhat by the conventional method).

We have, therefore, using (9.5-2),

$$\pi^{n/2} = C_n \int_0^{\infty} e^{-r^2} n r^{n-1} dr$$

Set $r^2 = t$. Then $dr = \frac{1}{2}t^{-1/2} dt$, and

$$\begin{aligned}\pi^{n/2} &= \frac{nC_n}{2} \int_0^\infty e^{-t} \frac{t^{(n-1)/2} dt}{t^{1/2}} \\ &= \frac{nC_n}{2} \int_0^\infty e^{-t} t^{(n/2-1)} dt \\ &= \frac{nC_n}{2} \Gamma\left(\frac{n}{2}\right) = C_n \Gamma\left(\frac{n}{2} + 1\right)\end{aligned}$$

Therefore,

$$C_n = \frac{\pi^{n/2}}{\Gamma\left(\frac{n}{2} + 1\right)}$$

It is easy to show that

$$C_n = \frac{2\pi}{n} C_{n-2}$$

and we can compute the table (of the volume of a unit sphere in n dimensions)

n	C_n	
1	2	2.00000 . . .
2	π	3.14159 . . .
3	$4\pi/3$	4.18879 . . .
4	$\pi^2/2$	4.93480 . . .
5	$8\pi^2/15$	5.26379 . . .
6	$\pi^3/6$	5.16771 . . .
7	$16\pi^3/105$	4.72477 . . .
8	$\pi^4/24$	4.05871 . . .
$\vdots$	$\vdots$	$\vdots$
$\vdots$	$\vdots$	$\vdots$
$2k$	$\pi^k/k!$	$\longrightarrow 0$

From the table we see that the volume of the unit sphere, or equivalently the coefficient of r^n , comes to a maximum at $n = 5$ and falls off rather rapidly toward zero as n approaches infinity.

For $n = 2k$, the volume of an n -dimensional sphere of radius r is

$$\left(\frac{\pi^k}{k!}\right) r^{2k} = \frac{(\pi r^2)^k}{k!}$$

From this we see that once

$$k > \pi r^2$$

then increasing k (or equivalently n) will *decrease* the volume. Indeed, given any radius, no matter how large, the dimension of the space n can be increased until the volume of the sphere is arbitrarily small. [For odd-dimensional spaces, we have, from the original definition of $\Gamma(n)$, equation (9.4-1), that

$$V_n(r) = C_n r^n = \frac{\pi^{n/2} r^n}{\Gamma(n/2 + 1)}$$

will change smoothly with increasing n .]

We now consider the *fraction* of the volume of an n -dimensional sphere that is within a distance ϵ of the surface (ϵ is used here as any small positive number). We have

$$\frac{\text{shell}}{\text{volume}} = \frac{C_n r^n - C_n (r - \epsilon)^n}{C_n r^n} = 1 - \left(1 - \frac{\epsilon}{r}\right)^n$$

which approaches 1 as n gets large. Thus for a high-dimensional space, almost all the volume of a sphere is arbitrarily close to the surface. Indeed, no matter how thin a shell you wish to use, and how close you wish to get to 1 (say 99.44%), it is possible to find an n_0 such that for all greater n , the two conditions are met. There simply is not much volume inside a high-dimensional sphere; almost all of it is “on the surface.”

Let us now examine the angle between the vector from the origin $(0, 0, \dots, 0)$ to the point $(1, 1, \dots, 1)$ and any coordinate axis. The projection on each axis is, obviously, exactly equal to 1 (its coordinate value on that axis). The length of the vector is $\sqrt{n}$; hence the angle we want, θ , is given by

$$\cos \theta = \frac{1}{\sqrt{n}}$$

As $n \longrightarrow \infty$, this approaches 0; hence $\theta \longrightarrow \pi/2$. Thus for sufficiently large n , the diagonal line is “almost perpendicular” to each coordinate axis!

Exercises

9.5-1 Show for any family of convex similar figures that “almost all of the volume is on the surface.”

9.5-2 Show that the volume of a $2k$ -dimensional sphere divided by the corresponding volume of an enclosing cube is $(\pi/4)^k(1/k!)$.

9.6 A Paradox

The results of this section are not needed but are included to show how unreliable your intuition is about spheres in n -dimensional Euclidean space.




Figure 9.6-1 A paradox

Suppose, as shown in Figure 9.6-1, that we have a 4×4 square centered about the origin $(0, 0)$ and have four unit circles in each of the four corners. Now consider the radius of the circle about the origin that is tangent on the inside to the four circles. It has a radius

$$r_2 = \sqrt{(1 - 0)^2 + (1 - 0)^2} - 1 = \sqrt{2} - 1 = 0.414 \dots$$

Next consider the same situation in three dimensions. We have a $4 \times 4 \times 4$ cube with eight unit spheres in the corners. The inner sphere has the radius

$$r_3 = \sqrt{3} - 1 = 0.732 \dots$$

Finally, consider the similar situation in n dimensions. We have a $4 \times 4 \times \dots \times 4$ cube with 2^n unit spheres in the corners, each one touching all its n neighboring spheres. The spheres are packed in properly. The distance from the origin to the center of a sphere is given by the expression

$$\sqrt{(1 - 0)^2 + (1 - 0)^2 + \dots + (1 - 0)^2} = \sqrt{n}$$

Again subtracting the radius of the corner unit sphere, we get the radius of the inner sphere in n dimensions as

$$r_n = \sqrt{n} - 1 \quad (9.6-1)$$

When $n = 10$, this is

$$r_{10} = \sqrt{10} - 1 = 3.16 \dots - 1 = 2.16 \dots > 2$$

and the inner sphere reaches outside the cube! Does this seem impossible? The spheres are convex surfaces to be sure, the distance is surely the correct Euclidean distance, and the radius of the corner sphere is surely 1—so there is no escape from the conclusion. Contrary to any normal intuition, the inner sphere for $n \geq 10$ reaches outside the cube.

To compound the paradox further, consider the volume of this inner sphere *relative* to the whole cube as a function of n . We have, for the special case of n an even number $n = 2k$, and using (9.6-1),

$$\begin{aligned} \text{ratio} &= \frac{\text{volume of sphere}}{\text{volume of whole cube}} = C_n \frac{(\sqrt{n} - 1)^n}{4^n} \\ &= C_{2k} \frac{(\sqrt{2k} - 1)^{2k}}{4^{2k}} \\ &= \frac{\pi^k (\sqrt{2})^{2k} (\sqrt{k})^{2k}}{k! 4^{2k}} \left[\left(1 - \frac{1}{\sqrt{2k}} \right)^{\sqrt{2k}} \right]^{\sqrt{2k}} \\ &= \frac{\pi^k 2^k k^k}{4^{2k} k!} \left[\left(1 - \frac{1}{\sqrt{2k}} \right)^{\sqrt{2k}} \right]^{\sqrt{2k}} \end{aligned} \quad (9.6-2)$$

Using Stirling's approximation (9.3-1) for $k!$,

$$k! \simeq k^k e^{-k} \sqrt{2\pi k}$$

and the result from the calculus,

$$\lim_{x \rightarrow \infty} \left(1 - \frac{1}{x} \right)^x = e^{-1}$$

We get for (9.6-2),

$$\begin{aligned} \text{ratio} &\simeq \frac{\pi^k 2^k k^k}{4^{2k} k^k e^{-k} \sqrt{2\pi k}} e^{-\sqrt{2k}} \\ &\simeq \left(\frac{\pi e}{8} \right)^k \frac{e^{-\sqrt{2k}}}{\sqrt{2\pi k}} \end{aligned} \quad (9.6-3)$$

But

$$\frac{\pi e}{8} = 1.06747 \dots$$

Taking $\log_e$ of both sides, we get, as k approaches infinity, $\log_e(\text{ratio}) = (0.065288 \dots)k - \sqrt{2k} - \frac{1}{2} \log_e k - \frac{1}{2} \log_e 2 \longrightarrow \infty$. Therefore, we see that the factor

$$\left(\frac{\pi e}{8}\right)^k$$

will go to infinity faster than the other two factors in the product (9.6-3),

$$e^{-\sqrt{2k}} \quad \text{and} \quad \frac{1}{\sqrt{2\pi k}}$$

can decrease it.

We conclude that the volume of the *inner* sphere becomes arbitrarily larger than the volume of the cube which contains all the $2^{2k} = 2^n$ unit spheres in the corners.

The case of the odd dimension $n = 2k - 1$ does not change matters; it only makes the details more messy.

9.7 Chebyshev's Inequality and the Variance

If a random variable X is discrete, then the mean square is given by the expectation

$$E\{X^2\} = \sum_{i=-\infty}^{\infty} x_i^2 p(x_i)$$

and if the variable is continuous, it is given by

$$E\{X^2\} = \int_{-\infty}^{\infty} x^2 p(x) dx$$

For any $\epsilon > 0$, since the integrand is positive, we have

$$\int_{-\infty}^{\infty} x^2 p(x) dx \geq \int_{|x| \geq \epsilon} x^2 p(x) dx$$

and this can be written as

$$E\{X^2\} \geq \int_{|x| \geq \epsilon} x^2 p(x) dx \geq \epsilon^2 \int_{|x| \geq \epsilon} p(x) dx$$

It follows therefore that

$$E\{X^2\} \geq \epsilon^2 \text{Prob}\{|X| \geq \epsilon\}$$

or

$$\text{Prob}\{|X| \geq \epsilon\} \leq \frac{E\{X^2\}}{\epsilon^2} \quad (9.7-1)$$

This is the famous *Chebyshev inequality* (expressed either in continuous or discrete notation).

The *variance*, or dispersion, of a random variable X , labeled $V\{X\}$, is the mean square measured about the mean value

$$E\{X\} = a$$

Thus the variance of a random variable X is

$$\begin{aligned} V\{X\} &= E\{(X - a)^2\} \\ &= E\{X^2\} - 2aE\{X\} + E\{a^2\} \\ &= E\{X^2\} - 2a^2 + a^2 \\ &= E\{X^2\} - a^2 = \sigma^2 \end{aligned}$$

Clearly, since $E\{1\} = 1$, the variance has the property that for any constant

$$V\{c\} = 0$$

We also have

$$V\{cX\} = c^2 V\{X\}$$

If X_1 and X_2 are *independent* random variables, and if we write

$$\begin{aligned} E\{X_1\} &= a_1 & V\{X_1\} &= \sigma_1^2 \\ E\{X_2\} &= a_2 & V\{X_2\} &= \sigma_2^2 \end{aligned}$$

then we find that

$$E\{(X_1 - a_1)(X_2 - a_2)\} = E\{(X_1 - a_1)\}E\{X_2 - a_2\} = 0$$

The variance of a sum of two independent random variables X_1 and X_2 is

$$\begin{aligned} V\{X_1 + X_2\} &= E\{(X_1 + X_2 - a_1 - a_2)^2\} \\ &= E\{(X_1 - a_1)^2\} + 2E\{(X_1 - a_1)(X_2 - a_2)\} \\ &\quad + E\{(X_2 - a_2)^2\} \\ &= V\{X_1\} + V\{X_2\} = \sigma_1^2 + \sigma_2^2 \end{aligned} \quad (9.7-2)$$

It follows by induction that the variance of a sum of any number of independent random variables is the sum of their variances.

Exercises

9.7-1 Carry out the proof of the Chebyshev inequality for the discrete variable case.

9.7-2 The entropy is the expected value of $\log(1/p_i)$. Examine the corresponding formula for the variance.

$$\text{Ans. } \Sigma p_i \log^2(1/p_i) - H^2$$

9.8 The Law of Large Numbers

If we make a large number of *Bernoulli trials* (meaning that success or failure is a random variable X with values 1 or 0, respectively, and each trial is independent of all other trials), then we can ask for the number of successes *relative* to the total number of trials, and hope that this ratio will approach the probability of a single trial. This is simply hoping that the frequency of successes in a long run of trials will approach the probability of a success for a single trial. To get the formalism going, we consider the random variables X_i , which is 1 if the i th trial is successful and 0 if it is not. What will the average be? We set up the expression for the expectation of the average

$$E\left\{\frac{1}{n}(X_1 + X_2 + \dots + X_n)\right\} = \frac{1}{n} \sum_i E\{X_i\} = a$$

where a is the expected value of a single trial.

We next look at the variance of the average

$$\begin{aligned} V\left\{\frac{1}{n} \sum X_i\right\} &= E\left\{\frac{1}{n^2} \left[\sum (X_i - a)\right]^2\right\} \\ &= \frac{1}{n^2} E\left\{\sum (X_i - a)^2\right\} \end{aligned}$$

by (9.7-2), since the cross products drop out. But these are identical trials; hence

$$\begin{aligned} V\left\{\frac{1}{n} \sum X_i\right\} &= \frac{1}{n^2} \sum V\{X_i\} \\ &= \frac{n\sigma^2}{n^2} = \frac{\sigma^2}{n} \end{aligned}$$

We now apply Chebyshev's inequality (9.7-1),

$$\begin{aligned} \text{Prob} \left\{ \left| \frac{1}{n} \sum (X_i - a) \right| > \epsilon \right\} &\leq \frac{1}{\epsilon^2} E\left\{ \frac{1}{n^2} \sum (X_i - a)^2 \right\} \\ &\leq \frac{\sigma^2}{n\epsilon^2} \end{aligned} \quad (9.8-1)$$

Reversing things, we can write this as

$$\text{Prob} \left\{ \left| \frac{1}{n} \sum (X_i - a) \right| < \epsilon \right\} \geq 1 - \frac{\sigma^2}{n\epsilon^2}$$

Theorem: Weak law of large numbers. If $X_1, X_2, \dots, X_n$ are n independent identically distributed random variables with mean a and variance σ^2 , then given any $\epsilon > 0$, and $\delta > 0$ however small, there is an integer n_0 such for all $n > n_0$,

$$a - \epsilon \leq \frac{1}{n} \{X_1 + X_2 + \dots + X_n\} \leq a + \epsilon \quad (9.8-2)$$

with probability greater than $1 - \delta$. The theorem states that the average can be brought, *probably*, as close as you please to the expected value a .

The proof follows whenever we pick n_0 such that [from (9.8-1)]

$$\frac{\sigma^2}{n_0 \epsilon^2} < \delta$$

or

$$n_0 > \frac{\sigma^2}{\delta \epsilon^2}$$

Since this theorem is widely misunderstood, we pause to comment on it. In the first place, since they are Bernoulli trials, that means that the trials are independent—the *trials have no memory* of what has happened in the past. Therefore, the very concept that the law of large numbers means that in time any run of bad luck will later be “compensated for” by a run of good luck must be false. What happens is that any fluctuation from the average will be “swamped” (probably) in the long run, not that there will be compensation.

Furthermore, it does not mean that the sum comes close to the expected value, which is, of course, np . For example, in 10 tosses of a fair coin, you expect 5 heads, but a value as large as 8 is not surprising. In the million tosses we do not expect to be as close as 3 (as we were in the 10 tosses with 8 heads) from the expected value, but it is *percentagewise* that we come closer and closer—probably! The misunderstanding comes when you do not watch where the n is—it is in the denominator and when it is there, probably the differences between the number of actual successes and the expected number of successes *divided by the n* will decrease as n gets larger and larger. But if you talk about the difference between the number of observed successes and the expected number of successes, it can grow indefinitely large with larger and larger n . In this it resembles the earlier Stirling approximation (Table 9.2-1).

Because of its central role in the proof of Shannon’s theorem, let us illustrate the theorem by an experiment. Let the probability of an event (success) be

$$p = \frac{1}{10}$$

In N Bernoulli trials, what is the probability of k successes? The k successes can occur in any k of the N trials, namely

$$C(N, k)$$

different ways; therefore, the probability for the k successes (and, of

course, $N - k$ failures) is

$$C(N, k)p^k(1 - p)^{N-k} \quad (9.8-3)$$

The expected number of successes is

$$pN$$




Figure 9.8-1 Law-of-large-numbers example

In Figure 9.8-1 we have plotted horizontally k/N so that the distribution of (9.8-3) is centered at the same place, namely p , the expected value of a single trial.

To compensate for the compression of the horizontal axis, we plot N times the function (9.8-3), namely

$$NC(N, k)p^k(1 - p)^{N-k} \quad (9.8-4)$$

Thus the area under the curve is held constant. Figure 9.8-1 gives the distribution for $N = 10, 30, 100$. We see the peaking up of the distribution near $p = \frac{1}{10}$ and the falling away of the probability at places far

away from the neighborhood of $p = \frac{1}{10}$. In Figure 9.8-2 we have shrunk the vertical axis and give the cases $N = 100, 300, 1000$. (Only a few points are plotted and the curve for $N = 1000$ would, in fact, be very smooth indeed if we plotted all the points.)




Figure 9.8-2 Law-of-large-numbers example

Thus we see how, as the number of trials N increases, the distribution centers around the expected value p , and the variance decreases like $1/\sqrt{N}$. At $N = 1000$, for example, a deviation of k/N of as much as 0.05 is very unlikely indeed. The probability of being at all far from p can be made as small as you please by picking a sufficiently large N .

We are going to use this theorem when we encode a message in a very large extension. We can expect that the distribution of the difference between the number of actual errors in the transmission of the block of symbols and the expected number of errors, when divided by N , will get to be a very narrow spike indeed—provided that the N is large enough. Always probably!

We can use the law of large numbers in an interesting way to explain more about the nature of n -dimensional space. Suppose that we consider the vector from the origin to the point with coordinates ± 1 chosen at random for each of its coordinates. Now let us choose another such random diagonal line. What is the angle between them? Each vector is of length $\sqrt{n}$. The cosine of the angle is therefore given by

$$\cos \theta = \frac{\sum_{k=1}^n (\pm 1)(\pm 1)}{\sqrt{n}\sqrt{n}}$$

But uniform random products $(\pm 1)(\pm 1)$ are ± 1 uniformly; hence

$$\cos \theta = \frac{\sum_{k=1}^n (\pm 1)}{n}$$

But “almost surely” for large n the sum divided by n (the average of ± 1 ’s) is near zero; hence, as in Section 9.5, the angle approaches $\pi/2$, a right angle.

But there are 2^n such random diagonal lines; hence in n -dimensional space, for sufficiently large n , *there are almost 2^n almost perpendicular lines!* Now 2^n is very much larger than n ; for $n = 10$, we have $2^{10} = 1024$. Thus even at 10 dimensions n -dimensional space is very “vast” indeed!

Exercise

9.8-1 Make similar plots for $p = \frac{1}{2}$, $N = 5, 10, 20$.


9.9 Other Metrics

The preceding material has been developed using the Pythagorean distance, L_2 :

$$d^2 = \sum_{i=1}^n (x_i - y_i)^2$$

since that is the familiar distance measure. But coding theory often uses the Hamming distance, L_1 :

$$d = \sum_{i=1}^n |x_i - y_i|$$



Circle in L_1

Figure 9.9-1


With this distance function the circle (in two dimensions) looks like Figure 9.9-1.

The other main distance is the limit of the L_n distance as $n \longrightarrow \infty$:

$$d_n = \sqrt[n]{\sum |x_i - y_i|^n}$$

In the limit we see that the tendency of the high power in the sum is to select out the largest term and ignore the rest; hence we have the Chebyshev, L_∞ distance:

$$d = \lim_{n \rightarrow \infty} \sqrt[n]{\sum |x_i - y_i|^n} = \max_i |x_i - y_i|$$

Circle in L_∞ **Figure 9.9-2**

In this distance the circle becomes as shown in Figure 9.9-2.

Many of the phenomena of this chapter apply independent of the distance function used, but it is evident that the four-circle paradox changes as the distance function changes.

It is important to realize that the nature of n -dimensional space depends on the distance function used, but that the “vastness” of the n -dimensional space remains.

Shannon's Main Theorem

10.1 Introduction

In Chapter 6 we introduced the concept of entropy, and in Section 6.8 proved that there are encoding methods that come arbitrarily close to the entropy of the input. In particular, we showed [equation (6.8-2)] that by a suitably large extension of length n we can achieve the inequalities for the average code length L_n :

$$H_r(S) \leq \frac{L_n}{n} < H_r(S) + \frac{1}{n}$$

Thus the encoding process can be done arbitrarily well, and the entropy $H_r(S)$ provides the limit on what can possibly be done.

We next looked at the channel itself, and defined a channel capacity C , equation (8.1-1). In the present chapter we prove the fundamental result that even in the presence of noise, we can get arbitrarily close to the channel capacity. More specifically, we can get arbitrarily close to the rate, and keep the reliability of sending messages through the channel arbitrarily close to 1.

We will prove the theorem for the binary symmetric channel, since in this case the proof is easiest to understand, and it also provides the most common model used in practice. We will merely indicate how the underlying ideas can be extended to prove the theorem for more general channels. The rigorous proof in the general case is complex in

all its details and does not shed much more light on why the theorem is true in general.

10.2 Decision Rules

Suppose that you are at the receiving end of the channel and have received the symbol b_j . Which a_i should you assume was sent? That, of course, depends on the channel, that is, on the $P(b_j | a_i)$ and also on the source probabilities $p(a_i) = p_i$.

To get a feeling for what is involved, consider the particular channel where we have the channel matrix $P_{i,j} = P(b_j | a_i)$:

$$\begin{array}{ccc}
 & b_1 & b_2 & b_3 \\
 & & \uparrow & \\
 \begin{array}{l} a_1 \\ a_2 \\ a_3 \end{array} & \longrightarrow & \begin{pmatrix} 0.5 & 0.3 & 0.2 \\ 0.2 & 0.3 & 0.5 \\ 0.3 & 0.3 & 0.4 \end{pmatrix} & (10.2-1)
 \end{array}$$

Let $d(b_j)$ be the result of the *decision rule* you will use when b_j is received. Even if you have a uniform input and all the $p(a_i) = \frac{1}{3}$ for the source alphabet, the three sets of rules for this channel

<i>Rule 1</i>	<i>Rule 2</i>	<i>Rule 3</i>	
$d(b_1) = a_1$	$d(b_1) = a_1$	$d(b_1) = a_1$	
$d(b_2) = a_1$	$d(b_2) = a_2$	$d(b_2) = a_3$	(10.2-2)
$d(b_3) = a_2$	$d(b_3) = a_2$	$d(b_3) = a_2$	

are each defensible—arguments can be made in favor of each. Note that we are here considering one symbol at a time, not runs of symbols.

In the general case, there are s rules (one for each received symbol) each of which may have any one of the q input symbols, so there are in total

$$q^s$$

possible decision rules for picking an a_i when a b_j is received.

To settle the question of which set of decision rules to adopt, we shall invoke the widely used *maximum-likelihood rule* of statistics. This rule says that you should “take the symbol that is most probable, given

that you have observed what you have.” In mathematical terms we require (given equal frequency of the use of the input symbols, which is the reasonable way to use a binary symmetric channel)

$$d(b_j) = a^*$$

where a^* is defined by the conditions

$$P(a^* | b_j) \geq P(a_i | b_j) \quad \text{for all } i \quad (10.2-3)$$

In words, given that you have received b_j , no other a_i is more probable than is a^* . These are not the channel probabilities. To get to them we use Bayes’ theorem on each side of this inequality:

$$\frac{P(b_j | a^*)p(a^*)}{p(b_j)} \geq \frac{P(b_j | a_i)p(a_i)}{p(b_j)}$$

Under the assumption of the uniform probability of the input symbols, we have $p(a^*) = p(a_i)$, and we get

$$P(b_j | a^*) \geq P(b_j | a_i) \quad (10.2-4)$$

for all i . Thus we now have the maximum-likelihood condition (10.2-3) in the form used in defining the channel.

We can use this rule *even when* the source alphabet is not equally likely. This is not optimal in general, but—like giving away the optimality of Huffman encoding and using Shannon–Fano encoding—what we give away will not cost us too much in the long run. We will still be able to obtain the desired result. This indicates a certain robustness of the final theorem—it is not a delicately true theorem. We can deviate fairly far from the best possible at various stages and still find that the theorem can be proved.

The maximum-likelihood rule is still not unique, as you can see from the example (10.2-2). In the example for the channel (10.2-1), clearly for the maximum-likelihood rule you must take

$$d(b_1) = a_1$$

But given b_2 , you can pick any one of the following three:

$$d(b_2) = a_1$$

$$d(b_2) = a_2$$

as our rule. You will have finally, by the maximum-likelihood rule,

$$d(b_3) = a_2$$

Thus there are three possible maximum-likelihood rules (10.2-2) for this channel.

What is the probability of error using this channel and this decision rule? If we write the probability of an error as $P(E | b_j)$, then we have, since $P(a^* | b_j)$ is the probability that the symbol a^* we pick was the source of the received b_j ,

$$P(E | b_j) = 1 - P[d(b_j) | b_j] \quad (10.2-5)$$

as the probability of making an error given that we have received b_j . The *average error* P_E is given by

$$P_E = \sum_B P(E | b_j)p(b_j)$$

This may be written as, using (10.2-5),

$$P_E = \sum_B p(b) - \sum_B P(a^* | b)p(b)$$

Using Bayes' theorem (Section 7.3), we can reverse the conditional probabilities to get to the channel probabilities [remember that all $p(a) = 1/q$]

$$\begin{aligned} P_E &= 1 - \sum_B P(b | a^*)p(a) \\ &= 1 - \frac{1}{q} \sum_B P(b | a^*) \end{aligned} \quad (10.2-6)$$

Let us calculate this for the channel (using the underlined rule). From


$$\frac{1}{3} \begin{pmatrix} \underline{0.5} & 0.3 & 0.2 \\ 0.2 & 0.3 & \underline{0.5} \\ 0.3 & \underline{0.3} & 0.4 \end{pmatrix}$$

and using formula (10.2-6) we have

$$\begin{aligned} P_E &= 1 - \frac{1}{3}(0.5 + 0.3 + 0.5) \\ &= 1 - \frac{1}{3}(1.3) = 0.5666 \dots \end{aligned}$$

10.3 The Binary Symmetric Channel

We now illustrate Shannon's main theorem by showing how the error in the channel can be reduced by using an error-correcting code for the binary symmetric channel. For a single binary digit we have the standard diagram



Note that P = the probability of correct transmission. We naturally assume that $Q < \frac{1}{2}$, since if $Q \geq \frac{1}{2}$, we merely reverse the 0 and 1 at B . For $Q = \frac{1}{2}$ we have all noise and the channel capacity is zero.

For the n th extension of the channel (Section 8.2), we send n binary digits. For example, for the third extension (octal code) we have for single-error-correcting code (maximum-likelihood detection):

Messages	Unused <i>words</i>	Received
000	001	000
	010	001
111	011	010
	100	011
	101	100
	110	101
		110
		111

No error	P^3
1 error	$3P^2Q$ (correctable)
2 errors	$3PQ^2$
3 errors	Q^3

$$P_E \approx Q^3 + 3PQ^2 = Q^2(Q + 3P)$$

If $P = 0.99$ (99% reliable), we get

We see that the Hamming distance is in this case the maximum-likelihood detector. In general, for "white noise" the minimum Hamming distance provides the maximum-likelihood detection. In the case of two message points of the same minimum distance from the received symbol, we simply give up, or else toss a coin and guess (we will be right half the time). In the case of more than two message points the same distance away, we can still guess.

For the n th extension of the binary symmetric channel, we see that we have a uniform channel. How we use the channel in practice may be another matter.

10.4 Random Encoding

We are now ready to prove Shannon's main theorem, which states that there exists a method of encoding messages that is arbitrarily close to capacity and arbitrarily reliable in transmission (storage). We will prove the theorem *only* for the binary symmetric channel, both because that is the channel of main interest and because for this restricted case the proof can be made intelligible to unsophisticated beginners.

We will use block codes of n binary digits. These are simply the n th extension (Section 6.8) of the binary code. There are 2^n possible strings of n binary digits, but if we are to have protection against the noise of the channel we cannot use all of them. We must keep the encoded messages far apart (in the Hamming distance) so that the noise of the channel cannot cause errors (for almost all messages).

First, how many messages do we need to be near capacity? If we have M equally likely message [which is the way to use a uniform channel (Section 8.2)], then the information per symbol a_i is, since

$$p = \frac{1}{M} \tag{10.4-1}$$

$$I(a_i) = \ln \frac{1}{p} = \ln \frac{1}{1/M} = \ln M$$

Now we know from Section 8.5 that the capacity of a binary symmetric channel is

$$C = 1 - H(P) = 1 - H(Q)$$

for one binary digit, and hence for the n th extension we will have the

capacity

$$nC = n[1 - H(Q)]$$

(by Section 8.3). To be near capacity we must have

$$I(a_i) = \ln M = n(C - \varepsilon_1)$$

where ε_1 is a small positive number. This may be written

$$M = 2^{n(C - \varepsilon_1)} = \frac{2^{nC}}{2^{n\varepsilon_1}} \quad (10.4-2)$$

Looked at this way, we see that by increasing n the denominator can be made arbitrarily large, so that we are taking a small fraction of all possible messages of length n bits that can get through at capacity C —provided that n is made large enough!

Second, how do we design the code book so that all messages are equally likely and are far apart? Frankly, at this point, we do not know how to do this in general and we resort to *random encoding*. We know from the end of Chapter 9 that in n -dimensional space, random points are far apart, provided that n is large. Judged from the center of the cube, two random vertices will almost certainly be almost a quadrant apart. Thus, making random choices for the encodings is not as foolish as at first it may seem.

To get the encoding of the first message in the code book, we toss a well-balanced coin n times and record the corresponding sequence of 0's and 1's as the encoded message. We do the same for the second message in the code book. And so on through the whole M messages in the code book. There will be exactly nM tosses of the coin; hence we will get one particular code book from the ensemble of 2^{nM} possible code books.

An alternative and equivalent method of getting the code book is to write all possible 2^n binary sequences on slips of paper, put them in a container, stir, select the first one as the encoding of the first entry in the code book, return this slip, stir, draw the second slip as the second encoded line of the code book, and so on, until all M messages in the code book have been assigned. These are equivalent methods, and both run the risk of having the same encoded sequence emerge for two different messages. For large n , however, this is very, very unlikely, as we will have to prove if we are to show that the error rate can be made arbitrarily small.

We picture these sequences as message points in an n -dimensional Hamming space (Section 9.5). The probability of an error in any one position of the block of letters is $Q = 1 - P$, and we expect nQ errors. We consider a sphere about the chosen message point with a radius of size nQ . Now consider slightly increasing this sphere by a *relative* amount ϵ_2 . The radius is now (see Figure 10.4-1)

$$r = (Q + \epsilon_2)n \quad (10.4-3)$$




Figure 10.4-1 Sending end

(r is the radius and is *not* the number of symbols or the radix, and ϵ_2 is so small that $Q + \epsilon_2 < \frac{1}{2}$). The actual size of ϵ_2 with respect to ϵ_1 will be determined later.

By the law of large numbers (Section 9.8) there is, for sufficiently large n , an arbitrarily small probability that the received message will fall outside this new sphere. (Since almost all the volume of an n -dimensional sphere is near the surface, it is natural that we are mainly concerned with a shell of thickness $2n\epsilon_2$ centered about the distance nQ .) Thus if we use maximum-likelihood detection, *and* if we can prove that almost surely no two spheres overlap, then we will have the result that you, the receiver, can almost surely know what was sent.

For the receiver we must invert the picture, and draw it centered about the received message symbol b_j (see Figure 10.4-2). Drawing the same-size spheres about the b_j , we find the same probabilities for the a_i being in the spheres as we had for the b_j to be in the spheres about a_i . We need to set up a formula for the error that you cannot decode the message correctly.

You will be in error if the proper a_i is not in the sphere $S(r)$ about the received b_j , and you will also be in error if there is more than one possible a_i in the sphere since then you will not know how to decode. (We will decline to guess.) The following table gives the situation.




Figure 10.4-2 Receiving end

Cases	Correct a_i in $S(r)$	Other a_i in $S(r)$	Error
1	Yes	Yes	Error
2	Yes	No	Correct
3	No	Yes	Error
4	No	No	Error

From this table we set up the probability P_E of an error. The first term covers cases 3 and 4 (too much noise), while the second term covers case 1 (code points too close together). We have

$$P_E = P\{a_i \notin S(r)\} + P\{a_i \in S(r)\}P\{\text{at least one other code word is in } S(r)\} \quad (10.4-4)$$

where $\in$ means “is contained in” and $\notin$ means “is not contained in.” Since $P\{a_i \in S(r)\} \leq 1$ we can simplify (10.4-4) slightly by dropping this factor:

$$P_E \leq P\{a_i \notin S(r)\} + P\{\text{at least one other code word is in } S(r)\} \quad (10.4-5)$$

We now look at the second term on the right. The words “at least one other . . .” means “one or more code points in sphere in addition to the given a_i .” Now given two events E_1 and E_2 it is clear that

$$P(E_1 \text{ and/or } E_2) = P(E_1) + P(E_2) - P(E_1 E_2)$$

since a point in $E_1 E_2$ has been counted in both E_1 and E_2 and must be subtracted out. Since all probabilities are nonnegative,

$$P(E_1 \text{ and/or } E_2) \leq P(E_1) + P(E_2)$$

It is evident from this simple case that (summation over $A - a_i$ means all of A except a_i)

$$P\{\text{at least one other code word is in } S(r)\} \leq \sum_{A-a_i} P\{a \in S(r)\}$$

where the summation is over the $M - 1$ code words not sent. We now have, for (10.4-5),

$$P_E \leq P\{a_i \notin S(r)\} + \sum_{A-a_i} P\{a \in S(r)\} \quad (10.4-6)$$

In words, the first term is the probability that the sent word is not within the sphere about the received word (equivalently, that the received word is not within the sphere about the sent word). The second term is the sum over all the terms not sent, any one of which might be in the sphere of radius r , which is, *relatively* with respect to n , slightly larger than the Hamming distance corresponding to the expected number of errors. We hope that by using randomness the message points will be spread out and that there will be only a very small chance that any other symbol will be in the sphere. This, in turn, rests on the vastness of n -dimensional space; it is very unlikely that any two random points will be near to each other in a sufficiently high-dimensional space.

The first term of (10.4-6), which depends only on the single symbol a_i , can be made less (by the law of large numbers) than any preassigned positive number $\delta = \delta(\epsilon_2, n)$ simply by making the block length n so large that the extended sphere of radius $r = n(Q + \epsilon_2)$ covers the distances for all but the arbitrarily small number of possible cases. We have, therefore, for (10.4-5),

$$P_E \leq \delta + \sum_{A-a_i} P\{a \in S(r)\} \quad (10.4-7)$$

One way of looking at this equation is that the first term contributes to the error rate depending only on the particular code symbol sent; the second term depends on all the other symbols in the code, and is related to the rate of signaling, since it contains the $M - 1$ possible code symbols other than the correct one.

10.5 Average Random Code

All that we have done so far seems reasonable, except possibly the choice of a random code (although we hope that it has at least been made plausible by the earlier use of hash coding—Sections 5.11 to 5.14 and much of Chapter 9). It is the sort of thing one would think of in the normal course of research. But the next idea of Shannon's, that we should *average over all random codes* (but what else is there to do?) is not one of those things that comes to mind easily. Of course, once the idea is suggested, it is something to try.

We therefore now average over all codes. We use a wavy line for the average over all the code books. The code words come from the urn one at a time (or else are made up bit by bit by the random toss of a coin) and we chose [equation (10.4-2)] $M = 2^{n(C-\epsilon_1)}$ words. There are M symbols of n bits, so there are in all (as we earlier observed)

$$2^{nM}$$

possible code books we could use. With the random selection all these code books are equally likely; namely, each occurs with probability

$$\frac{1}{2^{nM}}$$

The average, weighted by this probability, gives the probable error

$$\widetilde{P_E}$$

The δ is a fixed constant for all the code books and does *not* depend on the whole code. Hence when we average over all the codes, equation (10.4-7) becomes (since all code points $a \neq a_i$ are equivalent *on the average*)

$$\begin{aligned} \widetilde{P_E} &\leq \delta + (M - 1) \widetilde{P\{a \in S(r)\}} \\ &\leq \delta + M \widetilde{P\{a \in S(r)\}} \quad (a \neq a_i) \end{aligned} \quad (10.5-1)$$

To evaluate $\widetilde{P\{a \in S(r)\}}$ ($a \neq a_i$), we fix our attention on any particular line of the code book, not that of a_i . Since the code books were made up randomly, the encoding of the line we are looking at will vary uniformly over all the space of sequences of n binary digits. What is the probability that this point will be in the sphere $S(r)$? This prob-

ability must be the ratio of the volume of the sphere to the total volume. The total volume is, of course, 2^n . Let $N(r)$ be the number of points inside the sphere $S(r)$. Then

$$\widetilde{P\{a \in S(r)\}} = \frac{N(r)}{2^n} \quad (a \neq a_i) \quad (10.5-2)$$

But for our binary symmetric channel, this is

$$\begin{aligned} N(r) &= 1 + C(n, 1) + C(n, 2) + \dots + C(n, r) \\ &= \sum_{k=0}^r C(n, k) \end{aligned}$$

Recall that [from (10.4-3)]

$$r = (Q + \epsilon_2)n \quad \text{and} \quad Q + \epsilon_2 < \frac{1}{2}$$

We can apply the binomial inequality of equation (9.4-6) using $Q + \epsilon_2 = \lambda$, to get

$$N(r) \leq 2^{nH(\lambda)}$$

Hence equation (10.5-2) becomes

$$\widetilde{P\{a \in S(r)\}} \leq 2^{-n[1-H(\lambda)]} \quad (a \neq a_i) \quad (10.5-3)$$

Assembling things, we have for equation (10.5-1),

$$\widetilde{P_E} \leq \delta + M2^{-n[1-H(\lambda)]} \quad (10.5-4)$$

For a binary symmetric channel, equation (8.5-3), we have

$$C = 1 - H(P) = 1 - H(Q)$$

Therefore, the exponent in equation (10.5-4) may be written as

$$\begin{aligned} 1 - H(\lambda) &= 1 - H(Q + \epsilon_2) \\ &= 1 - [H(Q) + H(Q)]_{\text{convex}} - H(Q + \epsilon_2) \\ &= C - [H(Q + \epsilon_2) - H(Q)] \end{aligned} \quad (10.5-5)$$

But the entropy function is convex downward and can be bounded at

any point Q by fitting a tangent line at Q :

$$H(Q + \epsilon_2) \leq H(Q) + \epsilon_2 \left. \frac{dH}{dQ} \right|_Q,$$

where (since $0 \leq Q \leq \frac{1}{2}$)

$$\frac{dH}{dQ} = \log\left(\frac{1}{Q}\right) - \log\left(\frac{1}{1-Q}\right) = \log\left(\frac{1-Q}{Q}\right) > 0$$

Therefore, (10.5-5) becomes

$$1 - H(\lambda) \geq C - \epsilon_2 \log \frac{(1-Q)}{Q} = C - \epsilon_3$$

where ϵ_3 is defined as

$$\epsilon_3 = \epsilon_2 \log \left(\frac{1-Q}{Q} \right)$$

Now (10.5-4) becomes

$$\widetilde{P}_E \leq \delta + M 2^{-n(C-\epsilon_3)}$$

From (10.4-2),

$$M = 2^{n(C-\epsilon_1)}$$

and

$$\widetilde{P}_E \leq \delta + 2^{n(C-\epsilon_1)} 2^{-n(C-\epsilon_3)} = \delta + 2^{-n(\epsilon_1-\epsilon_3)}$$

Thus as long as we pick ϵ_2 small enough so that

$$\epsilon_1 - \epsilon_3 = \epsilon_1 - \epsilon_2 \log \left(\frac{1-Q}{Q} \right) > 0$$

then the average error

$$\widetilde{P}_E$$

can be made arbitrarily small by choosing n sufficiently large.

We have now shown that the average code meets our requirements; hence there must be at least one code that does. Thus we have proved Shannon's result:

Theorem. There exist codes that are arbitrarily reliable and which can signal at rates arbitrarily close to the channel capacity of the binary symmetric channel.

What do we learn from the proof? We learn that theoretically suitable codes exist, but we also learn a number of other things. We see from the use of the central limit theorem that the messages will have to be so long that we are, on the average, correcting many errors in each message. We also see that we must send very long messages to get an efficient code. But long messages require a lot of information to be gathered before anything can be sent, meaning that in many circumstances nothing will be sent until enough has been gathered to do the encoding of all of it. Further, we see that the code book, being random and with large M , must be very large indeed! There can be no short description, as there is for the Hamming codes, for instance, telling how to encode and decode a message—it is random and “random” means that no short description is possible.

We see that it is a theoretical result, not a practical result. Yet this does not mean that it is worthless. We see the nature of the problem of efficient encoding—we must have long messages and we must correct many erroneous bits in each message if we are to make very efficient use of the channel. The theorem sets bounds but does not suggest much other than that very good codes must be very long, often rather impractical ones.

10.6 The General Case

The proof of Shannon's main theorem in the general case goes along the lines of the special case of the binary symmetric channel. The differences in the proof arise from the question of finding an equivalent metric to the Hamming distance (which was appropriate *only* for a channel with white noise). We also have the problem of counting the number of possible messages that lie within a “sphere” with the new distance function. Finally, we cannot use the simple formula for the capacity of the channel but must go back to fundamentals.

The metric we need is the probability of errors. If the “noise” is not white, the spheres become ovals with their longer directions in the

directions of the more probable errors, and the shorter in the directions of the less probable errors.

Once a suitable metric is found, we pick the “ovals of equal probability radius” about the received symbols, increase them as before, so that by the law of large numbers when n is sufficiently large, almost certainly the original symbol is inside. We choose the number of messages M so that we are close to capacity. We must also use a sufficiently large n so there is arbitrarily low probability of an overlap of the ovals. Otherwise, the proof is the same as for the binary symmetric channel.

10.7 The Fano Bound

To prove the converse of Shannon’s theorem, that we cannot signal arbitrarily reliably at a rate higher than the channel capacity C , we need an important bound due to Fano. For convenience we use the notation

$$\bar{P}_E = \sum_B P(a^*, b) = 1 - P_E$$

where, as before [equation (10.2-5)],

$$P_E = \sum_{B, A=a^*} P(a, b)$$

Consider the expression

$$\begin{aligned} H(P_E) + P_E \log(q - 1) &= P_E \log\left(\frac{q - 1}{P_E}\right) + \bar{P}_E \log\left(\frac{1}{\bar{P}_E}\right) \\ &= \sum_{B, A=a^*} P(a, b) \log\left(\frac{q - 1}{P_E}\right) \\ &\quad + \sum_B P(a^*, b) \log\left(\frac{1}{\bar{P}_E}\right) \end{aligned} \quad (10.7-1)$$

The equivocation $H(A | B)$ from equation (7.5-4) is of a similar form,

$$\begin{aligned} H(A | B) &= \sum_{B, A=a^*} P(a, b) \log\left[\frac{1}{P(a | b)}\right] \\ &\quad + \sum_B P(a^*, b) \log\left[\frac{1}{P(a^* | b)}\right] \end{aligned}$$

Subtracting equation (10.7-1) from this, we get

$$\begin{aligned} H(A | B) - H(P_E) - P_E \log(q - 1) \\ = \sum_{B, A-a^*} P(a, b) \log \left[\frac{P_E}{(q-1)P(a | b)} \right] + \sum_B P(a^*, b) \log \left[\frac{\bar{P}_E}{P(a^* | b)} \right] \end{aligned}$$

It is difficult to apply the Gibbs inequality here, so we go back to basics. To apply the inequality

$$\log_e x \leq x - 1$$

we must convert this inequality to $\log_2$ by writing it as

$$\frac{\log x}{\log e} < x - 1$$

Apply this to *each* sum. The $1/\log e$ will factor out of each term on the right-hand side. Ignoring the log factor, we get for the right-hand side,

$$\begin{aligned} &\leq \sum_{B, A-a^*} P(a, b) \left[\frac{P_E}{(q-1)P(a | b)} - 1 \right] + \sum_B P(a^*, b) \left[\frac{\bar{P}_E}{P(a^* | b)} - 1 \right] \\ &= \frac{P_E}{q-1} \sum_{B, A-a^*} p(b) - P_E + \bar{P}_E \sum_B p(b) - \bar{P}_E \\ &= \frac{P_E}{q-1} (q-1) \sum_B p(b) + \bar{P}_E \sum_B p(b) - (P_E + \bar{P}_E) \\ &= P_E + \bar{P}_E - (P_E + \bar{P}_E) = 0 \end{aligned}$$

Hence we have the important Fano inequality

$$H(A | B) \leq H(P_E) + P_E \log(q - 1) \quad (10.7-2)$$

Note first that we have *not* used any decision rule, although the results depend on it. Second, when do we have equality? Equality occurs when

$$\log_e x = x - 1$$

namely when $x = 1$. This corresponds, since we used it in both sum-

mations, to both

$$P(a | b) = \frac{P_E}{q - 1} \quad \text{for all } b \text{ and } a \neq a^*$$

and

$$P(a^* | b) = 1 - P_E \quad \text{for all } b$$

The second, however, is included in the first, since

$$\sum_A P(a | b) = 1 \quad \text{for all } b$$

10.8 The Converse of Shannon's Theorem

To prove the converse of Shannon's theorem, we use Fano's inequality of equation (10.7-2). We want to prove that the probability of error cannot be kept arbitrarily small if we exceed the channel capacity. That is, we cannot make

$$M = 2^{n(C+\epsilon)} \quad (\epsilon > 0)$$

$$P_E \rightarrow 0 \text{ as } n \rightarrow \infty$$

for M equally probable original messages. Suppose we try. Then for the n th extension of the alphabet, we have that the change in the entropy

$$H(A^n) - H(A^n | B^n)$$

is less than or equal to nC (the channel capacity measured in blocks of n symbols is *defined* as the upper bound). But $p = 1/M$, so

$$\begin{aligned} H(A^n) &= \sum \frac{1}{M} \log M \\ &= \log M = n(C + \epsilon) \end{aligned}$$

We have

$$H(A^n) - H(A^n | B^n) \leq nC$$

Rearranging, we have

$$n(C + \epsilon) - nC = n\epsilon \leq H(A^n | B^n)$$

Applying Fano's result (10.7-2), we obtain

$$n\epsilon \leq H(A^n | B^n) \leq H(P_E) + P_E \log(q - 1)$$

Now $H(P_E) \leq 1$ (from Section 8.5), so using $q - 1 < q = M$, we have

$$n\epsilon \leq 1 + P_E(nC + n\epsilon)$$

or

$$\frac{n\epsilon - 1}{nC + n\epsilon} \leq P_E$$

We can therefore write

$$P_E \geq \frac{n\epsilon - 1}{nC + n\epsilon} \geq \frac{\epsilon - 1/n}{C + \epsilon} \longrightarrow \frac{\epsilon}{C + \epsilon}$$

which is independent of n . As $n \rightarrow \infty$, P_E is bounded away from 0. Hence the error cannot be made to vanish if the rate exceeds the channel capacity.

As an illustration, suppose that we tried to send at a rate of twice the channel capacity. If we sent every other bit through the channel, we could get them through arbitrarily reliably. For the missing bits, the receiver tosses a coin and is right half the time; thus almost three-fourths of the message is correct and one-fourth erroneous.

Algebraic Coding Theory

11.1 Introduction

Shannon's main theorem, which used very long block codes, shows that we can signal at a rate arbitrarily close to the maximum rate and with arbitrarily few errors. The proof is based on random encodings which use very long blocks, and in any one of the long words there will probably be many errors to be corrected at the receiver, which means that to be efficient we must work at high error-correction rates. In Chapter 3 we showed how to design single-error-correcting codes as well as double-error-detecting codes. The theory of multiple-error-correcting codes is very elaborate, so we merely sketch the main ideas involved in how to design and implement some multiple-error-correcting codes.

The idea of a random code is appealing until you think about the terminal equipment. If the code words are long, then the encoding book must be very large, since by definition "random" means that there is no pattern and each encoding must be looked up because there can be no formula for computing the encoded message from the source information as there are for regular codes. The decoding book must be even larger, since although only M selected sequences can be sent, every one of the 2^n possible sequences of n bits (in a binary channel) must have an entry for decoding the received message. Otherwise, we must do a very large amount of computing to find the proper message to associate with the received message. Thus, in practice, large random codes are hopeless. (But remember the discussion of hash coding in Sections 5.11 to 5.14.)

We have, therefore, to look for highly structured codes, and to do this we need the theory of the algebra of finite fields (Galois fields). As noted above, the topic of algebraic codes is vast, and we can only touch the surface to indicate some of the ideas involved in the higher-error-correcting codes. This is not a book for experts who are going to design coding systems ideally fitted to complex situations; the book was organized for those who need some acquaintance with what can be done, and hopefully if a situation arises that requires more theory the reader will at least recognize the possibilities and go to other sources for more details (e.g., Refs. [B1], [B3], [B4], [C], [G], [Gu], [L1], [L2], [MS], [Mc], [P], [PW], [V], [W], [WJ]). However, it is only fair to warn the reader that in practice we seldom know the noise well enough to justify using many of the more elaborate codes that have been developed (and usually based on pure white noise). Occasionally, we do.

We will emphasize binary block codes, both because they are most widely used, and because the base r codes are somewhat messier in notation, although the ideas needed are usually not significantly different. We also deliberately review the Hamming codes of Chapter 3 so that you will remember them this time! You will see them in a very different light.

11.2 Error-Detecting Parity Codes Revisited

The single-error-detecting parity codes of Section 2.2 used a single parity check over all the positions. If we write a 1 for each *position* that is checked and a 0 otherwise, we get the matrix (in this case the trivial $m \times 1$ matrix)

$$M = (1, 1, 1, \dots, 1)$$

This matrix says that the parity check involves every position.

Consider an encoded symbol which is now labeled c ; it is a vector of n binary digits, say,

$$c = (0, 1, 0, \dots, 1)$$

Let the *transpose* of c , written c^T , be the corresponding column vector. We have (using mod 2 arithmetic, of course) the matrix product

$$Mc^T = 0 \tag{11.2-1}$$

for every encoded word to be sent. It is important to see why this is so. When inspected, equation (11.2-1) is simply a statement of what

the parity check is. The 1 digits in the matrix M select the positions of c that enter into the parity check.

Suppose next that a *single* error is made in sending the code symbol c . The error changes one binary digit. Recalling that we do our arithmetic modulo 2, the received word is

$$c + e$$

where e is the error vector with all 0's except a 1 in the position of the error.

Consider, finally, the expression (using mod 2 arithmetic)

$$\begin{aligned} M(c + e)^T &= Mc^T + Me^T \\ &= Me^T \neq 0 \quad (\text{syndrome for a single error}) \end{aligned}$$

The *syndrome* of the erroneous message (the result of applying the parity checks) in this case is a 1 whenever there is an odd number of errors and a 0 whenever there is an even number of errors, including no errors.

11.3 Hamming Codes Revisited

In Section 3.4 we examined a special case of single-error-correcting codes. These Hamming codes used the parity checks

$$\begin{array}{ll} 1 & 1, 3, 5, 7, 9, 11, 13, 15, \dots \\ 2 & 2, 3, 6, 7, 10, 11, 14, 15, \dots \\ 3 & 4, 5, 6, 7, 12, 13, 14, 15, \dots \\ 4 & 8, 9, 10, 11, 12, 13, 14, 15, \dots \\ & \text{etc.} \end{array}$$

In the further special case of $n = 7$, we write the matrix

$$M = \begin{pmatrix} 0 & 0 & 0 & 1 & 1 & 1 & 1 \\ 0 & 1 & 1 & 0 & 0 & 1 & 1 \\ 1 & 0 & 1 & 0 & 1 & 0 & 1 \end{pmatrix} \begin{array}{l} \leftarrow \text{third parity check} \\ \leftarrow \text{second parity check} \\ \leftarrow \text{first parity check} \end{array}$$

where we have put the parity checks in sequence *from the bottom up*. Each position in any parity check has a 1 in the corresponding row of the matrix.

Again let c be a code word. We again have *for all code words*,

$$Mc^T = 0 \tag{11.3-1}$$

since this is the definition of an encoded word. In technical jargon the matrix M annihilates every properly encoded word c . A single error e will change the original message I sent to the received message you get. It will now be

$$c + e$$

and we will have again

$$\begin{aligned} M(c + e)^T &= Mc^T + Me^T \\ &= Me^T = \text{the syndrome} \end{aligned} \quad (11.3-2)$$

Thus the syndrome of the error *depends on the error and not on the code word sent*. *The syndrome is the column where the error occurs*. Thus the columns of the parity-check matrix M are all the syndromes that can appear for any single error.

To generalize for *any* parity-check code, we can write a parity-check matrix M which has 1's in the places of each row where the corresponding parity check is applied and 0's where the parity check does not apply. We *assume* that the rank of this parity-check matrix M is equal to the number of rows in the matrix (using, of course, modulo 2 arithmetic). For example, the matrix (compare Section 3.4)

$$M = \begin{pmatrix} 1 & 1 & 1 & 1 \\ 0 & 0 & 1 & 1 \\ 1 & 1 & 0 & 0 \end{pmatrix}$$

has only rank 2, since if the bottom two checks are satisfied, the upper one is also satisfied. (The sum of all three rows is 0 0 0 0.) If the rows are not linearly independent, then at least one syndrome digit is determined from the others, and we are not going to get all the different syndromes we want—we will waste signaling capacity.

The encoding process *in the general parity-check case* assures us that

$$Mc^T = 0$$

for all legitimate code words. Thus once more, the error e gives rise to the received message $c + e$ and

$$M(c + e)^T = Me^T = \text{syndrome}$$

The syndrome again depends only on the error and is the corresponding column of the matrix.

If there is a single error, then you have the following decoding rule: The syndrome is exactly the corresponding column of the parity check

matrix. Therefore, by comparing the syndrome with the columns of the matrix until you find a match, you can locate the position of the error.

Thus you see that for any single-error-correcting code, given any syndrome you can find the corresponding column of the check matrix M (the syndrome “all zeros” means no error).

If we are to have multiple-error-correcting codes, then the syndrome is exactly the sum (modulo 2) of the correcting codes’ columns of the matrix corresponding to the error positions. This correspondence *must be unique* if you are to find the proper columns of the matrix and hence the corresponding multiple errors that induced the syndrome. If a syndrome does occur, and it is not the sum of the permitted number of columns, then there is nothing you can do to correct the errors that have occurred. You can only detect the presence of some errors.

To be specific, in a double-error-correcting code any possible syndrome that can arise from two errors must be the sum of a *unique pair* of columns of the matrix—there cannot be two different pairs of columns leading to the same syndrome, since if there were, you would have no way of knowing which pair of errors gave rise to the observed syndrome. Nor can a sum of any two columns be some other single column. In symbols you have

$$M(e_1 + e_2)^T = Me_1^T + Me_2^T = \text{syndrome}$$

and the syndrome must be unique for *every pair* of errors e_1 and e_2 , as well as for every single error (and cannot be all 0’s, since that means no error).

The same sort of argument applies to higher error-correcting codes.

Exercise

11.3-1 Discuss the problem of triple-error-correcting code design.

11.4 Double-Error-Detecting Codes Revisited

For the double-error-detecting code of Section 3.7, we have the matrix (for $n = 8$) where we have added the overall parity check as the top row of the matrix,

$$M = \begin{pmatrix} 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 0 & 0 & 0 & 1 & 1 & 1 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 & 1 & 0 \end{pmatrix}$$

Any syndrome with the first component equal to 1 must have the rest of the components as the syndrome of a single error, with the syndrome all 0's meaning that the additional parity check position is wrong (see last column of the matrix M). If the first component of the syndrome is 0, there is either zero or some other even number, 2, 4, 6, 8, . . . , of errors. In the first case the message is correct, and in the second case you do not know what to do since the syndrome is not a column of the matrix M . You can merely note the presence of an even number of errors for this code.

11.5 Polynomials Versus Vectors

Originally (as we did in Chapters 2, 3, and 10) the code sequences of 0's and 1's were regarded as vectors. The significant step forward occurs when this view is replaced by the view that they are *polynomials*—the digits of the code word are to be regarded as the coefficients of the powers of the variable (indeterminate) x . Thus the sequence

$$1, 0, 1, 1, 0$$

is now to be viewed as the polynomial

$$P_1(x) = x^4 + x^2 + x$$

In a sense we are using a *generator function* representation of the code sequence of 0's and 1's.

Another example is the sequence 0, 1, 1, 0, 1, which corresponds to the polynomial

$$P_2(x) = x^3 + x^2 + 1$$

The set of all polynomials of a given degree with integer coefficients permits addition and subtraction (which in our case of modulo 2 arithmetic are the same thing!) of any two of them. For example, using the two polynomials above, $P_1(x) + P_2(x) = x^4 + x^3 + x + 1$. Addition is *associative*, $(P_1 + P_2) + P_3 = P_1 + (P_2 + P_3)$. The zero element (0, 0, 0, 0, 0) corresponds to the zero polynomial and is identically zero. Since each nonzero polynomial is its own additive inverse, and the sum of any two polynomials is another polynomial of the same degree (at most), it follows that the set of all polynomials of degree n forms a *group*. Remember that the coefficients of the polynomials can be only 0 or 1.

In Section 2.8 we examined modular arithmetic. One of the purposes of modular arithmetic is to be sure that all the numbers that can arise will fall within the given range of numbers. This is accomplished in modular arithmetic by dividing all numbers by the modulus m and using only the m positive remainders $0, 1, \dots, (m - 1)$, which are all less than the modulus.

When we face the class of polynomials we find that the *product* of two polynomials has the degree that is the sum of the degrees of the individual polynomials. To keep the product within the maximum degree less than n (so that the polynomial will be identifiable with a code sequence), we must similarly reduce the degree. Thus we need a polynomial to use as a modulus. Just as with modular arithmetic where we divide all numbers by the modulus and identify the remainder as the number, so we will now divide all polynomials by the modulus polynomial of degree n and use the remainder as the product. As a result, the remainder will have n coefficients (the degree is $n - 1$), which we can then identify as a code symbol.

We saw in Section 2.8 why we want a prime number as our modulus. For a prime modulus, if the product of two numbers is zero, then at least one of the two numbers is zero. If the modulus is not a prime, it is possible that neither number is zero, although their product is. Similarly, we want a prime polynomial as our modulus—a prime polynomial being one that cannot be represented as the product (in the field of coefficients) of two nontrivial polynomials.

A polynomial of degree n might by chance have its coefficient of the highest power 0 so that the polynomial was really of lower degree. Since we have only 0 and 1 in the field of coefficients, we will use the standard words *monic polynomial* to mean a polynomial whose coefficient of the highest power of x is 1.

But the reader may well ask: Why do we want to multiply polynomials together? The answer is simply that we want to introduce as much structure into our theory as we can. Then, and only then, will the codes we finally get be highly structured, and as a result we can hope that the encoding and decoding equipment will also be highly structured (meaning easy to build in hardware or to program in software). In practice, we simply do not want to face random, or even lightly structured codes.

11.6 Prime Polynomials

A prime polynomial is a monic polynomial that cannot be factored into a product of lower-order polynomials. But since our field of coefficients uses arithmetic modulo 2, we cannot be sure of how old familiar results

will now appear. Thus we will experimentally explore what can happen—we will avoid going into all the technical details of finite fields to prove how factorization must occur.

For polynomials of degree zero there is only the single, trivial polynomial (if you care to call this degenerate case a polynomial), namely

$$P = 1$$

It corresponds in normal arithmetic to the number 1.

For monic polynomials of degree 1 there are only two polynomials,

$$\begin{array}{c} x \\ x + 1 \end{array}$$

both of which are prime polynomials (just as we disregard the use of 1 as a factor in factoring integers, we disregard the trivial polynomial 1 as a factor).

For monic polynomials of degree 2 we have four possibilities (each of the coefficients of x and 1 may be 0 or 1):

$$\begin{array}{ll} x^2 & = x \cdot x \\ x^2 + 1 & = (x + 1)(x + 1) \\ x^2 + x & = x(x + 1) \\ x^2 + x + 1 & = \text{prime} \end{array}$$

The fact that

$$x^2 + 1 = (x + 1)^2$$

may come as a surprise at first, but multiplying out the right-hand side, we get

$$x^2 + 2x + 1$$

Since $2 = 0 \pmod{2}$, we have

$$x^2 + 2x + 1 = x^2 + 1$$

Is the polynomial

$$x^2 + x + 1$$

a prime? Are you sure? We try dividing by the two lower-degree polynomials. If $x^2 + x + 1$ is composite, these are the only possible factors. Clearly, x is not a factor. We then try

$$\begin{array}{r}
 x \\
 x + 1 \overline{) x^2 + x + 1} \\
 \underline{x^2 + x} \\
 0 1 \\
 0 0 \\
 \underline{ 1} = \text{remainder}
 \end{array}$$

Yes, $x^2 + x + 1$ is a prime polynomial.

For further practice we examine the eight possible cubics

$$\begin{array}{ll}
 x^3 & = x \cdot x \cdot x \\
 x^3 + 1 & = \\
 x^3 + x & = x(x^2 + 1) \\
 x^3 + x + 1 & = \\
 x^3 + x^2 & = x^2(x + 1) \\
 x^3 + x^2 + 1 & = \\
 x^3 + x^2 + x & = x(x^2 + x + 1) \\
 x^3 + x^2 + x + 1 & =
 \end{array}$$

where we have given the four obvious factorizations, and left the others blank.

If a cubic can be factored, then one factor must be of first degree. We have given those divisible by x , so we need only try to factor $x + 1$. We have for $x^3 + 1$,

$$\begin{array}{r}
 \\
 1 \overline{) 1 0 1} \\
 \underline{1 1} \\
 1 0 \\
 \underline{1 1} \\
 1 1 \\
 \underline{1 1} \\
 0 = \text{remainder}
 \end{array}$$

We shall not go on finding the prime polynomials of higher and higher degrees. Evidently, for each particular polynomial we need only try lower-order prime polynomials up to half the degree of the given polynomial, and we are beginning to see how to do some of the factorizations in our head (see also Exercise 11.6-1). In any case, we are

not trying to show how to make the best codes but are content to show how coding theory develops toward higher error-correcting codes which have a good deal of structure in them, so that the terminal gear is practical to build in software or hardware.

The idea behind the proof that there are always prime polynomials of every degree goes as follows. We know that there are 2^n distinct polynomials of degree n . Given how many prime polynomials there are of each lower degree, we can compute how many distinct products of degree n we can form. The difference between these two numbers is the number of prime polynomials of degree n there must be. By suitable bounds on each number we can obtain a bound on the next number.

Exercises

11.6-1 Prove the theorem: If and only if the sum of the coefficients of a polynomial $= 0 \pmod{2}$, then $x + 1$ is a factor.

11.6-2 Carry out the factorization of polynomials of degree 4.

11.6-3 Carry out the factorization of polynomials of degree 5.

11.7 Primitive Roots

We need to make one more digression before returning to error-correcting codes. In normal algebra, the n th roots of unity are the roots of $x^n - 1 = 0$. Their explicit form is given by

$$e^{2\pi i k/n} \quad (k = 0, 1, 2, \dots, n - 1)$$

Note that $(e^{2\pi i k/n})^n = \cos 2\pi k + i \sin 2\pi k = 1$ and hence they are all roots of unity.

Some of these roots, for example $k = 1$, are such that their successive powers generate all the other roots. Indeed, if k is relatively prime to n , this will happen. To prove this result, take the successive powers, $m = 1, 2, \dots, n - 1, n$, of this root:

$$e^{2\pi i k/n}$$

These m th powers are

$$e^{2\pi i k m/n}$$

If any two of these were equal, then we would have

$$e^{2\pi i k m_1/n} = e^{2\pi i k m_2/n}$$

or

$$e^{2\pi i k (m_1 - m_2)/n} = 1$$

This means that

$$\frac{k(m_1 - m_2)}{n} = \text{some integer}$$

Since k is relatively prime to n , this means that they have no factors in common and n must divide $m_1 - m_2$. However, $m_1 - m_2$ is less than n , so $m_1 = m_2$. Therefore, if k is relatively prime to n , the successive powers of that root generate all the roots of unity and it is a *primitive root*.

The reason that primitive roots are important is that a primitive root can serve as a *generator* of the whole set of numbers we need. However, we must examine the problem in our modulus 2 arithmetic.

11.8 A Special Case

The perfect Hamming codes have $n = 2^m - 1$ digits. For the case $n = 7 = 2^3 - 1$, we have the matrix (Section 11.3)

$$M = \begin{pmatrix} 0 & 0 & 0 & 1 & 1 & 1 & 1 \\ 0 & 1 & 1 & 0 & 0 & 1 & 1 \\ 1 & 0 & 1 & 0 & 1 & 0 & 1 \end{pmatrix} \quad (11.8-1)$$

We know that any rearrangement of the columns of the matrix M will not fundamentally change the code; it will only affect how and where we find the column corresponding to the syndrome that emerges when a single error occurs.

To illustrate the new approach, suppose that we use the prime polynomial

$$x^3 + x + 1$$

as our modulus. Let a be a root of this polynomial. Thus

$$a^3 + a + 1 = 0$$

or

$$a^3 = a + 1$$

This means that the powers of a which are 3 or higher can be reduced, using this equation, to sums of lower powers. Thus every polynomial in a can be reduced to a linear combination of 1, a , and a^2 . Now consider a column 3-vector which has components 1, a , a^2 from the bottom up. The column vector gives the coefficients of the corresponding power of a . Note that multiplication by a shifts each component up one position, and that $a^3 = a + 1$ is used whenever a^3 occurs, so that only 1, a , and a^2 appear in the vector. For the various powers of a we have the corresponding 3-vectors

$$1 = \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}$$

$$a = \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix}$$

$$a^2 = \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix}$$

$$a^3 = a + 1 = \begin{pmatrix} 0 \\ 1 \\ 1 \end{pmatrix}$$

$$a^4 = a^2 + a = \begin{pmatrix} 1 \\ 1 \\ 0 \end{pmatrix}$$

$$a^5 = a^3 + a^2 = a^2 + a + 1 = \begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix}$$

$$a^6 = a^3 + a^2 + a = a^2 + 1 = \begin{pmatrix} 1 \\ 0 \\ 1 \end{pmatrix}$$

$$a^7 = a^3 + a = \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}$$

as we should. Thus we have shown that a is a primitive root of the prime polynomial, since its successive powers generate all seven possible different nonzero 3-tuples.

If we write the columns of this new matrix in the form (for convenience we are reversing the order of the columns),

$$\begin{matrix} a^6 & a^5 & a^4 & a^3 & a^2 & a & 1 \\ \begin{pmatrix} 1 & 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 & 1 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & 1 \end{pmatrix} \end{matrix}$$

we find that this is merely the matrix M (11.8-1) rearranged. In this new arrangement of the columns you compare the syndrome you compute at the receiving end with the successive powers of a (which you can either store or generate as needed). When you get a match between the syndrome and a column, you have that column as the corresponding power of a that matches. (A slightly different, but equivalent, decoding matrix would result if we used the prime polynomial $x^3 + x^2 + 1$.)

How do I encode? Since we want the syndrome 0, 0, . . . 0 to mean no error, I must require that the sent polynomial be exactly divisible by the modulus polynomial $x^3 + x + 1$. Thus I put the message in the positions corresponding to a^6 , a^5 , a^4 , a^3 , and temporarily put zeros in the other three positions corresponding to a^2 , a , and 1. When I divide this polynomial by the modulus polynomial, I will get a remainder. If this remainder is put into the last three positions of the polynomial I wrote (or if you wish is added to the polynomial), then the result will be exactly divisible by the modulus polynomial. This, then, is the encoding process—supply the last three positions by using the remainder upon division by the modulus polynomial; thus the whole polynomial is equal to zero modulo the prime polynomial.

To decode at the receiving end, you simply divide the received polynomial by the modulus polynomial. The syndrome (remainder) that results when it is written as a power of a gives the column where the error occurred. We now have a different way of viewing the earlier Hamming codes. This approach is suggestive of what to do next.

Example of Encoding. Given the code above, I encode the message

$$1 \quad 0 \quad 0 \quad 1 \quad - \quad - \quad -$$

as follows. The message corresponds to the polynomial

$$a^6 + a^3$$

This is the fourth position from the right (which is where the error occurred). You add a 1 in this fourth position to get the corrected message, and then drop the last three places to get the original message.

Exercises

11.8-1 Find the decoding matrix for the prime polynomial $x^3 + x^2 + 1$.

11.8-2 Repeat the encoding and decoding of the example for the message 1 0 1 0 and an error in the a^2 position.

11.8-3 Find the decoding matrix for $x^4 + x + 1$.

11.8-4 Find the decoding matrix for $x^4 + x^3 + 1$.

11.8-5 Show that the syndrome for two errors is the sum of the separate syndromes.

11.9 Shift Registers for Encoding

We have developed the theory of algebraic codes in terms of these codes because they are easy to realize in either hardware or software. The apparent complexity of much of the previous division process for the modular reduction of polynomials can be handled simply by means of *shift registers with feedback paths*. Suppose that we have the 7-bit polynomial code of Section 11.8, and I want to encode the message 1 0 0 1. The encoded message will be a polynomial of the form

$$1.x^6 + 0.x^5 + 0.x^4 + 1.x^3 + *x^2 + *x + *1$$

(where * indicates the coefficient to be determined) and is to be divisible by the modulus polynomial $x^3 + x + 1$ (the one we used in Section 11.8).

I will first set the three unknown coefficients equal to zero and, as before, divide by the modulus polynomial to get the remainder. When I add the remainder $a^2 + a$ to the original message polynomial $a^6 + a^3$, I will have the encoded message

$$a^6 + a^3 + a^2 + a = (a^3 + a + 1)(a^3 + a)$$

and this, as shown, is exactly divisible by the modulus polynomial $a^3 + a + 1$.

stages, since nothing of interest has happened during the first four shifts until the leading digit gets to the left-hand position. The diagram starts there at the top line. The next line shows the digit fed back and the third row (just above the double line) shows the result of the addition. (We show the adding of 1 in the left position to match the division process in detail; it is not fed back in practice.) The shift occurs at the double line, and we next see the shifted digits. The set of three lines between the double lines gives the details before the next shift. In total there are, in this case, four sets of three lines. The end result is the remainder. A comparison will show how the shift register does the same thing as the division process.

Figure 11.9-3 shows a practical encoder. The digits of the message still come from the right, but now they also go out on the left through the dotted switch. When the remainder is computed, I have all the original message digits sent out and I then change the position of the switch so that I can now shift out the remainder to form the entire coded message (omitting the first digit, which is, of course, always zero). The encoding is that simple!




Figure 11.9-3 Encoder

Shift registers are not hard to build in practice, even with the additional ability to logically add (exclusive OR) in each position. The rest is similarly easy to build in hardware.

In software it is also not difficult to program. I test for the 1 in the left position and if I find it, I logically add the pattern of 1's required by the feedback paths (the modulus polynomial). Then I shift and repeat. The details are not difficult to supply.

Thus the encoding for this code is fairly simple. We next turn to the decoding process.

Exercises

11.9-1 Write a program to encode a message following the model in this section.

11.9-2 Draw a schematic of a shift register $x^4 + x^3 + 1$.

11.10 Decoding Single-Error-Correcting Codes

At the receiving end you do the same as at the encoding end, *except* that you also have to store the message so that you can make the needed corrections when they occur.

After the division process you have the remainder—the syndrome of the error. The all-zero remainder means “the message is correct and therefore send on the first 4 bits.”

If the remainder is not zero, it is some power of a , say a^k . How can you find which power? A table lookup might be appropriate for a reasonably short code, but for a very complex one this might require too much storage [although the price of storage is currently (1985) falling very rapidly].

Looking back at Section 11.8 you see that $a^7 = 1$; hence $a^6 = 1/a$. You can compute from the syndrome (remainder), which is a^k :

$$\begin{aligned}(a^6)a^k &= a^{k-1} && \text{modulo } (a^3 + a + 1) \\(a^6)a^{k-1} &= a^{k-2} \\&\vdots\end{aligned}$$

until you recognize $0 \ 0 \ 1 = 1$ as the remainder. You then know k .

For example, suppose that I had the message of Section 11.9,

$$1 \ 0 \ 0 \ 1 \ 1 \ 1 \ 0$$

and an error occurred in the a^2 term. Thus you receive

$$1 \ 0 \ 0 \ 1 \ \underline{0} \ 1 \ 0$$

The division process will go as follows (simplified):

$$\begin{array}{r}
 1 \ 0 \ 0 \ 1 \ 0 \ 1 \ 0 \\
 \underline{1 \ 0 \ 1 \ 1} \\
 1 \ 0 \ 0 \ 1 \\
 \underline{1 \ 0 \ 1 \ 1} \\
 1 \ 0 \ 0 \longrightarrow \text{remainder}
 \end{array}$$

Therefore, you start with this remainder, which is some unknown power of a , say the k th.

$$a^k = \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix}$$

Multiply by $a^6 = 1/a$. This is equivalent to shifting each component down one position. When a 1 is shifted out the bottom, you use $1/a = 1 + a^2$. Thus you have

$$a^6 a^k = a^{k-1} = \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix}$$

until $a^0 = 1$ is recognized:

$$a^6 a^{k-1} = a^{k-2} = \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix} = 1$$

In this case $k = 2$ —you need to add a 1 in the a^2 position.

If you shifted out *backwards* you could use the powers of a as a test, and when the match occurred add 1 to the digit being shifted out.

You could also start with the highest power of a and compute down until you get a match. In this case you could shift out forwards.

Notice that in software this needs to be done *only* if there is an error—otherwise nothing need be done beyond getting the syndrome and testing it for all 0's. In real-time coding systems, you must have the capacity to do the computing whether or not it is used.

11.11 A Double-Error-Correcting Code

To partially illustrate a double-error-correcting code, we need to go to a higher Hamming code with four parity checks and 15 total positions.

We examine the polynomial

$$x^4 + x + 1$$

as a possible candidate for the modulus polynomial. Clearly, x is not a factor of the polynomial. If the polynomial is to be divisible by $x + 1$, the sum of the coefficients must be 0 (by Exercise 11.6-1, a small modification of the old factoring rule still applies here). We have, therefore, only to try as the factor the single irreducible quadratic $x^2 + x + 1$ to find that we have a prime polynomial.

Next, we need to test to see if it meets the primitive condition. We compute the successive powers of a :

$$1 = \begin{pmatrix} 0 \\ 0 \\ 0 \\ 1 \end{pmatrix}$$

$$a = \begin{pmatrix} 0 \\ 0 \\ 1 \\ 0 \end{pmatrix}$$

$$a^2 = \begin{pmatrix} 0 \\ 1 \\ 0 \\ 0 \end{pmatrix}$$

etc.

We find that $a^{15} = 1$, as it should, and no earlier power equals 1. Thus it is a primitive polynomial and we have the parity-check matrix:

$$\begin{matrix} a^{14} & a^{13} & a^{12} & a^{11} & a^{10} & a^9 & a^8 & a^7 & a^6 & a^5 & a^4 & a^3 & a^2 & a & 1 \\ \begin{pmatrix} 1 & 1 & 1 & 1 & 0 & 1 & 0 & 1 & 1 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 1 & 1 & 1 & 1 & 0 & 1 & 0 & 1 & 1 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 1 & 1 & 1 & 0 & 1 & 0 & 1 & 1 & 0 & 0 & 1 & 0 \\ 1 & 1 & 1 & 0 & 1 & 0 & 1 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 1 \end{pmatrix} \end{matrix}$$

If we want to correct more errors than this single-error-correcting code does, then we must add more rows to the matrix. We will add four more rows, but how?

At present, two different double errors can produce the same result for the syndrome, that is,

$$a^k + a^m = a^{k'} + a^{m'}$$

For example, columns a and a^2 give the same syndrome as columns a^{12} and a^{14} . We need more rows. How about using the square of a as a generator? Unfortunately, the two syndromes (let a_1 be some power of a and a_2 be some other power of a)

$$a_1 + a_2 = s_1 = \text{first syndrome}$$

and

$$a_1^2 + a_2^2 = s_2 = \text{second syndrome}$$

are simply related by the equation

$$s_1^2 = (a_1 + a_2)^2 = a_1^2 + 2a_1a_2 + a_2^2 = a_1^2 + a_2^2 = s_2$$

We next try using the cubes for the second syndrome,

$$a_1^3 + a_2^3 = s_2$$

We rearrange this:

$$\begin{aligned} s_2 &= a_1^3 + a_2^3 = (a_1 + a_2)(a_1^2 + a_1a_2 + a_2^2) \\ &= s_1(a_1^2 + a_1a_2 + a_2^2) \\ &= s_1(a_1a_2 + s_1^2) \end{aligned}$$

so that

$$a_1 + a_2 = s_1$$

$$a_1a_2 = s_1^2 + \frac{s_2}{s_1} \quad (\text{if } s_1 \neq 0)$$

In all this computing when we want to do a multiplication (or division) we represent the quantities as a power; hence we add (or subtract) the exponents; when we want to do an addition we use the column vector representations and add the corresponding components.

Thus we constantly go between the two forms of representation. The matrix itself gives the needed table for the conversions.

Given the sum and product of two numbers, we almost automatically think of the quadratic equation and its roots,

$$a^2 - (s_1)a + \left(s_1^2 + \frac{s_2}{s_1}\right) = 0$$

The two roots give the individual syndromes of the errors.

If there is only one error, then the syndromes are

$$a_1 = s_1$$

$$a_1^3 = s_2$$

and

$$s_1^2 + \frac{s_2}{s_1} = 2s_1^2 = 0$$

The quadratic equation becomes, therefore,

$$a^2 + s_1a = 0 \quad \text{or} \quad a = s_1$$

$$a = 0$$

For no errors, we have

$$s_1 = s_2 = 0$$

and the quadratic equation is

$$a^2 = 0 \quad a = 0$$

$$a = 0$$

Thus for the lower rows of the matrix we use a^3 rather than a as the generator of the columns. The columns are therefore those corresponding to

$$1 \quad a^3 \quad a^6 \quad a^9 \quad a^{12} \quad a^{15} \quad a^{18} \quad \dots$$

(in reverse order). We get the complete matrix M , therefore,

$$\begin{array}{cccccccccccccccc}
 & a^{14} & a^{13} & a^{12} & a^{11} & a^{10} & a^9 & a^8 & a^7 & a^6 & a^5 & a^4 & a^3 & a^2 & a^1 & a^0 \\
 \left(\begin{array}{cccccccccccccccc}
 1 & 1 & 1 & 1 & 0 & 1 & 0 & 1 & 1 & 0 & 0 & 1 & 0 & 0 & 0 \\
 0 & 1 & 1 & 1 & 1 & 0 & 1 & 0 & 1 & 1 & 0 & 0 & 1 & 0 & 0 \\
 0 & 0 & 1 & 1 & 1 & 1 & 0 & 1 & 0 & 1 & 1 & 0 & 0 & 1 & 0 \\
 1 & 1 & 1 & 0 & 1 & 0 & 1 & 1 & 0 & 0 & 1 & 0 & 0 & 0 & 1 \\
 \hline
 1 & 1 & 1 & 1 & 0 & 1 & 1 & 1 & 1 & 0 & 1 & 1 & 1 & 1 & 0 \\
 1 & 0 & 1 & 0 & 0 & 1 & 0 & 1 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\
 1 & 1 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 \\
 1 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 1
 \end{array} \right)
 \end{array}$$

as the parity-check matrix for double-error correction. The syndrome s_1 is from the upper rows and the syndrome s_2 is from the lower rows. The two syndromes, through the quadratic equation, give us the two roots necessary to locate the two errors.

How do we do the encoding? How to satisfy all the parity checks at the same time is not clear, but depends on the use of the famous Chinese remainder theorem—which it is not appropriate to go into in this very elementary introduction to multiple-error correction. The elementary nature of the book prevents us from going deeply into the more difficult parts of coding theory.

As a concrete example that is easy to follow (and contains the general case as well as evades the encoding process) take the valid code of all 0's. Next we assume that during transmission the 0's in positions corresponding to a^3 and a^4 are changed to 1's. The syndrome of the received message will be the sum of those two columns:

$$\begin{pmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 1 \\ 0 \\ 1 \\ 0 \end{pmatrix} + \begin{pmatrix} 0 \\ 0 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{pmatrix} = \begin{pmatrix} 1 \\ 0 \\ 1 \\ 1 \\ 0 \\ 1 \\ 0 \\ 1 \end{pmatrix} \quad \begin{array}{l} s_1 = \begin{pmatrix} 1 \\ 0 \\ 1 \\ 1 \end{pmatrix} = a^7 \\ s_2 = \begin{pmatrix} 0 \\ 1 \\ 0 \\ 1 \end{pmatrix} = a^8 \end{array}$$

We next calculate

$$s_1^2 + \frac{s_2}{s_1} = (a^7)^2 + \frac{a^8}{a^7} = a^{14} + a = \begin{pmatrix} 1 \\ 0 \\ 0 \\ 1 \end{pmatrix} + \begin{pmatrix} 0 \\ 0 \\ 1 \\ 0 \end{pmatrix} = \begin{pmatrix} 1 \\ 0 \\ 1 \\ 1 \end{pmatrix} = a^7$$

Hence the quadratic equation is

$$x^2 + a^7x + a^7 = 0$$

Try $x = 1$, $1 + a^7 + a^7 = 1$,

$$\begin{pmatrix} 0 \\ 0 \\ 0 \\ 1 \end{pmatrix} + \begin{pmatrix} 1 \\ 0 \\ 1 \\ 1 \end{pmatrix} + \begin{pmatrix} 1 \\ 0 \\ 1 \\ 1 \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \\ 0 \\ 1 \end{pmatrix} \neq 0$$

Try $x = a_j$;

$$a^2 + a^8 + a^7 = \begin{pmatrix} 0 \\ 1 \\ 0 \\ 0 \end{pmatrix} + \begin{pmatrix} 0 \\ 1 \\ 0 \\ 1 \end{pmatrix} + \begin{pmatrix} 1 \\ 0 \\ 1 \\ 1 \end{pmatrix} = \begin{pmatrix} 1 \\ 0 \\ 1 \\ 0 \end{pmatrix} \neq 0$$

Try $x = a^2$;

$$a^4 + a^9 + a^7 = \begin{pmatrix} 0 \\ 0 \\ 1 \\ 1 \end{pmatrix} + \begin{pmatrix} 1 \\ 0 \\ 1 \\ 0 \end{pmatrix} + \begin{pmatrix} 1 \\ 0 \\ 1 \\ 1 \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \\ 1 \\ 0 \end{pmatrix} \neq 0$$

Try $x = a^3$;

$$a^6 + a^{10} + a^7 = \begin{pmatrix} 1 \\ 1 \\ 0 \\ 0 \end{pmatrix} + \begin{pmatrix} 0 \\ 1 \\ 1 \\ 1 \end{pmatrix} + \begin{pmatrix} 1 \\ 0 \\ 1 \\ 1 \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \\ 0 \\ 0 \end{pmatrix} = 0$$

Try $x = a^4$:

$$a^8 + a^4 + a^7 = \begin{pmatrix} 0 \\ 1 \\ 0 \\ 1 \end{pmatrix} + \begin{pmatrix} 1 \\ 1 \\ 1 \\ 0 \end{pmatrix} + \begin{pmatrix} 1 \\ 0 \\ 1 \\ 1 \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \\ 0 \\ 0 \end{pmatrix} = 0$$

Try $x = a^5$;

$$a^{10} + a^{12} + a^7 = \begin{pmatrix} 0 \\ 1 \\ 1 \\ 1 \end{pmatrix} + \begin{pmatrix} 1 \\ 1 \\ 1 \\ 1 \end{pmatrix} + \begin{pmatrix} 1 \\ 0 \\ 1 \\ 1 \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \\ 1 \\ 1 \end{pmatrix} \neq 0$$

and so on. Hence the two errors are in positions a^3 and a^4 .

You can pick any two errors or a single error (or even no error) and carry out the computation in detail to see how the system works. The powers of a are given in the top four rows of the matrix.

11.12 Decoding Multiple-Error Codes

For multiple-error-correcting codes there is a simple decoding rule. You might think that you must effectively, as in Section 11.11, turn over the syndrome problem to a computer. For a double-error-correcting code, the computer must find the zeros of the quadratic. One very simple way to find the roots is to evaluate the quadratic at each position of the stored message as it is being shifted out. If the power of a satisfies the quadratic, then it is a root and is one of the positions to be corrected.

Evidently, you need do this only when there are errors. But in real-time signaling you need to store several messages at the same time. When such a system is in operation, you will be shifting out and correcting one message, computing the corrections of a second message, and working on the first stages of a third. Thus at a minimum there will be two message time delays. This can be important in the feedback paths of control circuits! It could cause instability problems.

11.13 Summary

In this chapter we have given the argument for why we want a highly structured code: so that we can then encode and decode using systematic methods. We have shown how the division process of dividing one of our polynomials by another polynomial can be implemented easily.

We showed how the syndrome is uniquely tied to the error and is independent of the message. We next got down to the necessity of using polynomials modulo a prime (nonfactorable) polynomial. We then saw that the additional condition that the polynomial have a primitive root is necessary so that all the elements can be written as powers

of it—sort of going to logarithms. Finally, we illustrated the methods by devising one particular double-error-correcting code.

We did not attempt to provide all the theory of finite fields necessary to justify all that can be done, nor to create a general theory; we merely showed how a code could be found. The general codes are presented so that special cases can be recognized, or even developed if necessary; more than this would make the book far too long and would not add much to those devoted to algebraic coding theory that are already available (see the References).

It is necessary to clearly state that almost all of coding theory was developed in the days when computing was relatively cheap and storage was more expensive. In the year 1985, we are seeing a great decrease in the cost of storage, and hence there is a need to reconsider all the theory that we have developed in the past. Methods of design, encoding, and decoding which depend more on table lookup and other forms of storage and less on computing need to be researched and developed. Again, you need to be warned that the theory is almost all designed to meet white noise, and in practice often the noise is not white and the theoretical benefits of a code may therefore not be realized in practice.

Bandwidth and the Sampling Theorem

A.1 Introduction

This appendix appears to have only a slight connection with the earlier material; certainly, the mathematical tools used are completely different. One reason for including it is that in practice the channel matrix (Section 7.2)

$$(P_{i,j})$$

is rarely available. But more important, to send digital information we normally use analog channels that are almost exactly linear. The capacity of analog signal channels is typically measured by “bandwidth”; therefore, it seems necessary in a book on coding and information theory to at least glance at the idea of *bandwidth* and related topics such as *power*.

The eigenfunctions (the proper functions) of a linear, time-invariant system are the complex exponentials

$$e^{2\pi i f t}$$

Eigenfunction means that the system does not change the form of the function, only the amplitude. *Time-invariant* means that the behavior of the system is not a function of time but only of the inputs. This is characteristic of most current communication systems. *Linear* means

that the sum of two functions put into the channel gives the sum of the individual outputs as the output. This is characteristic of most channels *provided* that we avoid overloading them. The frequency f is measured in rotations (when we do the calculus we use the angular frequency $\omega = 2\pi f$). The independent variable is t , usually time.

A *band-limited function* is a function whose frequencies are all limited to some band. A familiar example is a hi-fi system, which typically will have a lower limit at best of a few cycles per second, say 15 or maybe even a hundred, and an upper limit of around 18,000 cycles per second. International convention uses the symbol Hz (hertz) for “cycles per second.” Radio and television stations are limited to certain bands of frequencies so they do not interfere with each other’s broadcasts. Many other systems have a natural limit to their band of frequencies, although in most all cases the band is not as sharply cut off as the theory indicates.

A.2 The Fourier Integral

The Fourier integral is the natural mathematical tool for handling band-limited functions. The Fourier integral represents a function $g(t)$ as a sum of complex exponentials (the eigenfunctions) in the form

$$g(t) = \int_{-\infty}^{\infty} G(f)e^{2\pi ift} df$$

The function under the integral $G(f)$ is called the *transform* of $g(t)$. It is related to the original function by the interesting formula

$$G(f) = \int_{-\infty}^{\infty} g(t)e^{-2\pi ift} dt$$

which is the same as the original one except for the sign on the i ; one integral has a “ $+i$ ” and the other has a “ $-i$ ” term in the exponent.

The Fourier integral is like a glass prism that breaks up a light beam into its colors—frequencies. The function $G(f)$ gives the amount (amplitude) of the frequency f in the original time signal $g(t)$. The integral is simply the properly weighted sum of all possible frequencies; $G(f)$ is the weight at frequency f .

There are positive and negative frequencies to correspond to the pair of functions $\cos 2\pi ft$, $\sin 2\pi ft$, which have the exponential repre-

sensation

$$\cos 2\pi ft = \frac{e^{2\pi ift} + e^{-2\pi ift}}{2}$$

$$\sin 2\pi ft = \frac{e^{2\pi ift} - e^{-2\pi ift}}{2i}$$




Figure A.2-1 Band-limited function

As an example of a band-limited function, suppose that $G(f)$ is zero outside the band (Figure A.2-1),

$$-F \leq f \leq F$$

Thus we have

$$g(t) = \int_{-F}^F G(f) e^{2\pi ift} df$$

Suppose further that inside the band the function $G(f)$ is a constant $1/2F$, chosen so that the area under the function $G(f)$ is exactly 1. We have, therefore,

$$\begin{aligned} g(t) &= \int_{-F}^F \frac{1}{2F} e^{2\pi ift} df \\ &= \frac{1}{2F} \frac{e^{2\pi ift}}{2\pi it} \Big|_{-F}^F \\ &= \frac{e^{2\pi i Ft} - e^{-2\pi i Ft}}{2i} \frac{1}{2\pi Ft} \\ &= \frac{\sin 2\pi Ft}{2\pi Ft} \end{aligned}$$

This is the classic function $(\sin x)/x$ shown in Figure A.2-2. It has a *main lobe* about zero of width $2/F$ and *side lobes* of half that width which decrease in size like $1/t$. The larger F is (the wider the band), the narrower is the center lobe and the side lobes as well. This suggests that a function can peak up suddenly only if the function has a wide bandwidth.




Figure A.2-2 $f(x) = \frac{\sin x}{x}$

A.3 The Sampling Theorem

The famous *sampling theorem* connects the rate of sampling with the bandwidth that is necessary to reconstruct the function from the samples. Before stating the theorem, we need to remind readers of some of their experiences. When viewing a movie or a TV program, what you see is a sequence of still pictures flashed so frequently that they appear to be moving; typically, about 15 to 20 times a second is enough to fool the human eye. However, there are clues to the sampling. For example, when the stagecoach in the “Western” starts moving, the wheels appear first to speed up, then to slow down, stop, go backward, stop, go forward, and so on, depending on the speed. The reason is easy to see (Figure A.3-1). When the stage coach wheel has turned so fast that one spoke advances to the position of the next spoke, they will appear to be standing still. Indeed, when the spoke gets merely to half the position of the next spoke during one time interval between samples, the wheel will appear psychologically to stand still and to have twice the number of spokes that it really has. At a rate slightly faster than this, the spokes will appear to be going backward.

A.4 Bandwidth Versus Rapid Change

From the preceding section we see that the bandwidth of a signal is inversely proportional to the width of the main lobe of the function ($\sin 2\pi Ft)/(2\pi Ft)$ —it takes a wide band of frequencies to get a narrow spike. We are not in a position (lacking a great deal of mathematical theory) to prove that no shape of the amplitude $G(f)$ in the band can give a narrower spike than the rectangular shape, so we are reduced to making suggestive observations as to why this is so.

It is not rapid change itself that is impossible for a narrow-band function; we have only to select any particular frequency, other than zero frequency, and take a sufficiently large an amplitude to have as large a change as we want in as short a time as we want. Rather, we want a change from one extreme to another, from a maximum to a minimum, or the reverse. This means that at the ends of the interval we want a horizontal tangent. Thus we have, essentially, a $\cos 2\pi Ft$ ($0 \leq 2\pi Ft \leq \pi$), and the transition takes place in time $1/2F$. Again we have the same reciprocal relationship between frequency and rate of change of the signal.

The sampling theorem, of course, is saying the same thing; we must have the necessary bandwidth, and this requires at least two samples in the highest frequency present in order to reconstruct the original signal (assuming that the samples we have go from the time $-\infty$ to $+\infty$).

A.5 AM Signaling

Most people know at least a little about AM (amplitude modulation) radios. The frequency of broadcast that you tune the radio set to is called the “carrier frequency” and is a constant frequency carefully controlled by the sending station. All the radio station does is change the amplitude of the frequency they send, and of course in turn that of the signal you receive (see Figure A.5-1). This, however, amounts to a small change in frequency, because the product of two cosines is a pair of cosines each shifted in frequency from one of the original ones (see a few lines later).

The radio receiver, by tuning, filters out all the frequencies that are broadcast except those close to the carrier frequency. Then in one way or another the received signal is “rectified”—meaning that the negative parts are either cut out or else are actually reversed in sign and are made positive. When the result of this is “smoothed” a bit, it is, approximately, the original modulation curve that the station imposed on its carrier.




Figure A.5-1

How does this involve the bandwidth? Suppose that the musical note to be broadcast is the pure sinusoid

$$\cos 2\pi ft$$

and that the carrier frequency wave is

$$\cos 2\pi f_0 t$$

where f_0 is very much larger than f . The modulated signal that is sent is the product of the two cosines:

$$(\cos 2\pi ft)(\cos 2\pi f_0 t) = \frac{1}{2} \cos 2\pi(f_0 + f)t + \frac{1}{2} \cos 2\pi(f_0 - f)t$$

Enough information is contained in either of the two terms to reconstruct the original—provided that the carrier frequency f_0 is known (you select the carrier frequency you want when you select the station by your tuning). Thus there are “single-sideband” signaling systems which broadcast only one of the two terms and suppress the other.

Inspection of this formula shows that to receive a band of frequencies of the original signal, you need to receive the same bandwidth (but shifted up by the carrier frequency). Thus it is bandwidth that is necessary to transmit the information—and the wider the band available, the more different signals you can transmit. Put simply, large bandwidth gives high fidelity, and lack of bandwidth will suppress many frequencies and give low fidelity.

The same is true in telephone conversations, although there is no carrier in the simplest cases. The telephone design allows the transmission of frequencies up to around 3600 Hz and pretty well suppresses those above 3600 Hz. This bandwidth gives medium fidelity. Higher-quality phone conversations (“quality” in tone, not in content!) would require wider bandwidth, and it is a fact of nature that bandwidth costs money and effort to achieve. The bandwidth originally chosen for telephones was an engineering compromise *at that time* between quality and cost, and it is not easily upgraded when the installed physical plant is so large.

A.6 FM Signaling

Signaling using AM, because of the fading and other properties it has, is not as good as FM (frequency modulation). In AM the amplitude of the transmitted signal is modulated; in FM the frequency is modu-

lated, meaning that the information is sent by directly varying the frequency. By using automatic gain control at the receiving end, the amplitude of the signal is kept approximately constant in spite of fading.

If the receiver has a filter (tuner) that lets less and less of the signal through when the frequency is farther and farther away from the carrier frequency (typically, this is much higher than the AM carrier frequency), then the energy (or power, if you prefer to measure things as energy per unit time) changes correspondingly. Thus loosely speaking, FM works with the signal

$$A \cos 2\pi[f_0 + f(t)]t$$

where $f(t)$ is small with respect to the f_0 of the carrier. Again it is bandwidth that is necessary to convey the information.

By using a much wider band than that used in AM, FM has much more ability to suppress noise and more faithfully reproduce the original signal. For the same bandwidth both AM and FM are equally effective.

A.7 Pulse Signaling

Typically, *digital signals* are sent as one of two states. One of two frequencies can be used, or one of two voltages (pulses) to signal. In the case of pulses, the corresponding representation in the characteristic modes (the eigenfunctions) of the linear, time-invariant signaling system,

$$e^{2\pi if t}$$

is given by a Fourier integral. But a linear system can give various gain and phase shifts to each different frequency. Therefore, the original pulse is apt to have quite a different shape at the detecting end. Further distortion comes from the inevitable “noise” of all real systems.

The wider the bandwidth, the more rapidly the system can change, and the more rapidly pulses can be sent down the system. Not that the pulses move faster, but they are spaced more densely in time. The inevitable distortion of the signaling system must not be allowed to grow too large before detection, reshaping, amplification, and retransmission. By spacing the repeaters sufficiently closely together, the degradation of the original pulse is kept small and hence the reliability is made high.

Sometimes it is the *change of state* of the system that indicates a 1 and no change a 0, rather than pulses for 1's and no (or negative) pulses for 0's.

There are many elaborate systems for pulse signaling; especially difficult are the problems of timing synchronization, and these lead to involved systems. Error detection and correction are often incorporated into the pulse system.

A.8 Bandwidth Generally

This appendix is intended to familiarize the reader with several aspects of channel capacity as equivalent to bandwidth. Without the use of a formidable amount of mathematics associated with the Fourier integral, the results can only be sketched. The aim of the appendix is to indicate that in some sense channel capacity as defined in information theory is related to bandwidth, which is the way channel capacity is usually rated in practice.

From Shannon's theorem we would expect that very sophisticated methods of encoding information could achieve great improvements in noise suppression without going to wide bandwidths. It is the complexity of the encoding and decoding that has stopped such efforts on a wide scale, such as for radio and TV. As we have seen, good encoding methods involve very long messages, and this has both severe storage problems and delays to contend with, let alone the processing of the signal. Thus practical signaling systems rarely come near to the ideal of Shannon's theorem. But the appearance of cheap microcomputers and cheap digital storage devices may change matters in the near future.

In TV it is necessary to achieve all the possible changes at each spot on the screen from one picture frame to the next, and this requires very wide bandwidths. But psychological tests indicate that human beings cannot *process* much more than 40 to 50 bits per second, if that much. Even allowing for the fact that various people seeing the same picture will process different parts of the picture, still the disparity between what is done and what seems to be the minimum necessary is almost a millionfold. It is in this way that Shannon's theory is most useful—it points to places where there are possibly large gains and steers the engineer away from places where there can be only small improvements at best.

There is only so much bandwidth in the spectrum of available frequencies that will pass through the earth's atmosphere, and already much of it is assigned to various uses. In a sense bandwidth is a "natural resource" that is definitely limited, not in the sense that the use from moment to moment consumes it, but in the sense that various competing uses of it cannot be made at the same time. From information theory we have a guide to where the bandwidth is used efficiently and where it is being wasted.

Because of the availability of cheap, reliable microcomputers, digital signaling methods that permit easy protection against noise (errors) are of increasing importance, and we have, therefore, neglected analog signaling systems in this brief discussion.

A.9 Continuous Signals

In actual practice, we usually send continuous (analog) signals (Section A.1), even when we have already encoded the message into binary digits. Thus it is necessary to look briefly at the information contained in an analog signal.

Most signaling systems use band-limited signals (Section A.2). Nature comes fairly close to this in many situations, although the absolutely sharp cutoff of the model does not occur in practice. Given a band-limited function of bandwidth W , the famous sampling theorem (Section A.3) says that it takes $2WT$ samples to reconstruct the signal of duration T . We assume that this theorem is true.

If we are sending the information in a continuous form of signals, each of duration T , each signal is represented by exactly $2WT$ numbers x_n taken at equally spaced intervals. These $2WT$ samples, numbers, can be viewed as a point in a $2WT$ -dimensional space, where we are now thinking of the space as continuous between the points.

It is natural to associate the corresponding energy E with

$$E = \frac{1}{2W} \sum_{n=1}^{2WT} x_n^2 \quad (\text{A.9-1})$$

The distance (radius) from the origin to this signal (using Euclidean distance) is

$$r = \left(\sum x_n^2 \right)^{1/2}$$

Hence we have

$$r = (2WE)^{1/2}$$

Since the signal power S is the energy E per unit time, we have

$$S = \frac{E}{T}$$

and the radius of the sphere of signals being sent is

$$r = (2WST)^{1/2} \quad (\text{A.9-2})$$

The noise added by the channel is also measured by its power, this time labeled N for “noise power.” Thus each sent signal is disturbed by the noise and at the receiving end lies in a sphere about the sent signal of radius

$$(2WNT)^{1/2} \quad (\text{A.9-3})$$

The whole (original) sphere (A.9-2) must be enlarged for this added noise and now has a radius of

$$r = [2WT(S + N)]^{1/2} \quad (\text{A.9-4})$$

As before (Section 3.6), we ask: How many spheres of radius of the noise (A.9-3) can fit in the enlarged sphere (A.9-4) of received messages? We divide the total space of received messages up into spheres of radius $= (2WNT)^{1/2}$ about the sent signals and get an upper bound on the total number of possible messages M that will be distinct at the received end. Using the formula for the volume of a $2WT$ -dimensional sphere, we get the bound

$$\begin{aligned} M &\leq \frac{\pi^{WT}[2WT(S + N)]^{WT}}{\Gamma(WT + 1)} \frac{\Gamma(WT + 1)}{\pi^{WT}(2WNT)^{WT}} \\ &\leq \left(1 + \frac{S}{N}\right)^{WT} \end{aligned}$$

The amount of information sent is the log of the number of equally likely sent messages; hence

$$I = \log M = WT \log \left(1 + \frac{S}{N}\right)$$

and the *rate* of information being sent per unit time is

$$W \log \left(1 + \frac{S}{N}\right)$$

This is the standard formula in the usual “signal-to-noise” form.

We have found an upper bound for the capacity of a channel with

bandwidth W and corresponding powers in both the signal and noise. It is much more difficult to get a lower bound since we will have to repeat the essentials of Shannon's main theorem and show that we can indeed pack in that many spheres (less a small amount that still leaves us arbitrarily close to capacity). The mathematical details in Chapter 9 and the proof of Shannon's theorem in Chapter 10 were both arranged so that this fact would seem to be probably true without proof.

Some Tables for Entropy Calculations

Five-place Tables of the Functions:
 $\log_2 (1/p)$, $p \log_2 (1/p)$, $H_2(p)$ for p 0.00(0.01)1.00

p	$\log (1/p)$	$p \log (1/p)$	$H(p)$
0.00	—	0	0
0.01	6.64386	0.06644	0.08079
0.02	5.64386	0.11288	0.14144
0.03	5.05889	0.15177	0.19439
0.04	4.64386	0.18575	0.24229
0.05	4.32193	0.21610	0.28640
0.06	4.05889	0.24353	0.32744
0.07	3.83650	0.26856	0.36592
0.08	3.64386	0.29151	0.40218
0.09	3.47393	0.31265	0.43647
0.10	3.32193	0.33219	0.46900
0.11	3.18442	0.35029	0.49992
0.12	3.05889	0.36707	0.52936
0.13	2.94342	0.38264	0.55744
0.14	2.83650	0.39711	0.58424
0.15	2.73697	0.41054	0.60984
0.16	2.64386	0.42302	0.63431
0.17	2.55639	0.43459	0.65770
0.18	2.47393	0.44531	0.68008
0.19	2.39593	0.45523	0.70147

p	$\log (1/p)$	$p \log (1/p)$	$H(p)$
0.20	2.32193	0.46439	0.72193
0.21	2.25154	0.47282	0.74148
0.22	2.18442	0.48057	0.76017
0.23	2.12029	0.48767	0.77801
0.24	2.05889	0.49413	0.79504
0.25	2.00000	0.50000	0.81128
0.26	1.94342	0.50529	0.82675
0.27	1.88897	0.51002	0.84146
0.28	1.83650	0.51422	0.85545
0.29	1.78588	0.51790	0.86872
0.30	1.73697	0.52109	0.88129
0.31	1.68966	0.52379	0.89317
0.32	1.64386	0.52603	0.90438
0.33	1.59946	0.52782	0.91493
0.34	1.55639	0.52917	0.92482
0.35	1.51457	0.53010	0.93407
0.36	1.47393	0.53062	0.94268
0.37	1.43440	0.53073	0.95067
0.38	1.39593	0.53045	0.95804
0.39	1.35845	0.52980	0.96480
0.40	1.32193	0.52877	0.97095
0.41	1.28630	0.52738	0.97650
0.42	1.25154	0.52565	0.98145
0.43	1.21754	0.52356	0.98582
0.44	1.18442	0.52115	0.98959
0.45	1.15200	0.51840	0.99277
0.46	1.12029	0.51534	0.99538
0.47	1.08927	0.51196	0.99740
0.48	1.05889	0.50827	0.99885
0.49	1.02915	0.50428	0.99971
0.50	1.00000	0.50000	1.00000
0.51	0.97143	0.49543	0.99971
0.52	0.94342	0.49058	0.99885
0.53	0.91594	0.48545	0.99740
0.54	0.88897	0.48004	0.99538
0.55	0.86250	0.47437	0.99277
0.56	0.83650	0.46844	0.98959
0.57	0.81097	0.46225	0.98582
0.58	0.78588	0.45581	0.98145
0.59	0.76121	0.44912	0.97650
0.60	0.73697	0.44218	0.97095
0.61	0.71312	0.43500	0.96480
0.62	0.68966	0.42759	0.95804
0.63	0.66658	0.41994	0.95067
0.64	0.64386	0.41207	0.94268

p	$\log (1/p)$	$p \log (1/p)$	$H(p)$
0.65	0.62149	0.40397	0.93407
0.66	0.59946	0.39564	0.92482
0.67	0.57777	0.38710	0.91493
0.68	0.55639	0.37835	0.90438
0.69	0.53533	0.36938	0.89317
0.70	0.51457	0.36020	0.88129
0.71	0.49411	0.35082	0.86872
0.72	0.47393	0.34123	0.85545
0.73	0.45403	0.33144	0.84146
0.74	0.43440	0.32146	0.82675
0.75	0.41504	0.31128	0.81128
0.76	0.39593	0.30091	0.79504
0.77	0.37707	0.29034	0.77801
0.78	0.35845	0.27959	0.76017
0.79	0.34008	0.26866	0.74148
0.80	0.32193	0.25754	0.72193
0.81	0.30401	0.24625	0.70147
0.82	0.28630	0.23477	0.68008
0.83	0.26882	0.22312	0.65770
0.84	0.25154	0.21129	0.63431
0.85	0.23447	0.19930	0.60984
0.86	0.21759	0.18713	0.58424
0.87	0.20091	0.17479	0.55744
0.88	0.18442	0.16229	0.52936
0.89	0.16812	0.14963	0.49992
0.90	0.15200	0.13680	0.46900
0.91	0.13606	0.12382	0.43647
0.92	0.12029	0.11067	0.40218
0.93	0.10470	0.09737	0.36592
0.94	0.08927	0.08391	0.32744
0.95	0.07400	0.07030	0.28640
0.96	0.05889	0.05654	0.24229
0.97	0.04394	0.04263	0.19439
0.98	0.02915	0.02856	0.14144
0.99	0.01450	0.01435	0.08079
1.00	0.00000	0.00000	0.00000

References

- [A] ABRAMSON, N. *Information Theory and Coding*. New York: McGraw-Hill, 1963.
- [B1] BERLEKAMP, E. R. *Algebraic Coding Theory*. New York: McGraw-Hill, 1968.
- [B2] BERLEKAMP, E. R., ed. *Key Papers in the Development of Coding Theory*. New York: IEEE Press, 1974.
- [Bl] BLAHUT, R. E. *Theory and Practice of Error-Control Codes*. Reading, MA: Addison-Wesley, 1983.
- [Bla] BLAKE, I., ed. *Selected Papers on Algebraic Coding Theory*. Stroudsburg, PA: Dowden, Hutchinson & Ross, 1973.
- [BM] BLAKE, J. F., and MULLIN, R. C. *An Introduction to Algebraic and Combinatorial Coding Theory*. New York: Academic Press, 1976.
- [C] CLARK, G. C., and CAIN, J. B. *Error-Correcting Coding for Digital Communications*. New York: Plenum, 1981.
- [DG] DAVISSON, L. D., and GRAY, R. M. *Data Compression*. Stroudsburg, PA: Dowden, Hutchinson & Ross, 1976.
- [G] GALLAGER, R. G. *Information Theory and Reliable Communication*. New York: John Wiley, 1968.
- [Ga] GATLIN, L. L. *Information Theory and the Living System*. New York: Columbia University Press, 1972.

- [Gu] GUIASU, S. *Information Theory with Applications*. New York: McGraw-Hill, 1977.
- [J] JAYANT, N. S. *Waveform Quantization and Coding*. New York: IEEE Press, 1976.
- [K] KNUTH, D. E. *The Art of Computer Programming*, Vol. 3: *Sorting and Searching*. Reading, MA: Addison-Wesley, 1973.
- [L1] LIDL, R., and NIEDERRETIER, H. *Finite Fields*. Reading, MA: Addison-Wesley.
- [L2] LIN, S., and COSTELLO, D. J., JR. *Error-Control Coding: Fundamentals and Applications*. Englewood Cliffs, N.J.: Prentice-Hall, 1983.
- [Mc] MCELIECE, R. J. *The Theory of Information and Coding*. Reading, MA: Addison-Wesley, 1977.
- [ME] MARTIN, N. F. G., and ENGLAND, J. W. *Mathematical Theory of Entropy*. Reading, MA: Addison-Wesley, 1981.
- [MS] MACWILLIAMS, J., and SLOANE, N. J. A. *The Theory of Error-Correcting Codes*. Amsterdam: North-Holland, 1977. (American edition: Elsevier, New York.)
- [P] PLESS, V. *Introduction to the Theory of Error-Correcting Codes*. New York: John Wiley, 1982.
- [PW] PETERSON, W. W., and WELDON, E. J., JR. *Error Correcting Codes*, 2nd ed. Cambridge, MA: MIT Press, 1972.
- [S] SLEPIAN, D., ed. *Key Papers in the Development of Information Theory*. New York: IEEE Press, 1974.
- [T] THOMPSON, T. M. *From Error-Correcting Codes through Sphere Packing to Simple Groups*. Washington, DC: Mathematical Association of America, 1983.
- [V] VITERBI, A. J., and OMURA, J. K. *Principles of Digital Communication and Coding*. New York: McGraw-Hill, 1979.
- [W] WAKERLY, J. *Error-Detecting Codes, Self-Checking Circuits and Applications*. Amsterdam: North-Holland, 1978.
- [WJ] WOZENCRAFT, J. M., and JACOBS, I. M. *Principles of Communication Engineering*. New York: John Wiley, 1965.

Index

A

Abramson, N., 253
Abstract theory, 1
Adjoint system, 133
Algebraic coding theory, 209
AM signaling, 242
AND, 30
Anti-Gray code, 99
ASCII code, 8, 10, 11, 27
 table of, 13
Average code length, 2, 63
Average error probability, 201
Average random code, 201

B

Backward conditional probabilities, 142
Bandwidth, 237
 example, 238
 generally, 246
 vs. change, 242
Bayes' formula (theorem), 141
Berlekamp, E.R., 253
Bernoulli trials, 183
Binary encoding, 7
Binary symmetric channel, 18, 142, 159, 195

 capacity of, 160-61

Binomial coefficient bound, 170-73

Binomial coefficients, 13

Bit, 7

 of information, 106

Blake, I., 253

Block code, 68

Burst error-detecting codes, 26

Byte, 10

C

Cain, F.B., 253

Carrier frequency, 224

Cauchy functional equation, 104

Channel, 138

 definition of, 139

 encoding, 2

 properties of, 140

 uniform, 154-55

Channel capacity, 103, 201

 of binary symmetric channel, 160-61

 definition of, 153

 encoding, 5

Channel relationships, 140

Chebyshev's inequality, 181-82

Clark, G., 253

Code compression, 18

Codes:

- anti-Gray, 99
- burst-detecting, 26
- comma, 54, 59, 69
- delta modulation, 100
- error-correcting, 34
- error-detecting, 20
- extensions of, 72
- Gray, 80, 97
- hash, 94
- hexadecimal, 10
- Huffman, 51
- instantaneous, 53-55
- miscellaneous, 79
- Morse, 12-13
- octal, 9
- predictive, 79
- radix r , 15
- random, 79
- resynchronization of, 101
- Shannon-Fano, 121-22, 126
- 2-out-of-5, 14, 23
- unique, 52
- van Duuren (3-out-of-7), 23
- variable-length, 12, 51
- weighted, 27

Coding:

- Huffman, 51
- Shannon-Fano, 121

Combinations, 13

- Comma codes, 54, 69

- Conditional mutual information, 150, 161

- Conditional probability, 80

- Continuous signals, 247

- Costello, D. J., 254

D

- Davidson, L.D., 253

- Decision rule, 192

- Decision tree, 53

- Decoding tree, 54, 56-57, 60, 62

- Delta modulation, 100

- Digital signaling, 6, 237

- Direct product, 53

- Double-error-correcting code, 228-34

- Double-error-detecting code, 47

E

- Efficient codes, 51, 56, 64

- Markov, 86

- Eigenfunctions, 237

- England, J., 254

- Entropy, 103, 107

- conditional, 130, 146

- joint, 146

- of Markov processes, 129

- received symbols, 145

- robustness, 136

- of source, 18

- source alphabet, 145

- Entropy function, 1, 107

- fundamental inequality, 115

- maximum of, 117

- properties of, 114-19

- table of, 250

- Equilibrium probabilities, 133

- Equivalent codes, 42

- Equivocation, 148

- Ergodic Markov processes, 85

- Error correction, 34

- Error-correcting codes, 35-36, 158

- perfect, 158

- Error-detecting codes, 20, 22, 147, 210

- probability, 24

- why, 20

- Error integral, 174

- Errors of interchange of symbols, 28

- Escape characters, 15

- Evolution, 3

- Excess redundancy, 22

- Extensions of a code, 53, 72, 124

- Markov process, 87

- examples, 126

F

- Fano bound, 205

- Finite automation, 53, 77

- Finite fields, 210

- FM signaling, 244

- FORTTRAN, 16, 49, 95

- compile, 17

- Forward conditional probabilities, 141

Fourier integral, 238
 Fundamental error formula, 199

G

Gallager, R.G., 253
 Galois fields, 4, 195
 Gamma function, 173
 Gatlin, L.L., 253
 Geometric approach, 44
 Gibbs inequality, 117
 Gray codes, 80, 97
 decoding, 99
 details of, 98
 Gray, R.M., 253
 code, 97
 Guiasu, S., 254

H

Hamming, distance, 45, 189
 meaning of, 46
 Hamming error-correcting codes, 39,
 210-11, 213, 220
 example, 41
 parity checks, 39, 210
 table of, 40
 Hamming-Huffman coding, 78
 Hartley, 106
 Hash coding, 94
 collisions, 95
 removals, 96
 Hexadecimal code, 10
 History, 2
 Huffman codes, 51, 63, 71, 114
 examples of, 65, 67-68
 noise in, 74
 radix r , 73
 use of, 77
 variability of, 68

I

Independent errors, 23
 Information source, 1, 5
 definition of, 104
 Instantaneous codes, 53
 construction of, 55

Interchange of symbols, 28
 ISBN book numbers, 32

J

Jacobs, I.M., 254
 Jayant, N.S., 254

K

Knuth, D.E., 254
 Kraft inequality, 57, 62, 121
 theorem, 57

L

Law of large numbers, 183
 theorem, 184
 Lin, S., 254
 Logical add (XOR), 23, 31

M

MacWilliams, J., 254
 Markov process, 80, 129, 131, 135
 ergodic, 84
 Martin, N., 254
 Maximum likelihood rule, 192, 196
 McEliece, R.J., 254
 McMillan inequality, 62
 Metalanguage, 16
 Memory source, 80
 Midpoint formula, 166
 Minimum distance, 45
 Minimum variance codes, 68, 75
 Model of signaling system, 4
 Modular arithmetic, 30
 Modulation, 242
 Modulo 2 arithmetic, 21
 Monic polynomial, 215
 Morse code, 12, 13, 51, 63
 table of, 13
 Mutual information, 148
 definition of, 148
 properties of, 149-51

N

nat (unit of information), 106
 n -dimensional codes, 37, 45

n -dimensional paradox, 179
 n -dimensional space, 176
 Euclidean, 176
 distance, 176
 n -dimensional spheres, 177
 coefficients, 177
 formula for, 177
 table of, 177
 Noise, 23
 entropy, 148
 in Huffman codes, 74
 Noise power, 248
 Noiseless coding theorem, 104, 125
 NORC computer, 34

O

Octal code, 9
 Omura, J.K., 254

P

Parity check, 21, 40
 computing, 22
 even, 21
 odd, 21
 Peterson, W.W., 254
 Pless, V., 254
 Polynomials, 214
 prime, 215
 primitive root, 219
 Predictive encoding, 79, 88
 encoder, 89
 decoder, 90
 run length, 91
 summary, 94
 Prime polynomials, 215
 Primitive roots, 219
 Probabilities of symbols, 1, 63
 Pulse signaling, 245

R

Radix r codes, 15, 72
 Random code, 196, 203
 average of, 201
 Random encoding, 79, 196

Random type fonts, 118
 Rectangular codes, 35
 Redundancy, 22, 36
 Reliable, 20
 Restrictions on codes, 21
 Resynchronization, 101
 Retransmission, 25

S

Sampling theorem, 240
 Single-error-detecting plus double-error
 correcting codes, 46, 228
 Shannon, 2
 Shannon's theorems:
 main, 191
 converse, 207
 general case of, 204
 noiseless, 125
 Shift registers, 224
 Shortened block codes, 60
 Side lobes, 240
 Slepian, D., 254
 Sloane, N.J.A., 254
 Source alphabet, 2, 6
 Source encoding, 2, 5
 Stirling's approximation, 165-70
 table of, 170
 STRETCH computer, 35
 Symbol, 1
 Symmetric channel, 18
 Syndrome, 38, 39, 212

T

Thompson, T.M., 254
 Transition graph, 81
 equilibrium, 83
 limiting matrix, 83
 matrix, 82
 probability, 82
 Trapezoid rule, 165
 Triangular codes, 37
 Two-out-of-five codes, 14, 23
 table, 14
 Two state devices, 6

U

Uniform channel, 154-55
 definition, 155
Uniform input, 156
Unique decoding, 52, 62

V

van Duuren code (3 out of 7), 23
Variable length codes, 13, 51
Variance, 181
 minimum, 75
 of sums, 184
Viterbi, A.J., 254

W

Wakerly, J. 254
Wallis integrals, 168
Weighted codes, 27
Weldon, E.J., Jr., 254
White noise, 23, 39
Wozencraft, J.M., 254

X

XOR (instruction), 30

Z

Zipf distribution, 71
 entropy of, 112